

# Microarchitectural Data Leakage

---

*via Automated Attack Synthesis*

Daniel Moghimi, Worcester Polytechnic University

Jun 23, 2020

Virtual Talk for Intel Product Security Incident Team (IPSIRT)

# About Me

- Daniel Moghimi (@danielmgmi)
- Security Researcher
- PhD Student @ WPI
  - Microarchitectural Security
  - Side Channels
  - Breaking Crypto Implementations
  - Trusted Execution Environment (Intel SGX)
- Contributed to:
  - ZombieLoad, Fallout, LVI,
  - MemJam, Spoiler, CacheZoom, CopyCat
  - Jackhammer, TPM-Fail



# Thanks...

- Berk Sunar @ WPI
- Moritz Lipp @ tugraz
- Michael Schwartz @ tugraz

# Disclaimers

- Our findings and reasonings are based on:
  - RE
  - Patents
  - Analysis
- You may know more than me how Intel CPU works!!!

# Today's Agenda

- Motivation: Meltdown-style Attacks
- Background: CPU Memory Subsystem
- Transynther, Automated Attack Synthesis
- MDS Root Cause Analysis and new subvariants
- Medusa attack and RSA key recovery

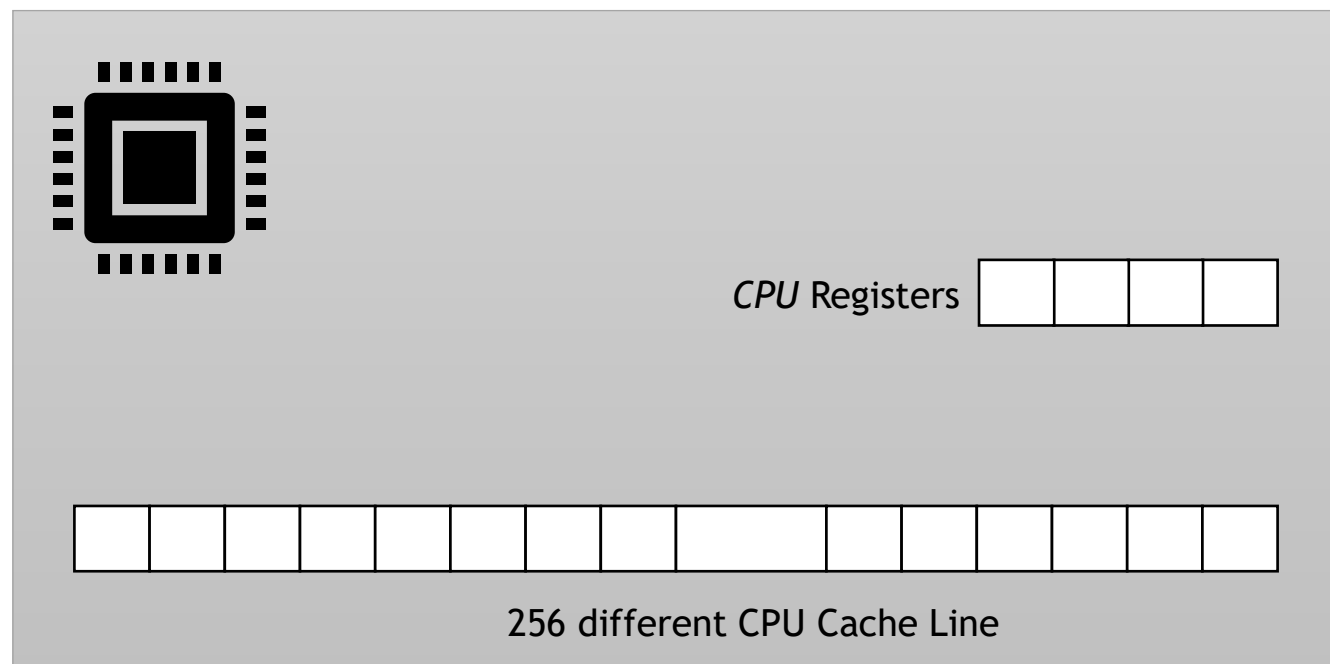

# 2018: Meltdown Attack?

```
char secret = *(char *) 0xffffffff81a0123;  
printf("%c\n", secret);
```




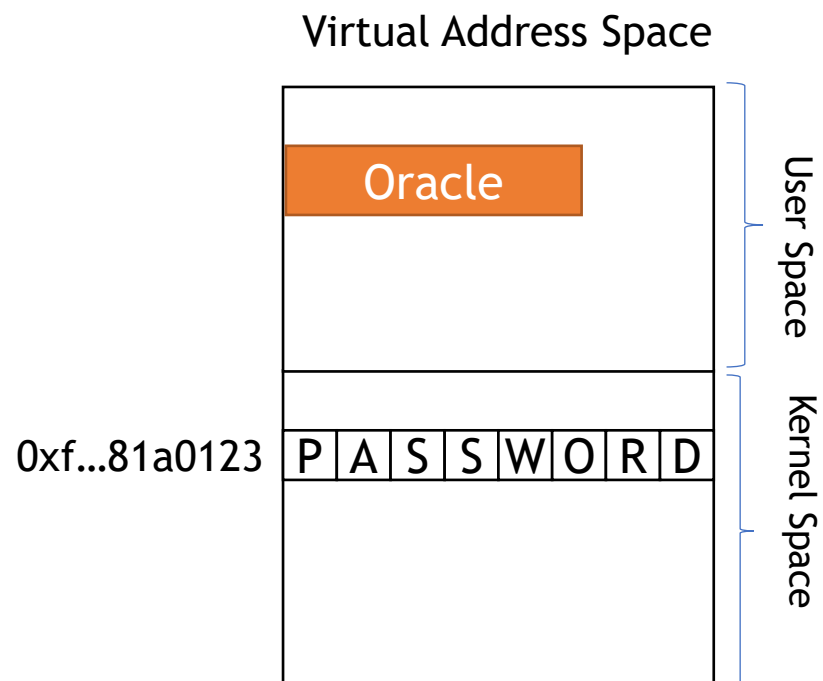
# 2018: Meltdown Attack?

```
char secret = *(char *) 0xffffffff81a0123;
```



# 2018: Meltdown Attack?


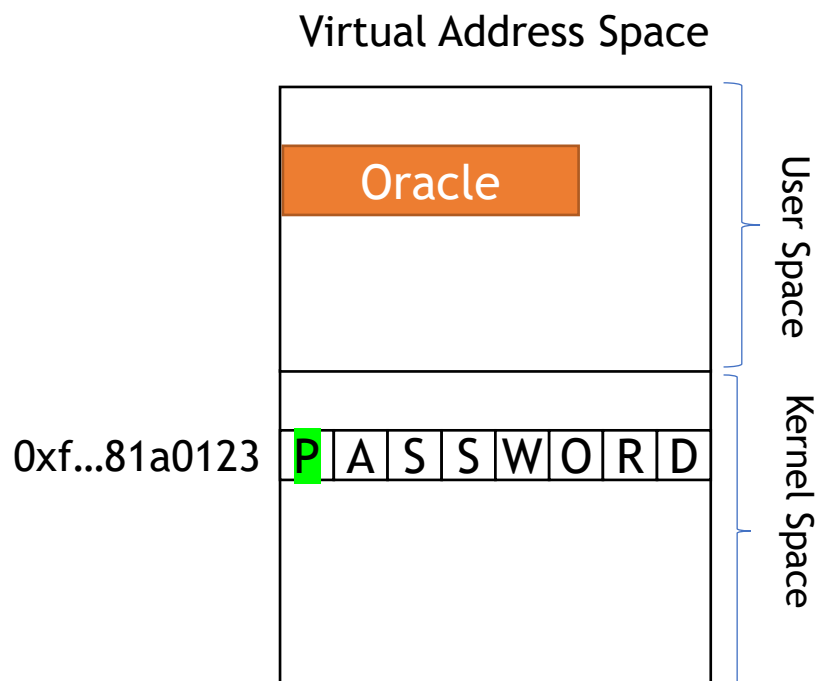
```
char secret = *(char *) 0xffffffff81a0123;
```





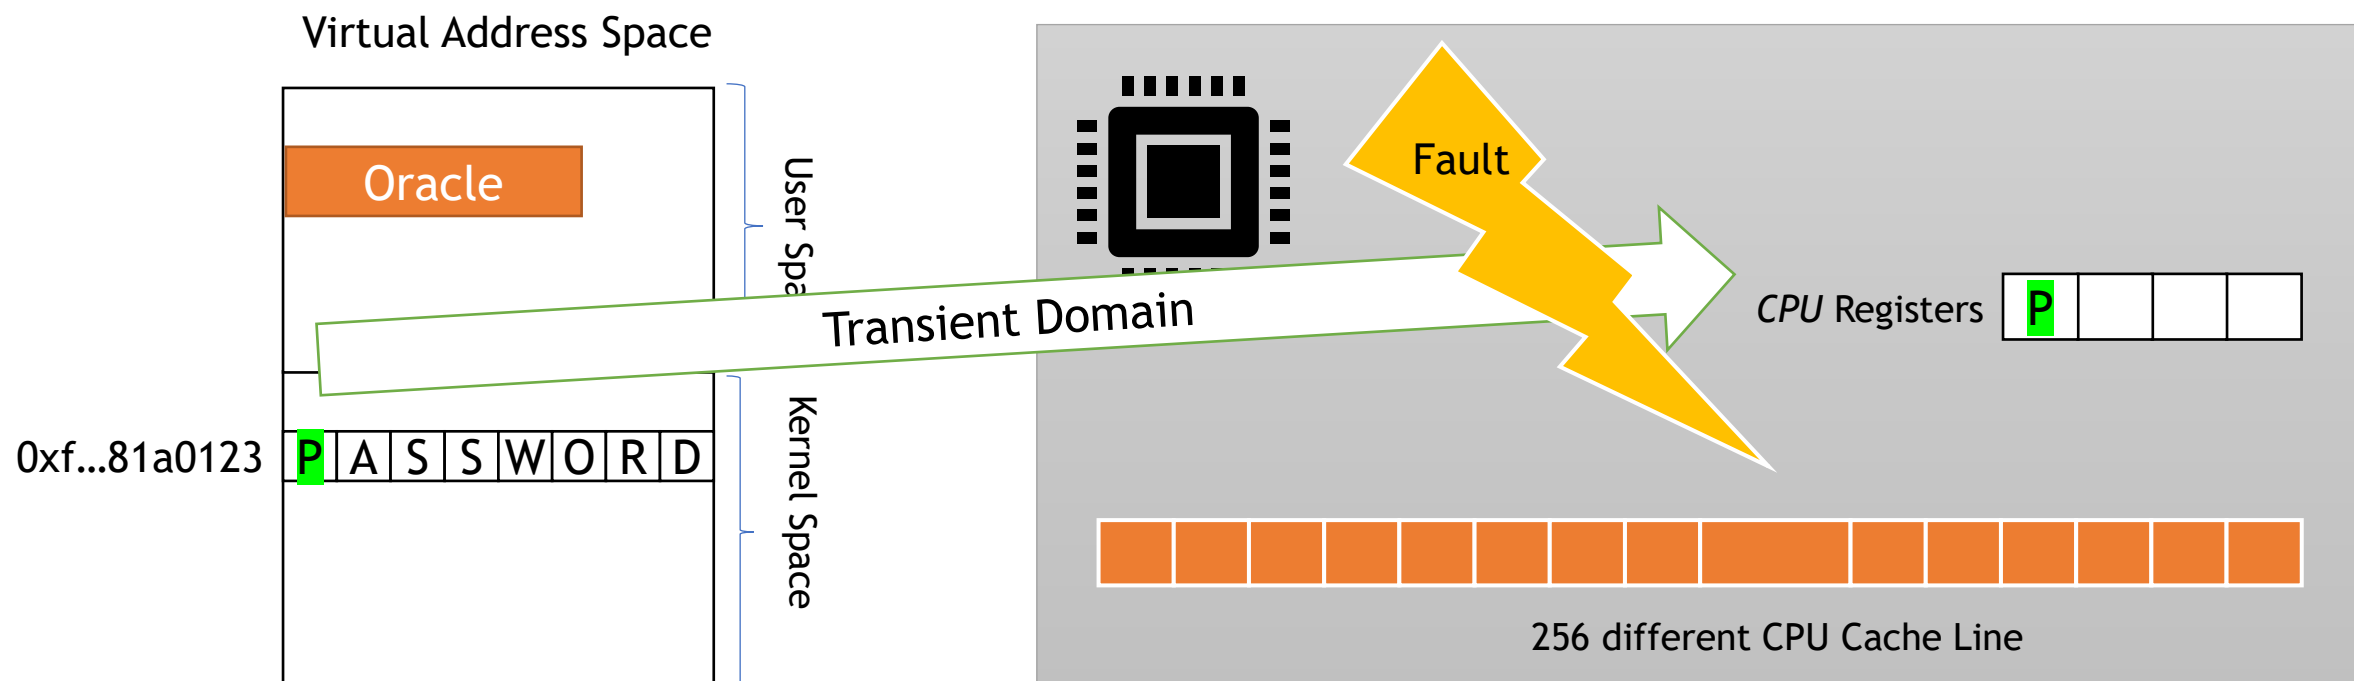
# 2018: Meltdown Attack? (Step 1)

➔ `char secret = *(char *) 0xffffffff81a0123;`



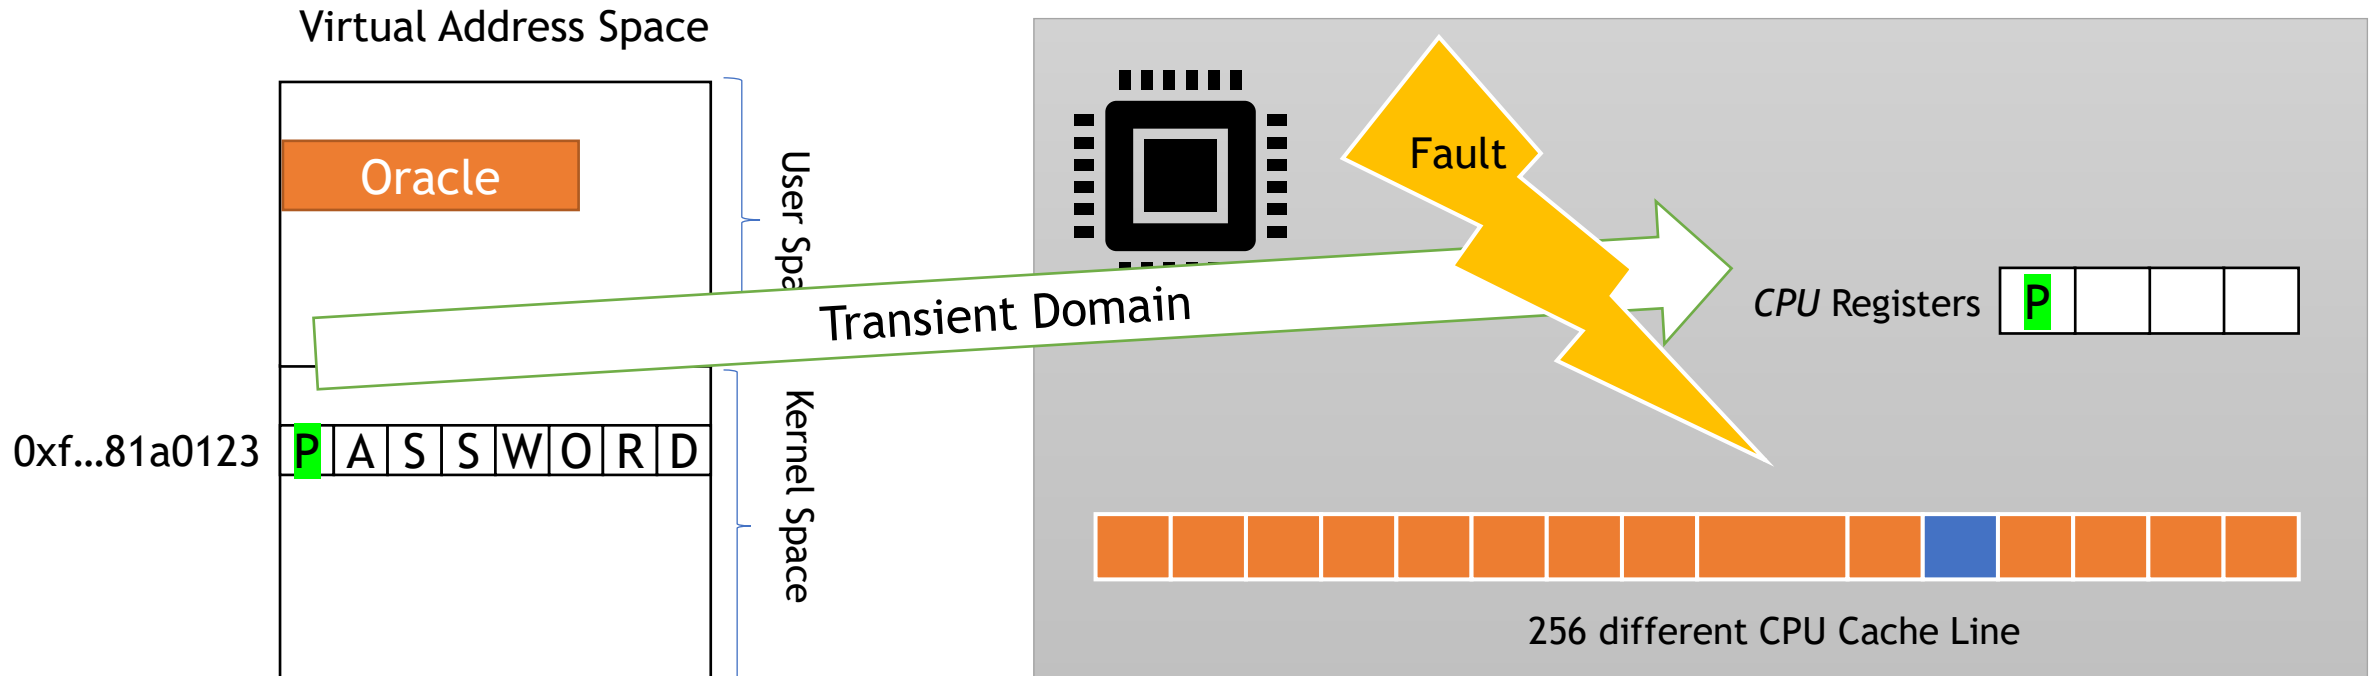
# 2018: Meltdown Attack? (Step 1)

➡ `char secret = *(char *) 0xffffffff81a0123;`




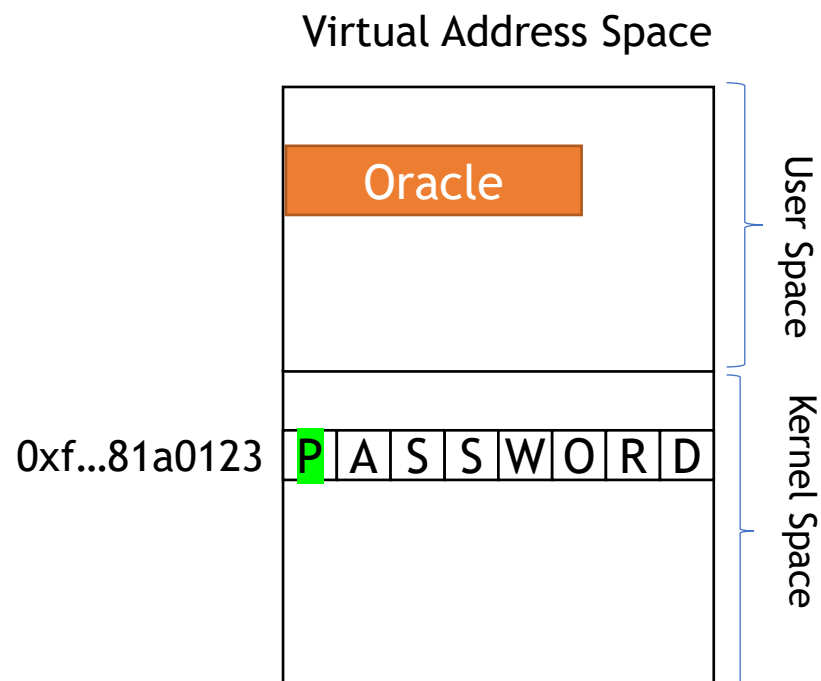
# 2018: Meltdown Attack? (Step 2)

➔ `char secret = *(char *) 0xffffffff81a0123;`  
`char x = oracle[secret * 4096];`





# 2018: Meltdown Attack? (Step 2)

```
char secret = *(char *) 0xffffffff81a0123;  
char x = oracle[secret * 4096];
```





# 2018: Meltdown Attack? (Step 3)

```
char secret = *(char *) 0xffffffff81a0123;  
char x = oracle[secret * 4096];
```





# 2018: Meltdown Attack? (Step 3)

```
char secret = *(char *) 0xffffffff81a0123;  
char x = oracle[secret * 4096];
```



# 2018: Meltdown Attack? (Step 3)

```
char secret = *(char *) 0xffffffff81a0123;  
char x = oracle[secret * 4096];
```


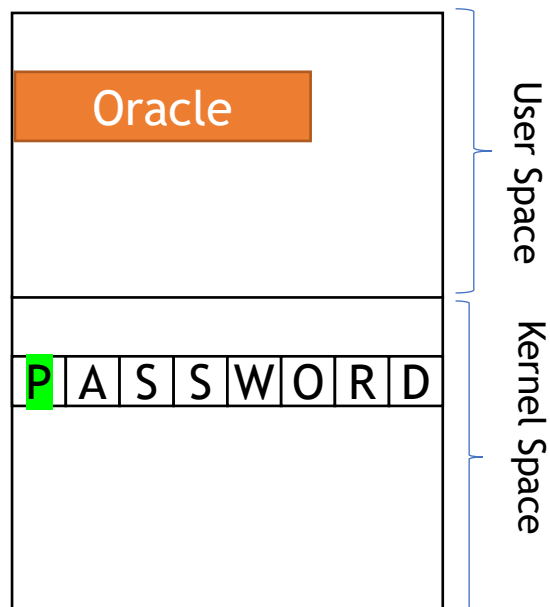


# 2018: Meltdown Attack? (Step 3)

```
char secret = *(char *) 0xffffffff81a0123;  
char x = oracle[secret * 4096];
```



Virtual Address Space





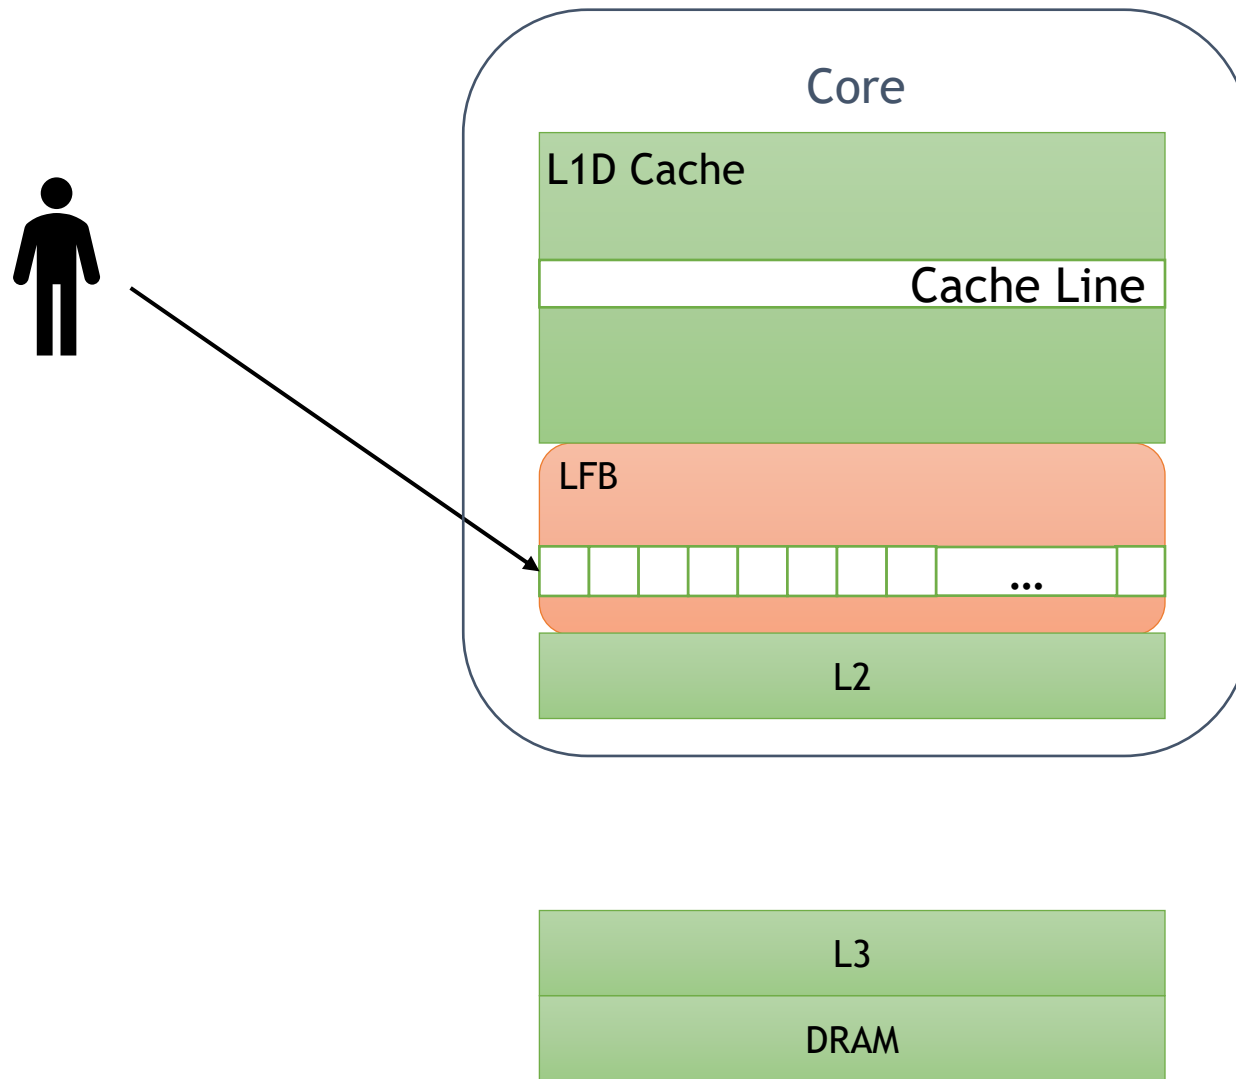
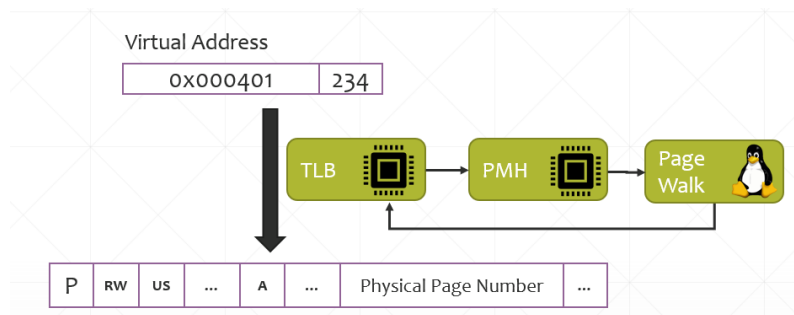
# Microarchitecture Data Sampling (MDS)

- Meltdown is fixed but you can still leak on the fix hardware.
- Which part of the CPU leak the data?!
- **Why does it leak?**

# ZombieLoad Attack

```
mov 0x401234, %rsi
```



```
mov (%rsi), %rax
```



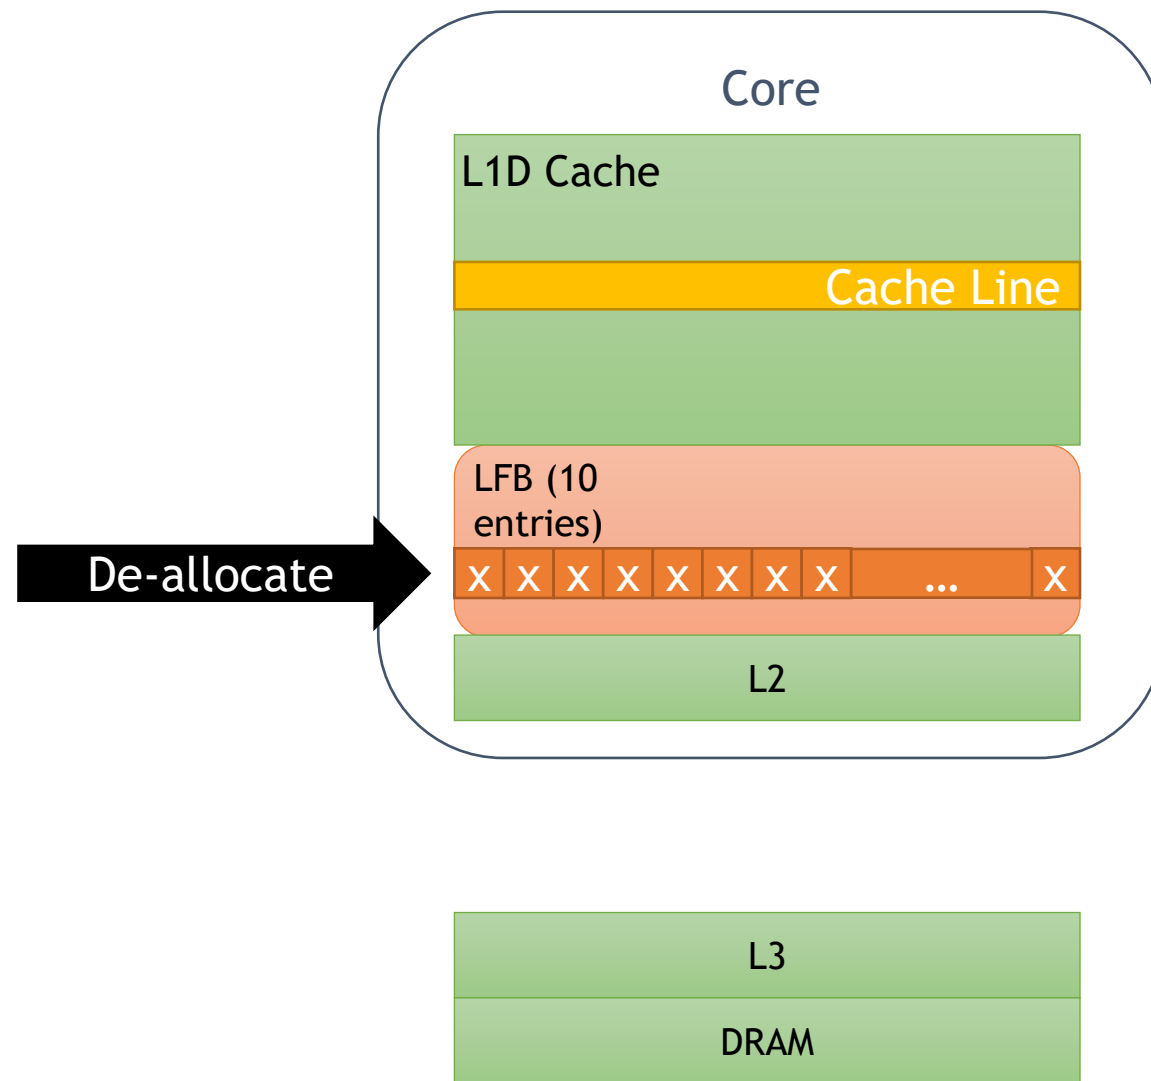
# ZombieLoad Attack

```
mov 0x401234, %rsi
```


```
mov (%rsi), %rax
```



# ZombieLoad Attack



# ZombieLoad Attack

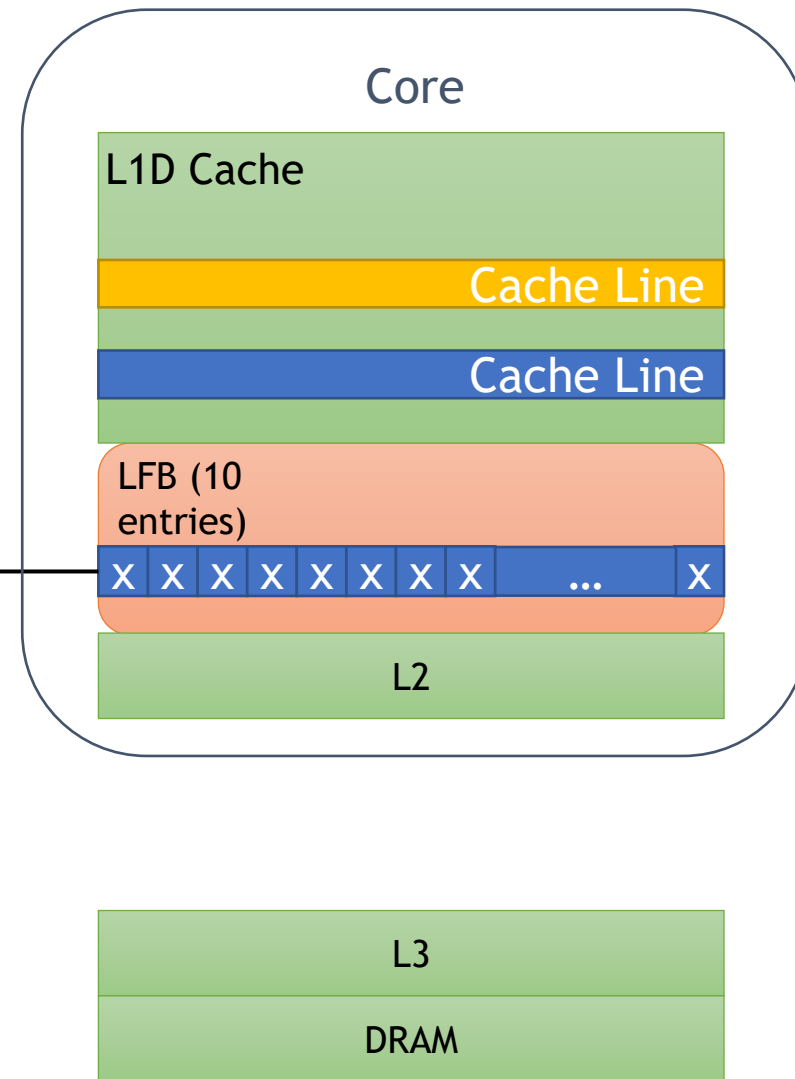


Variant 2: #RTM  
(AKA TAA)

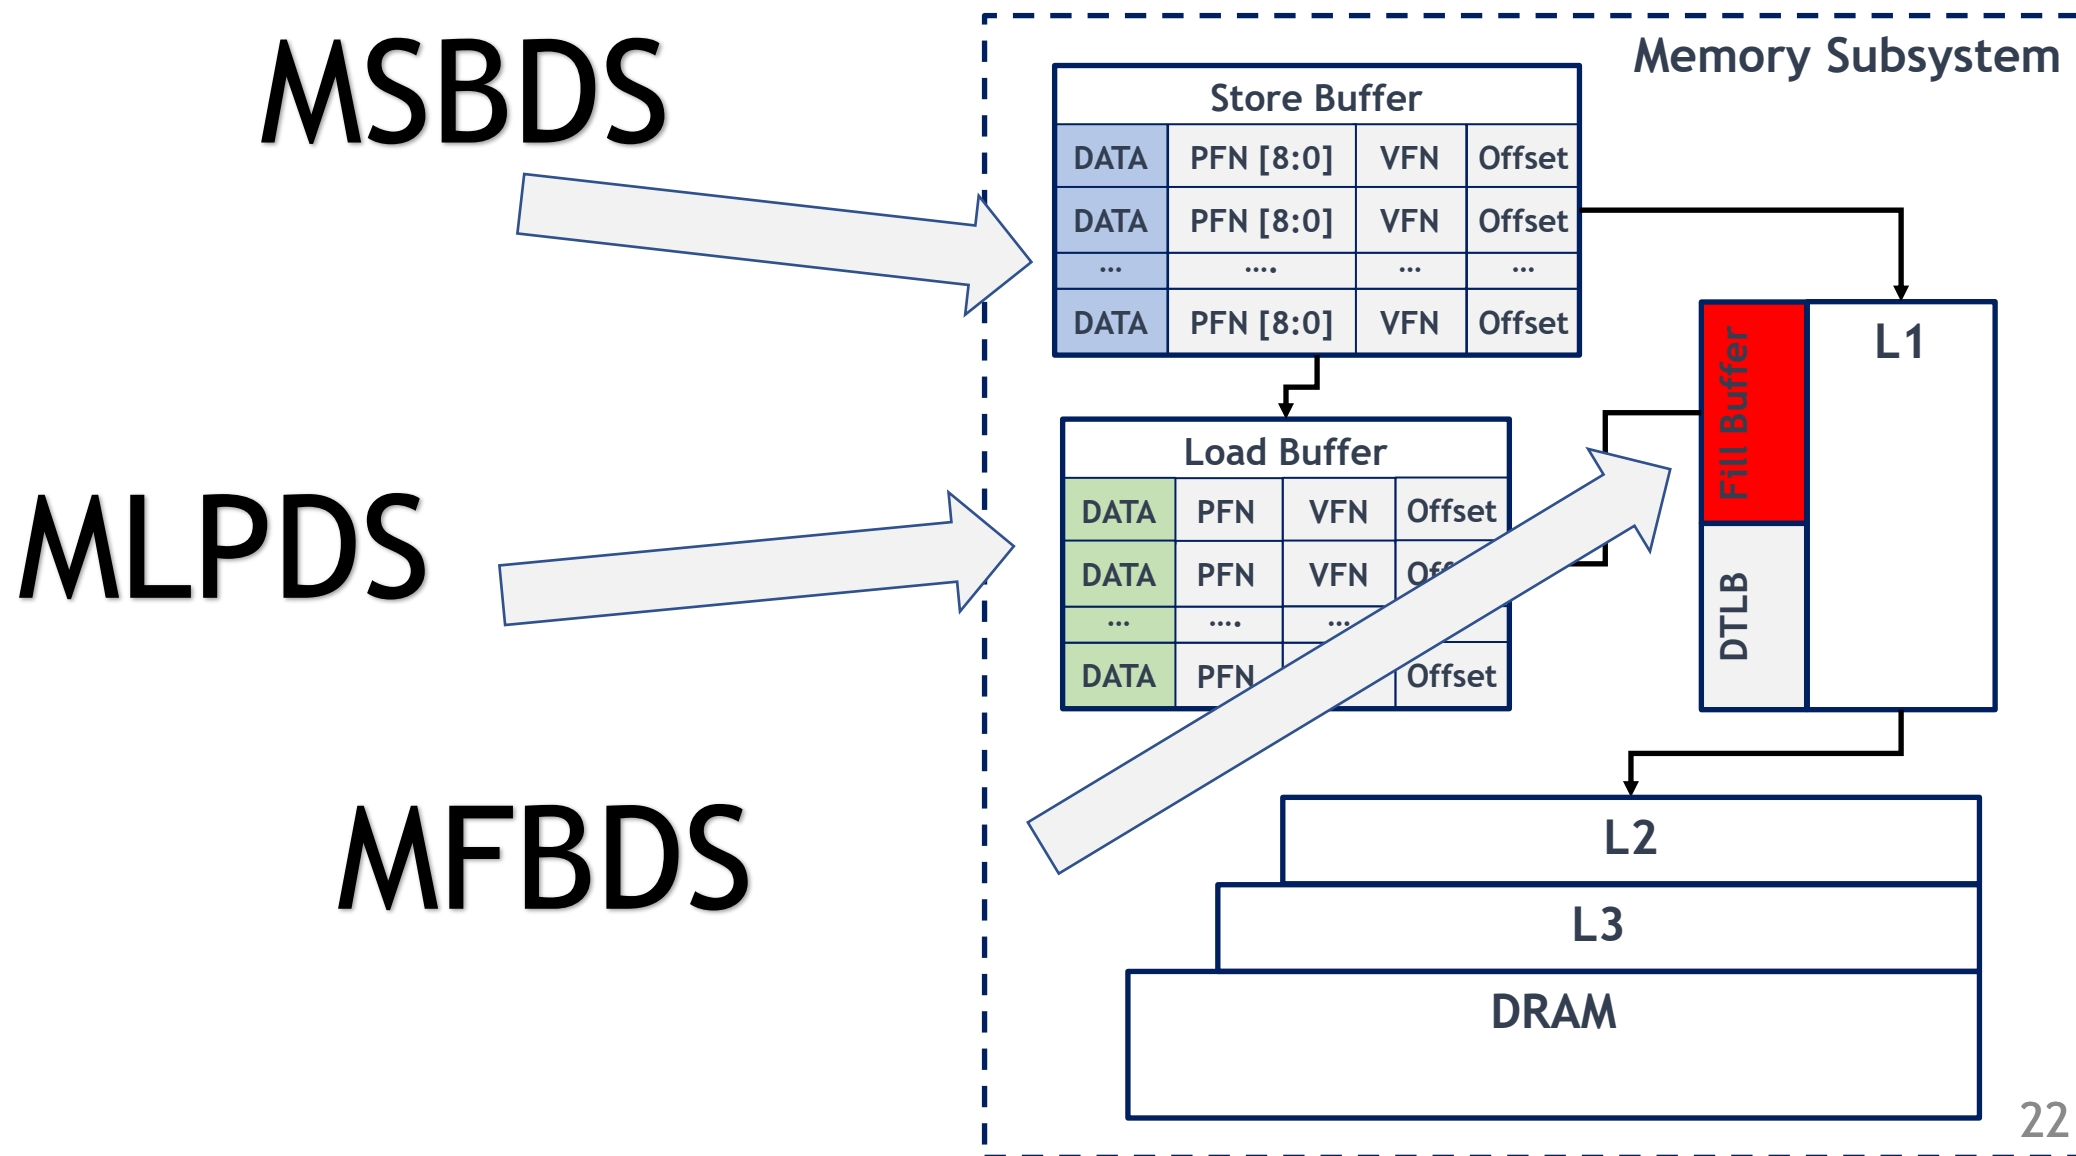


Variant 1: #GP


Variant 3: MC



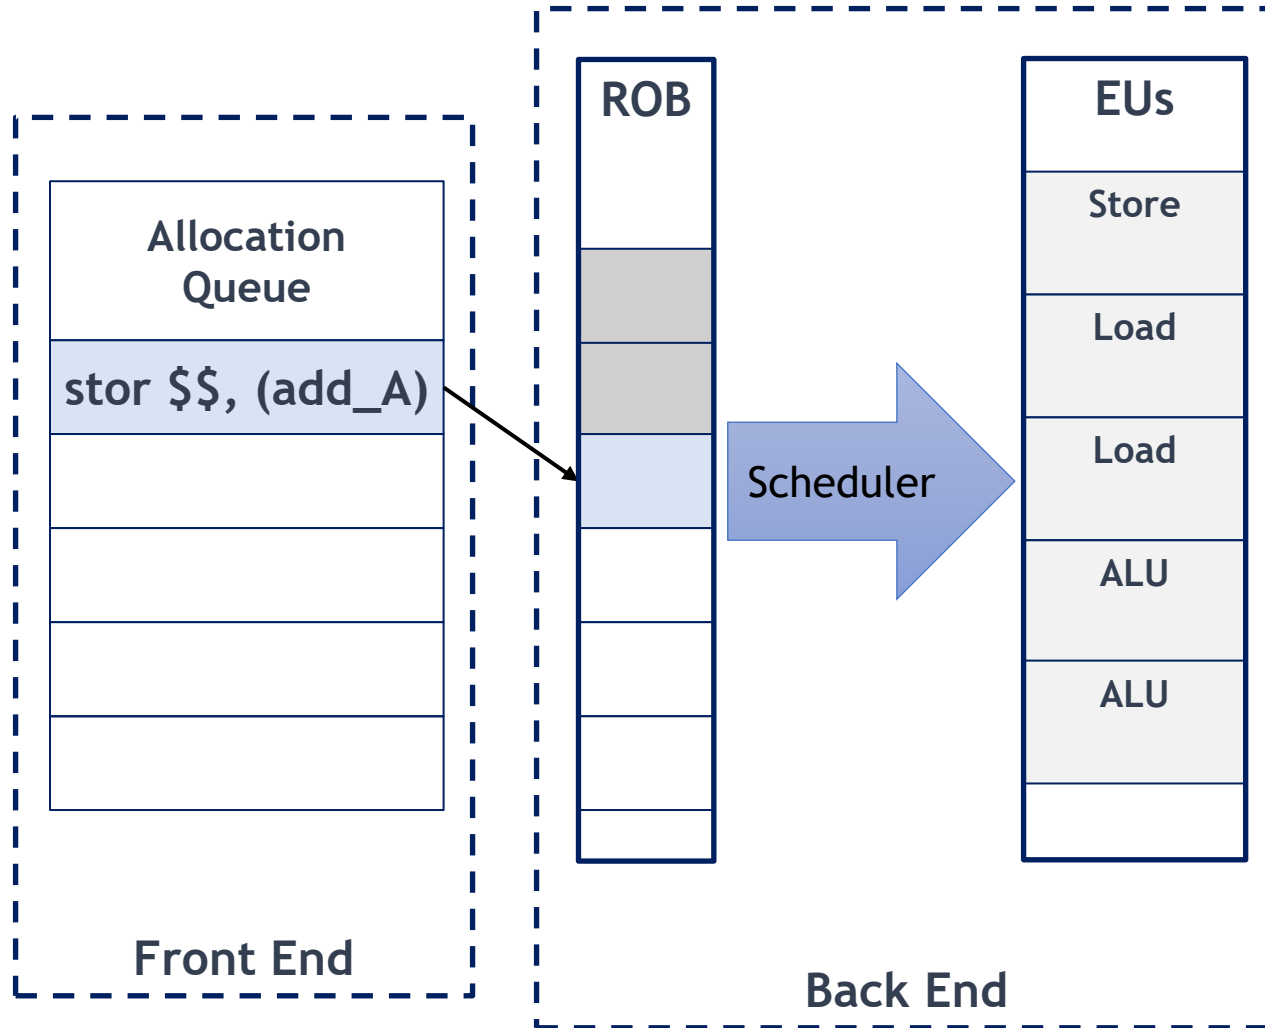
# CPU Memory Subsystem - Leaky Buffers



# CPU Memory Subsystem

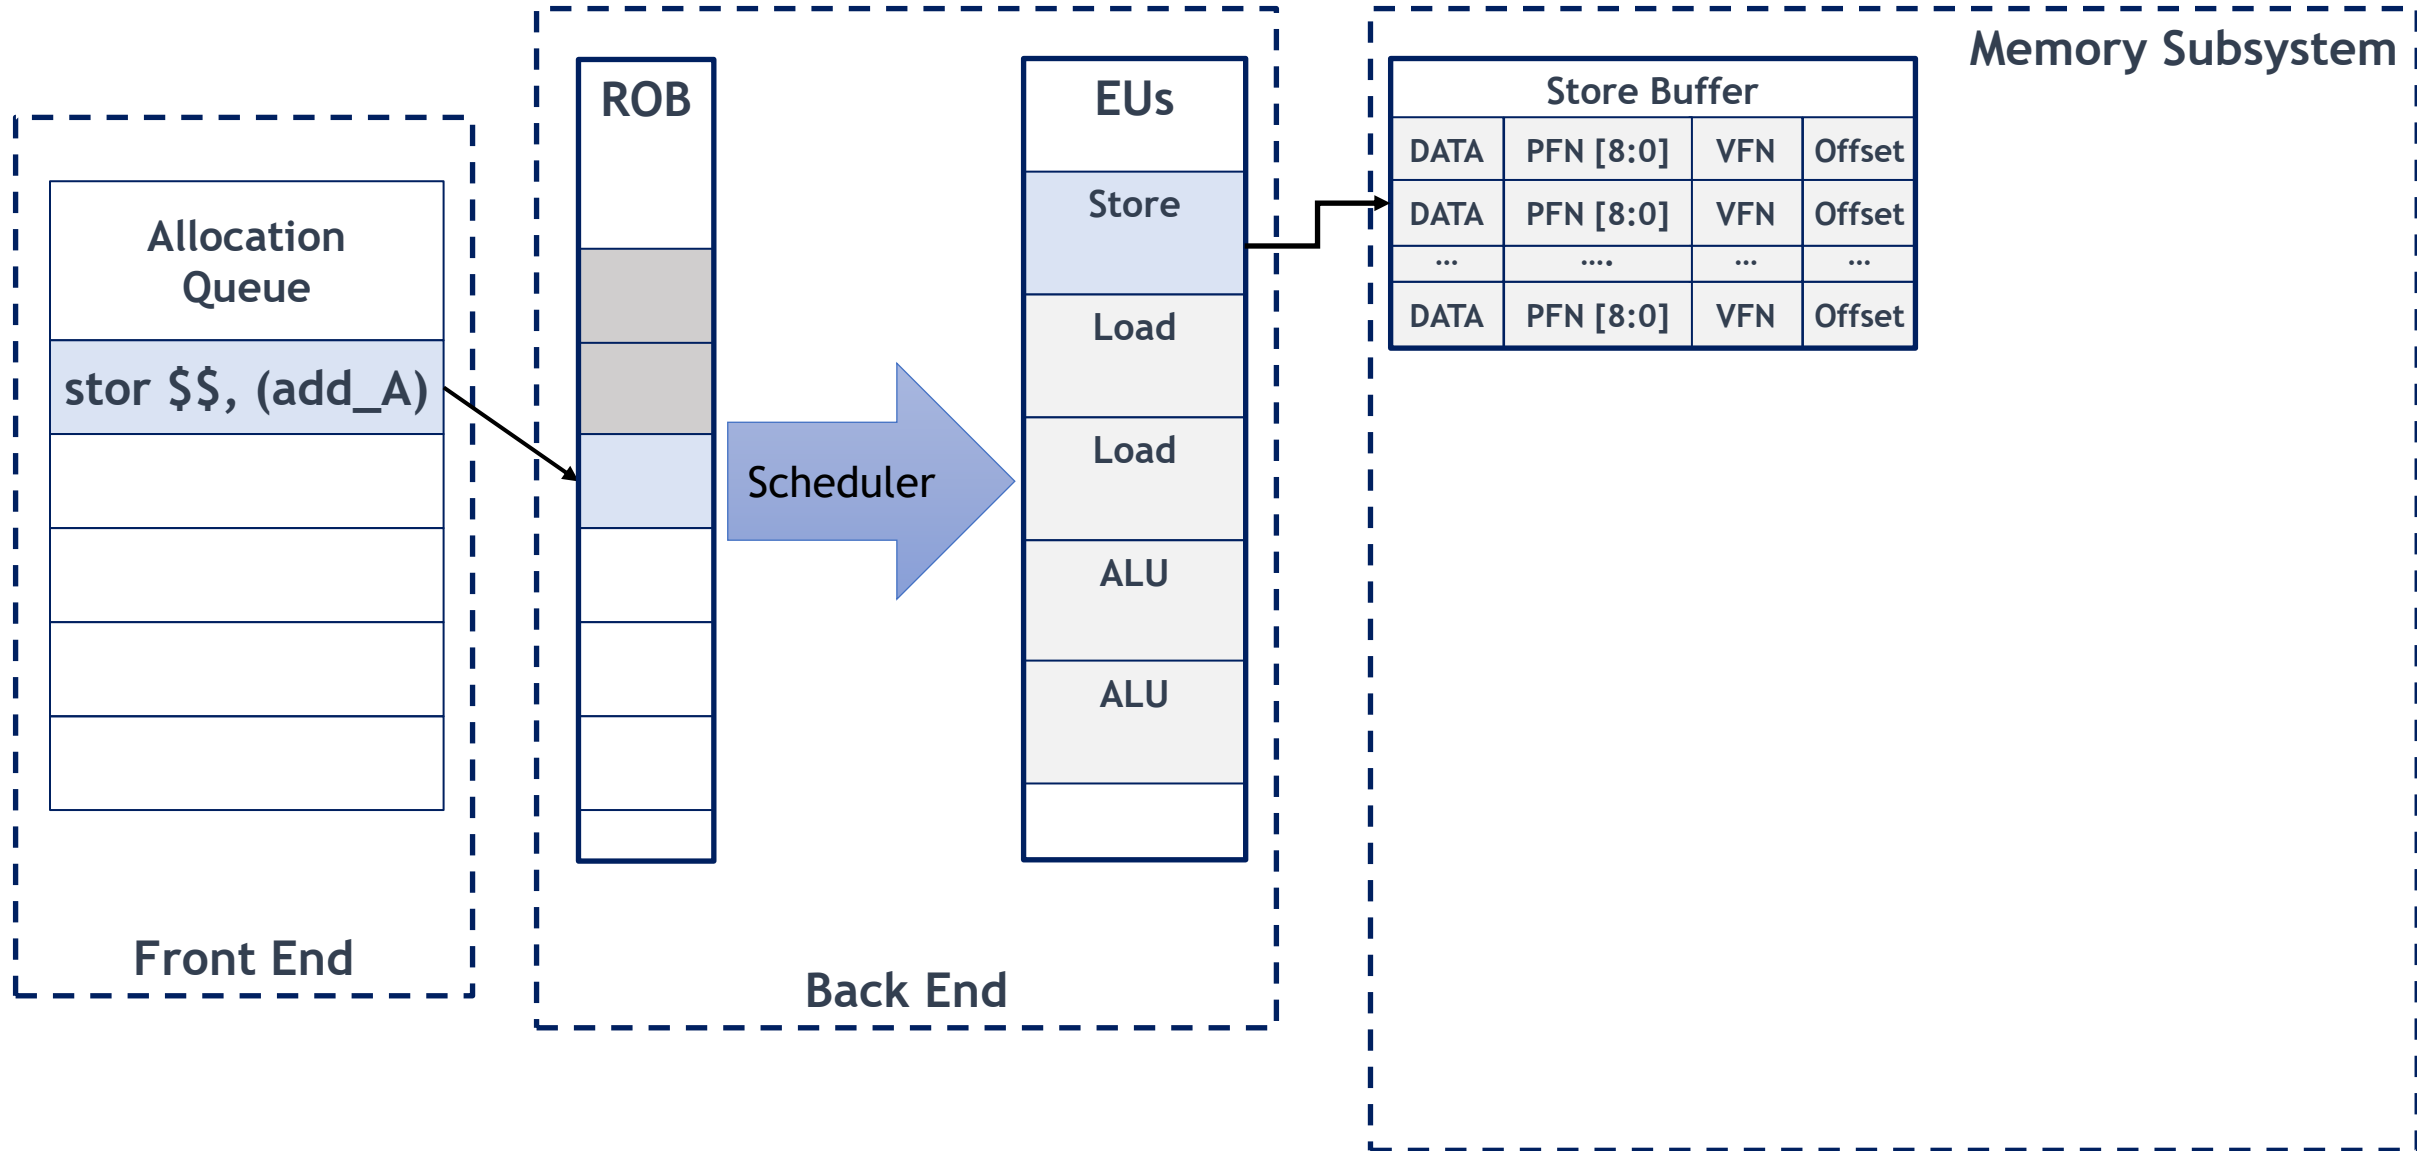


# CPU Memory Subsystem

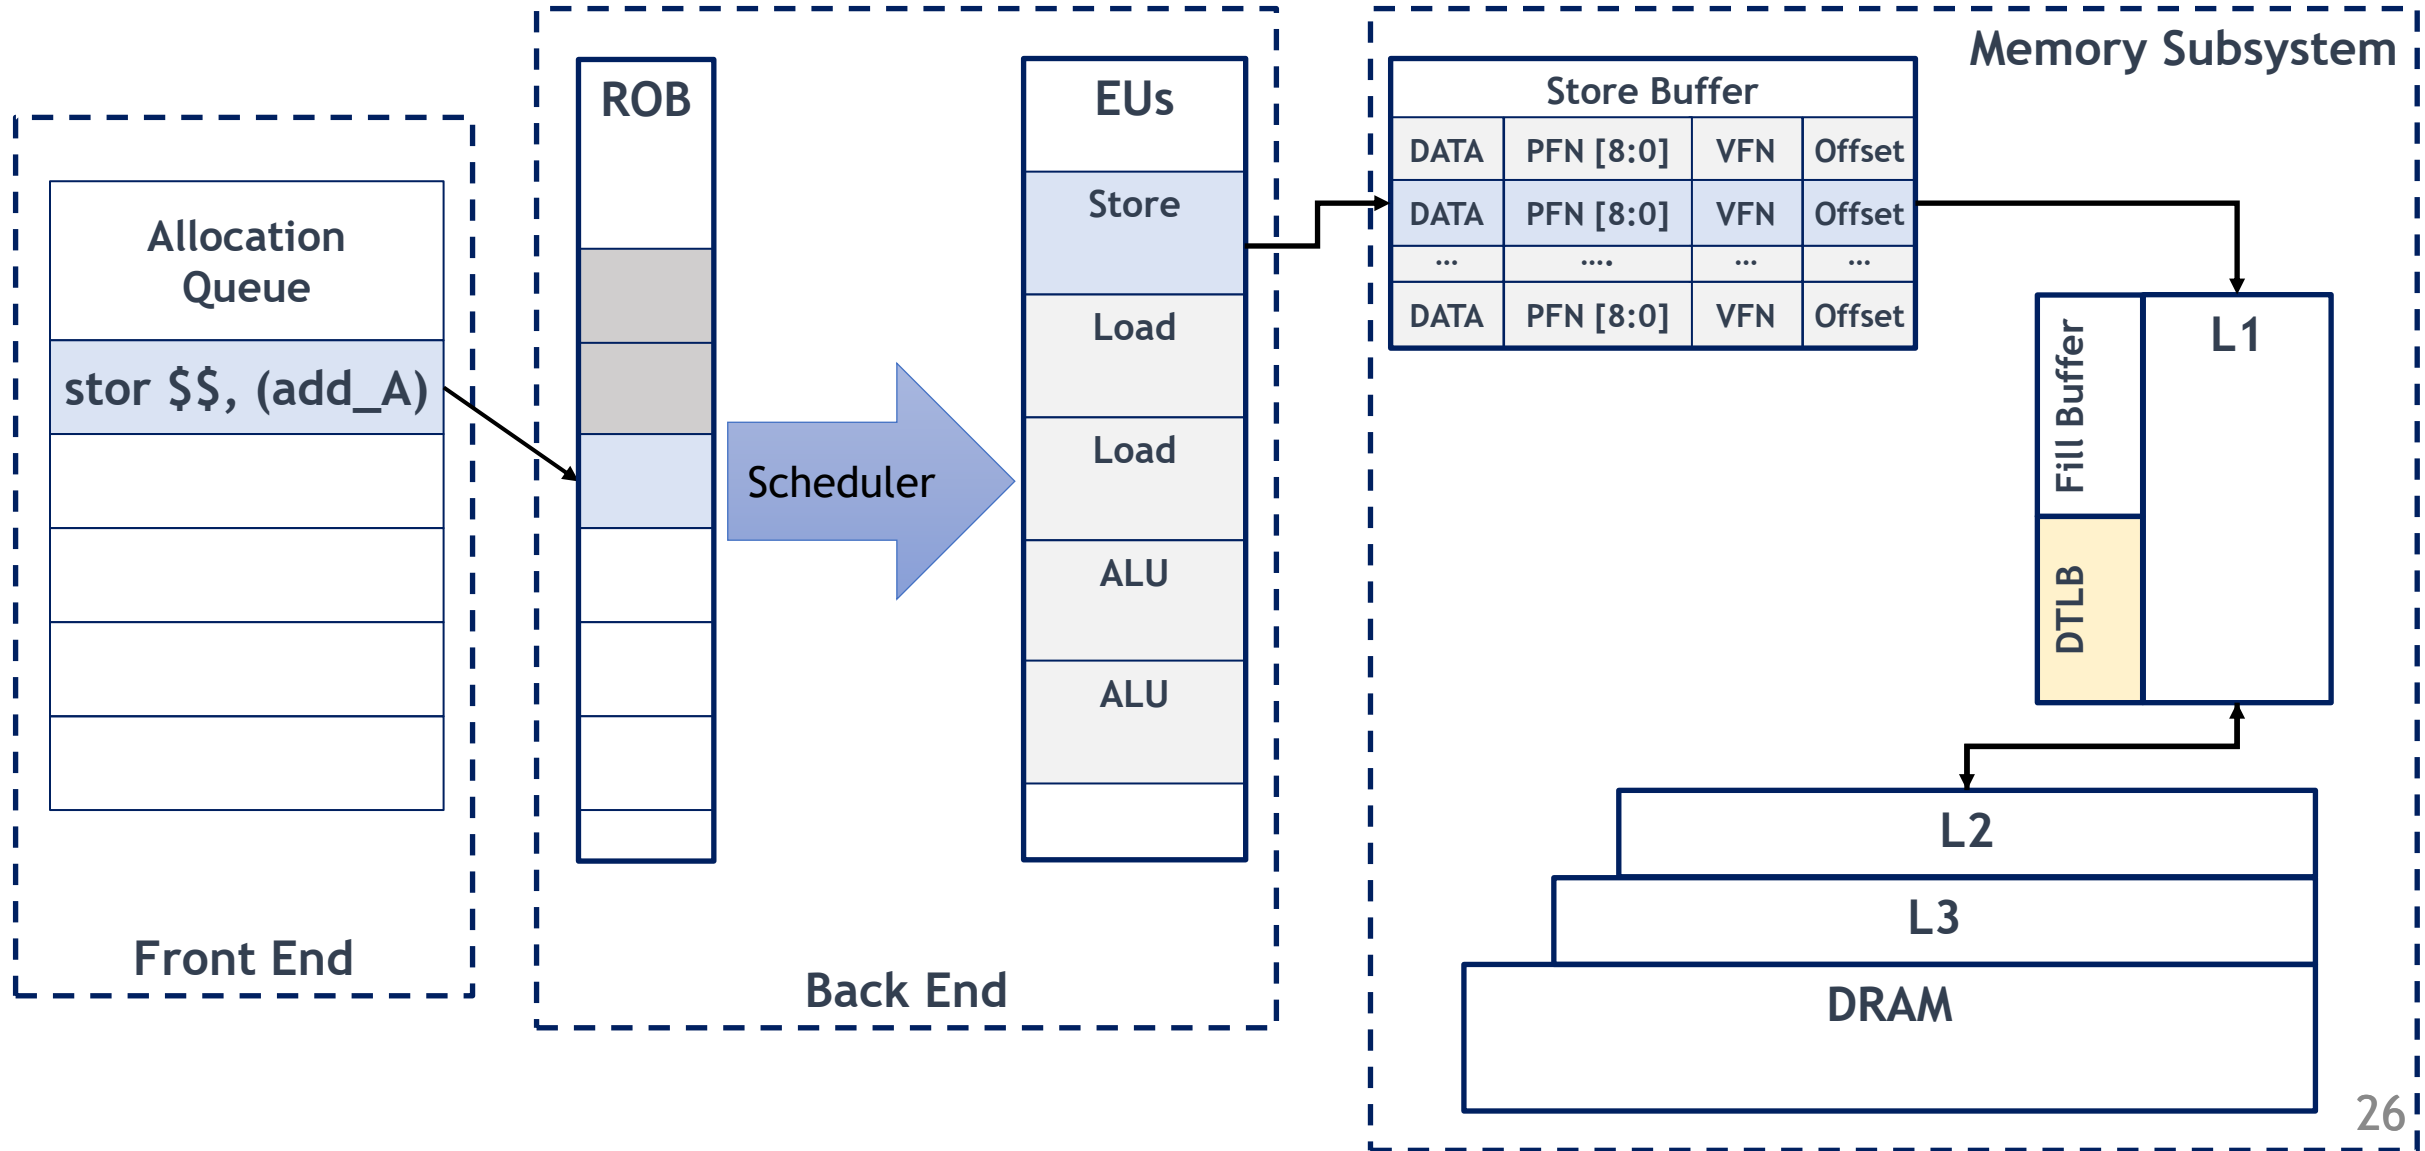




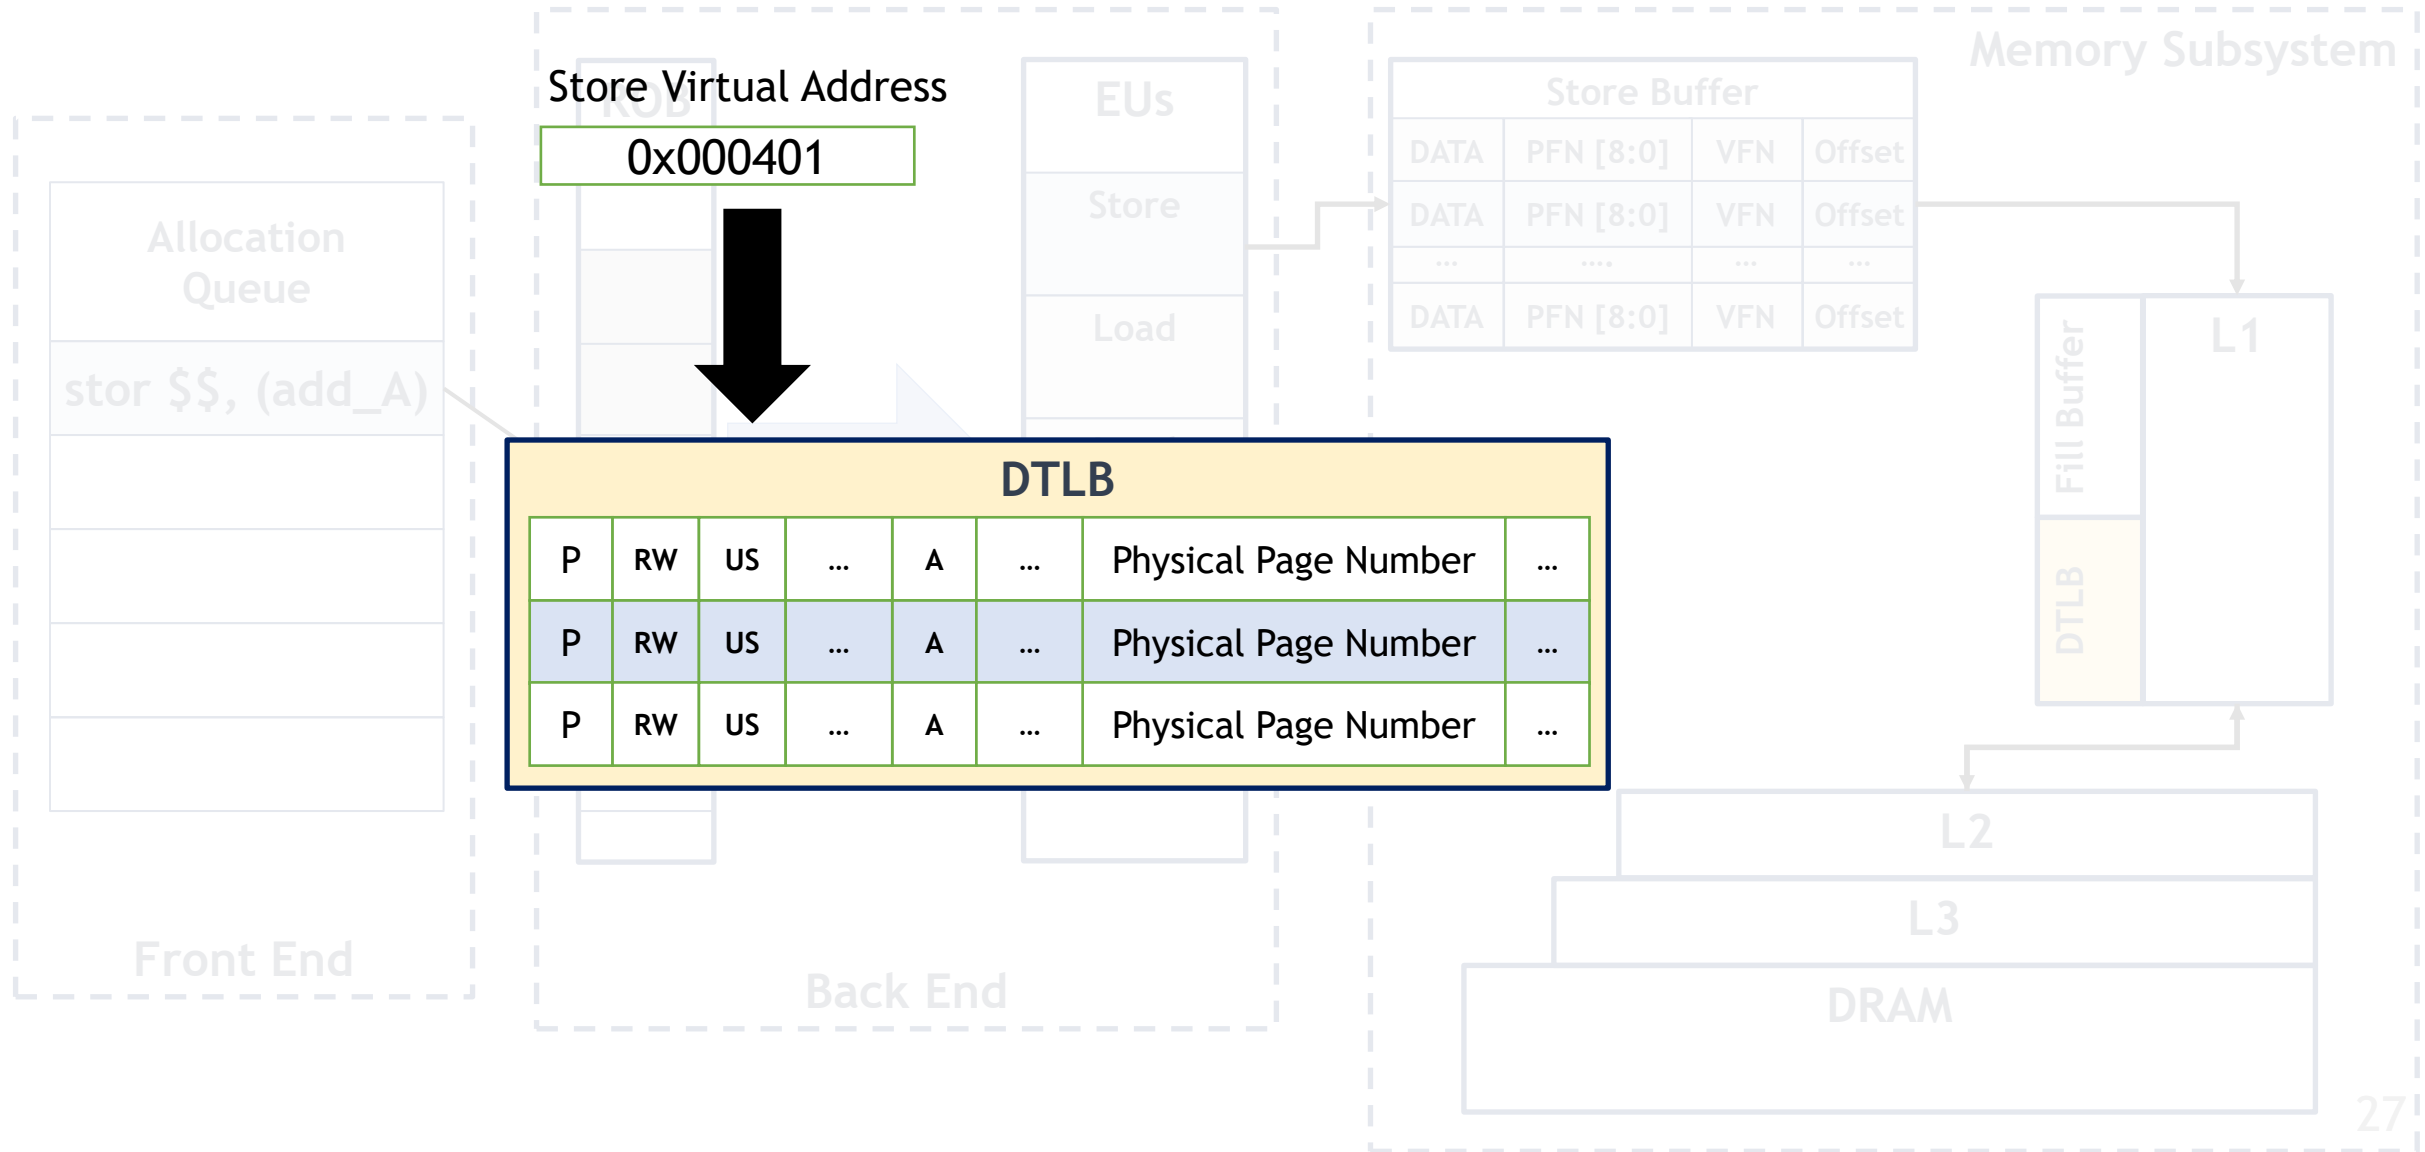
# CPU Memory Subsystem



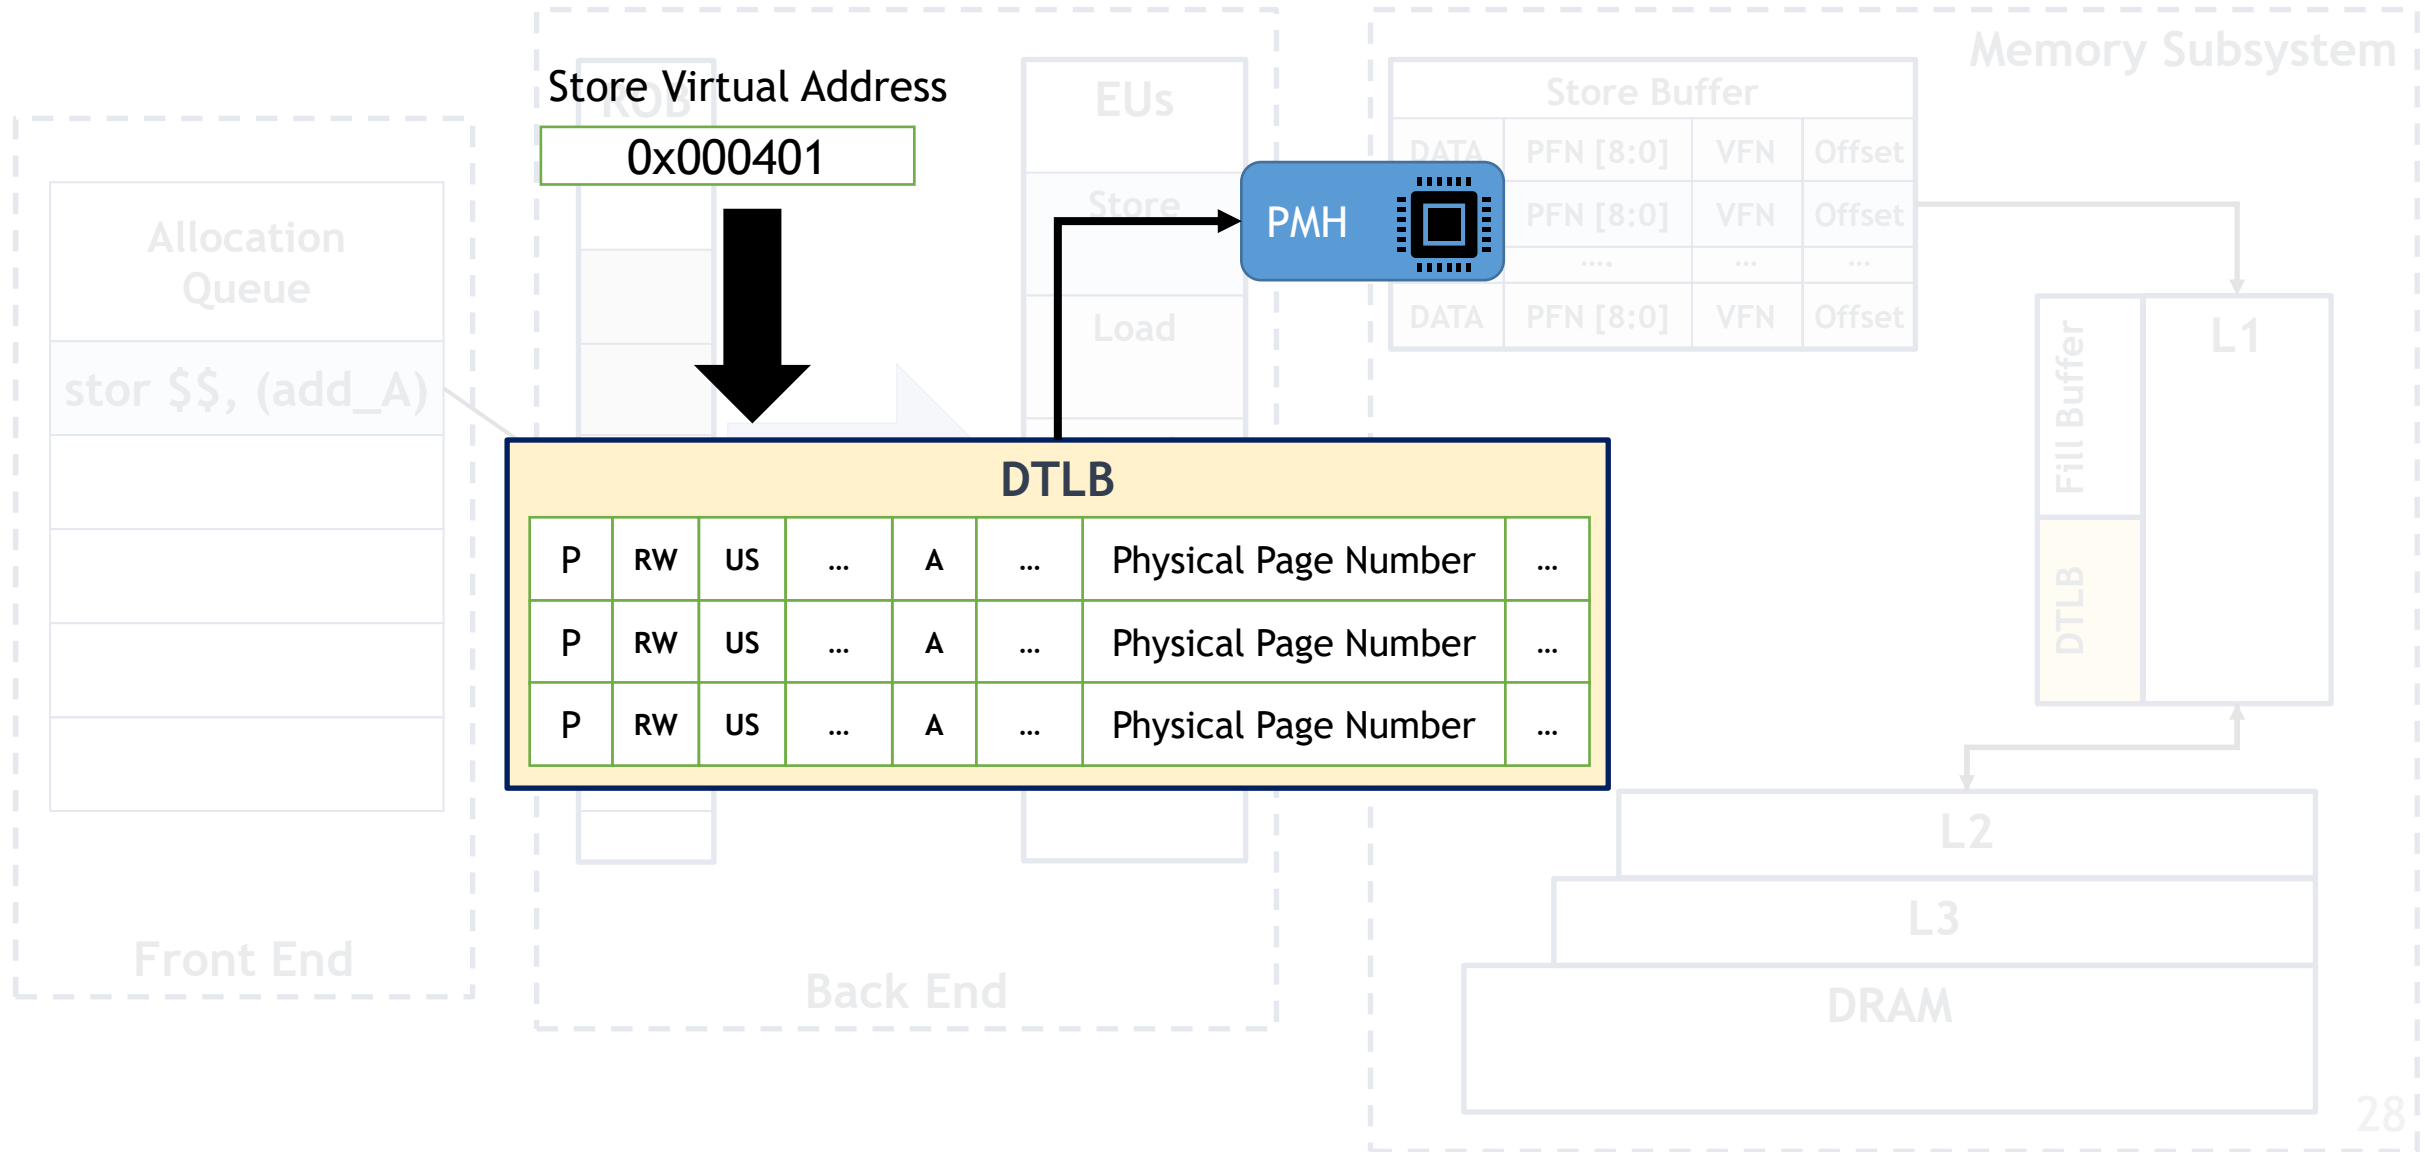
# CPU Memory Subsystem



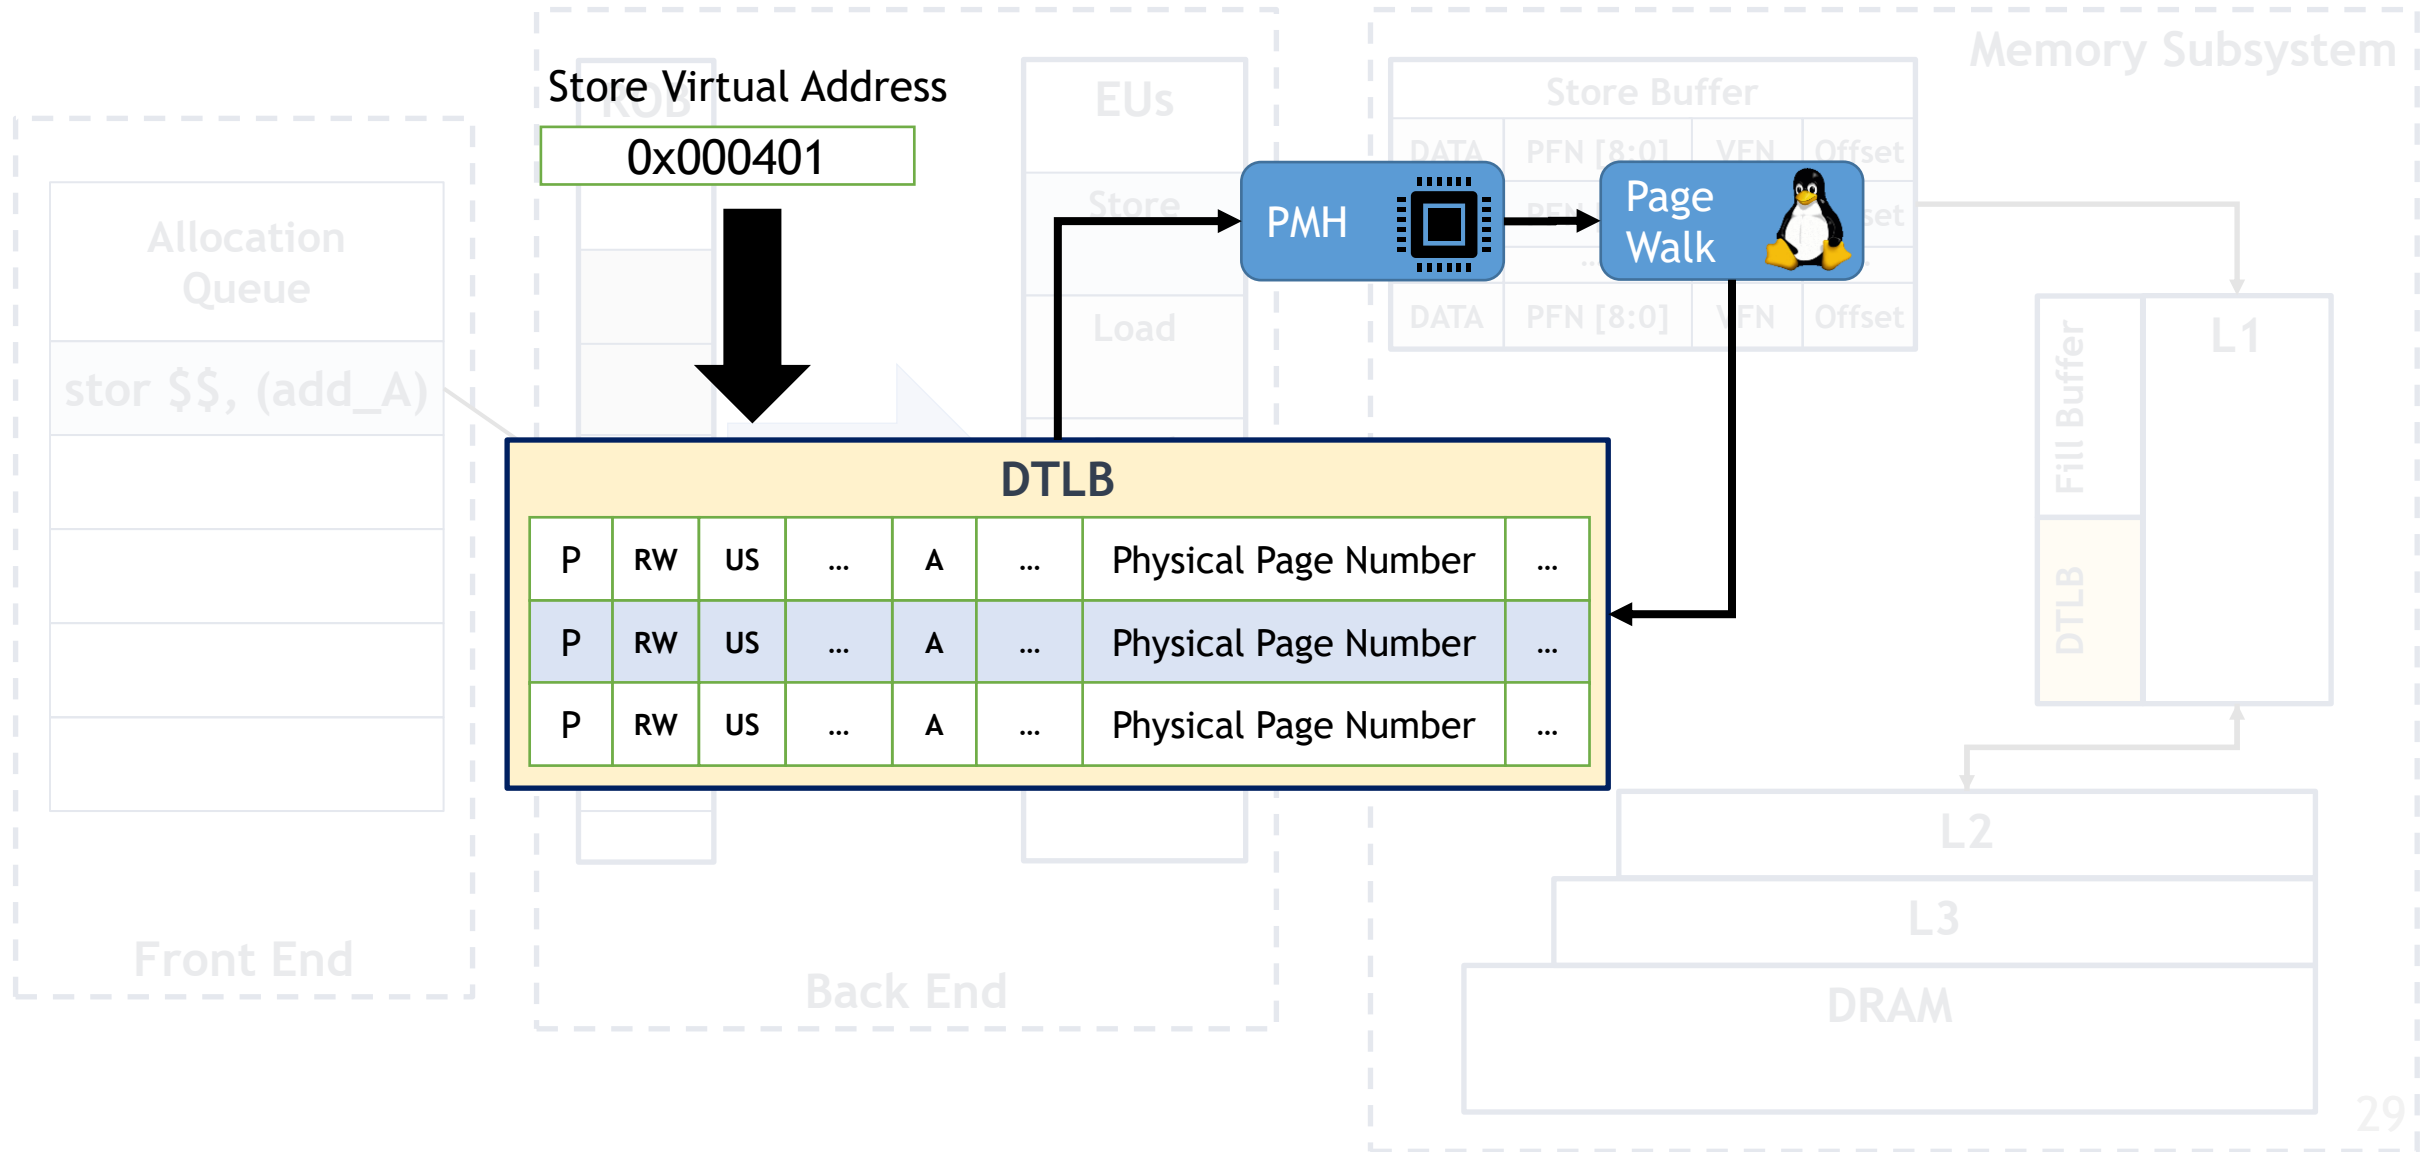
# CPU Memory Subsystem



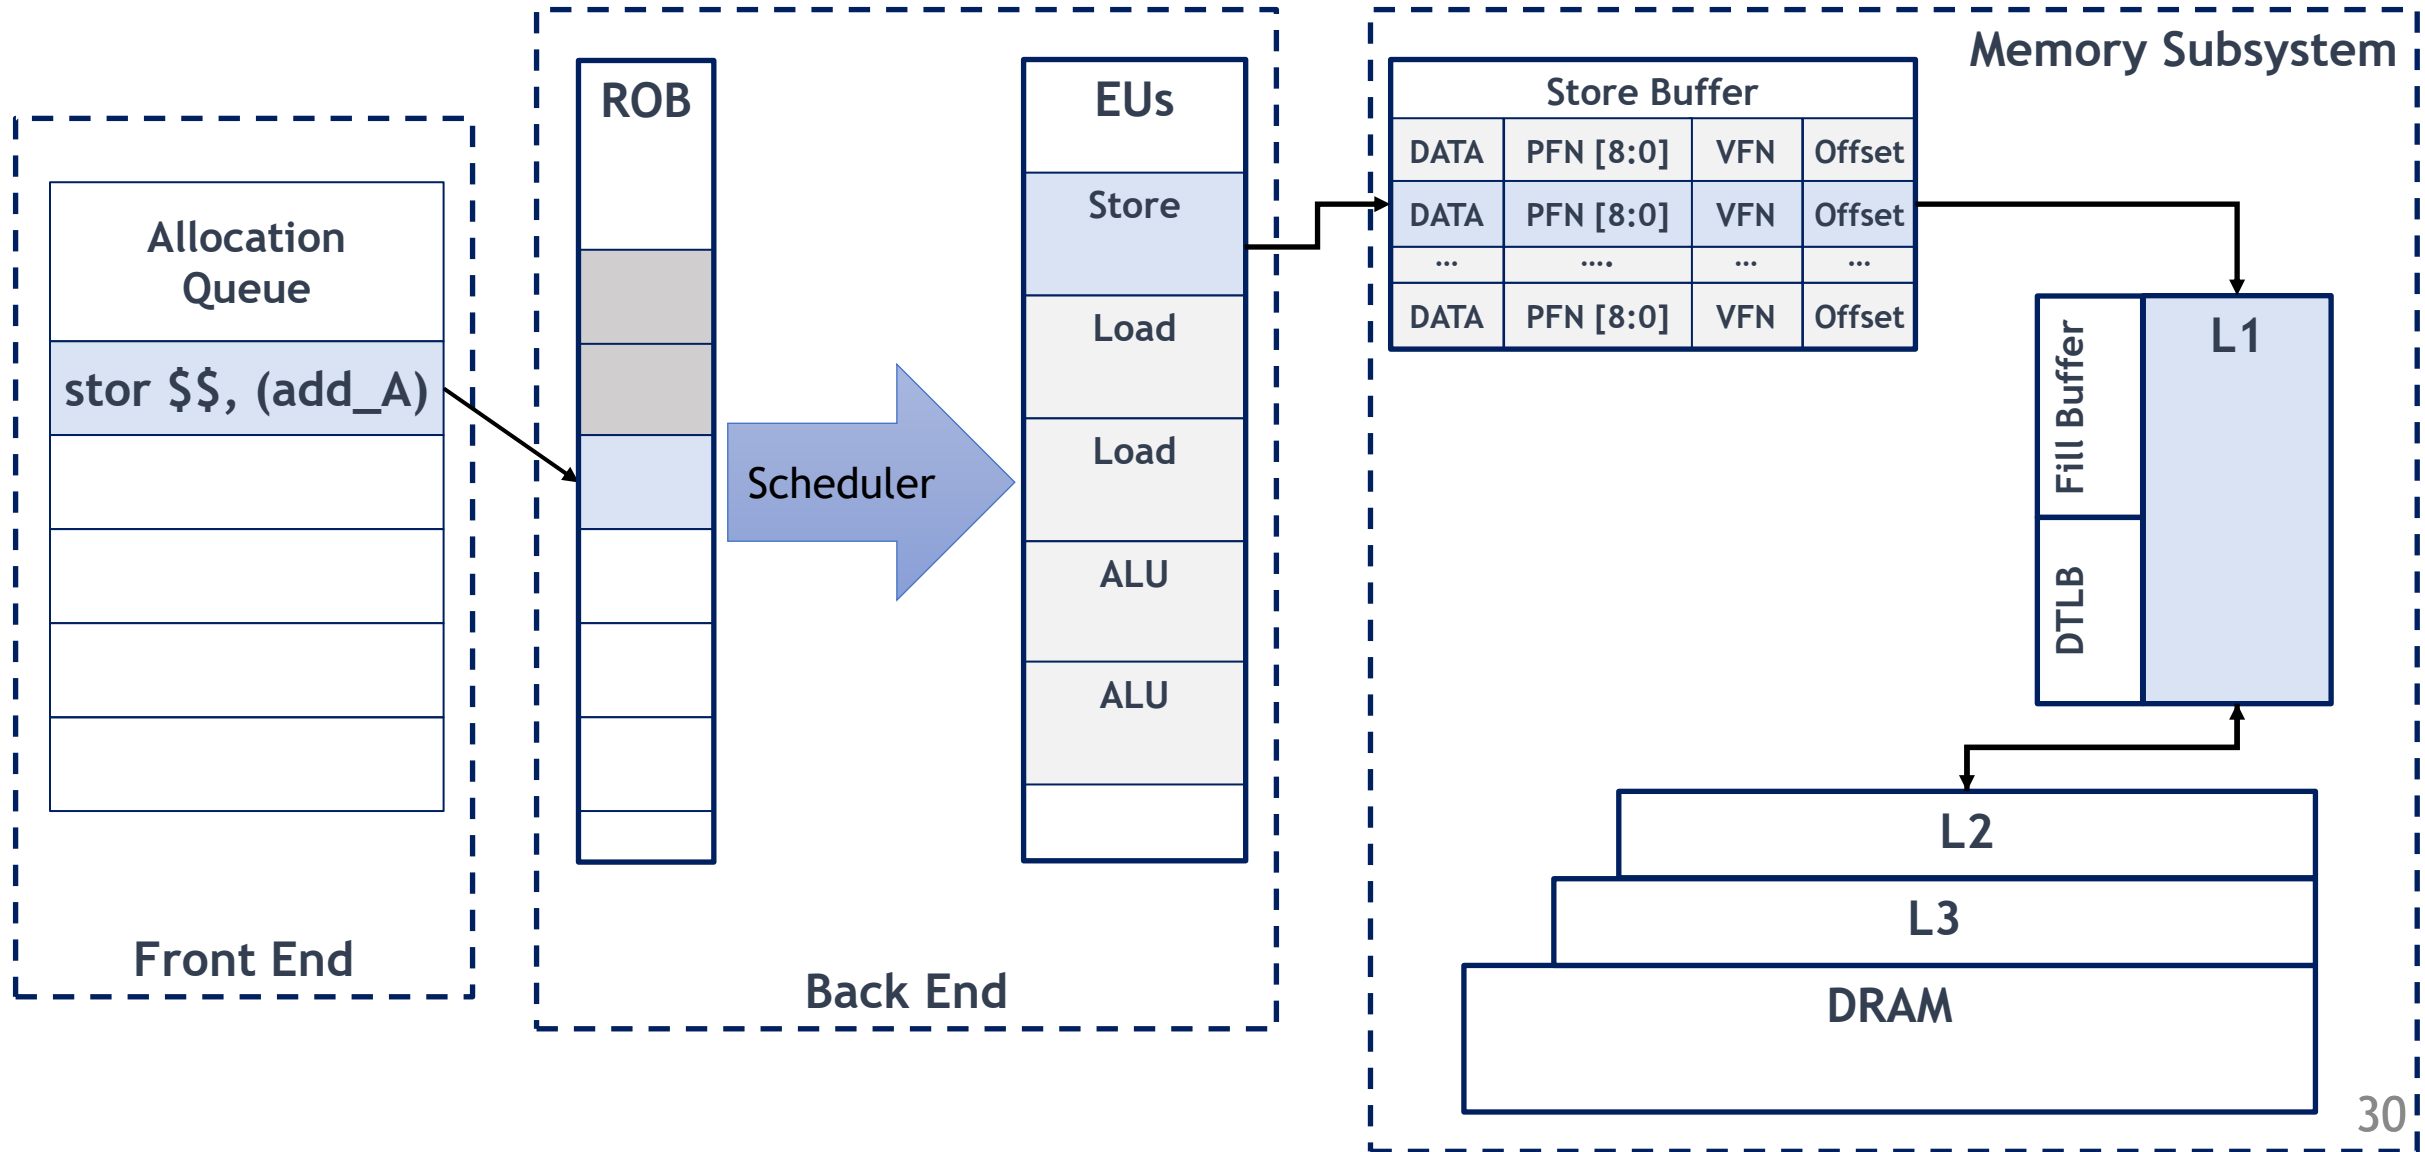
# CPU Memory Subsystem



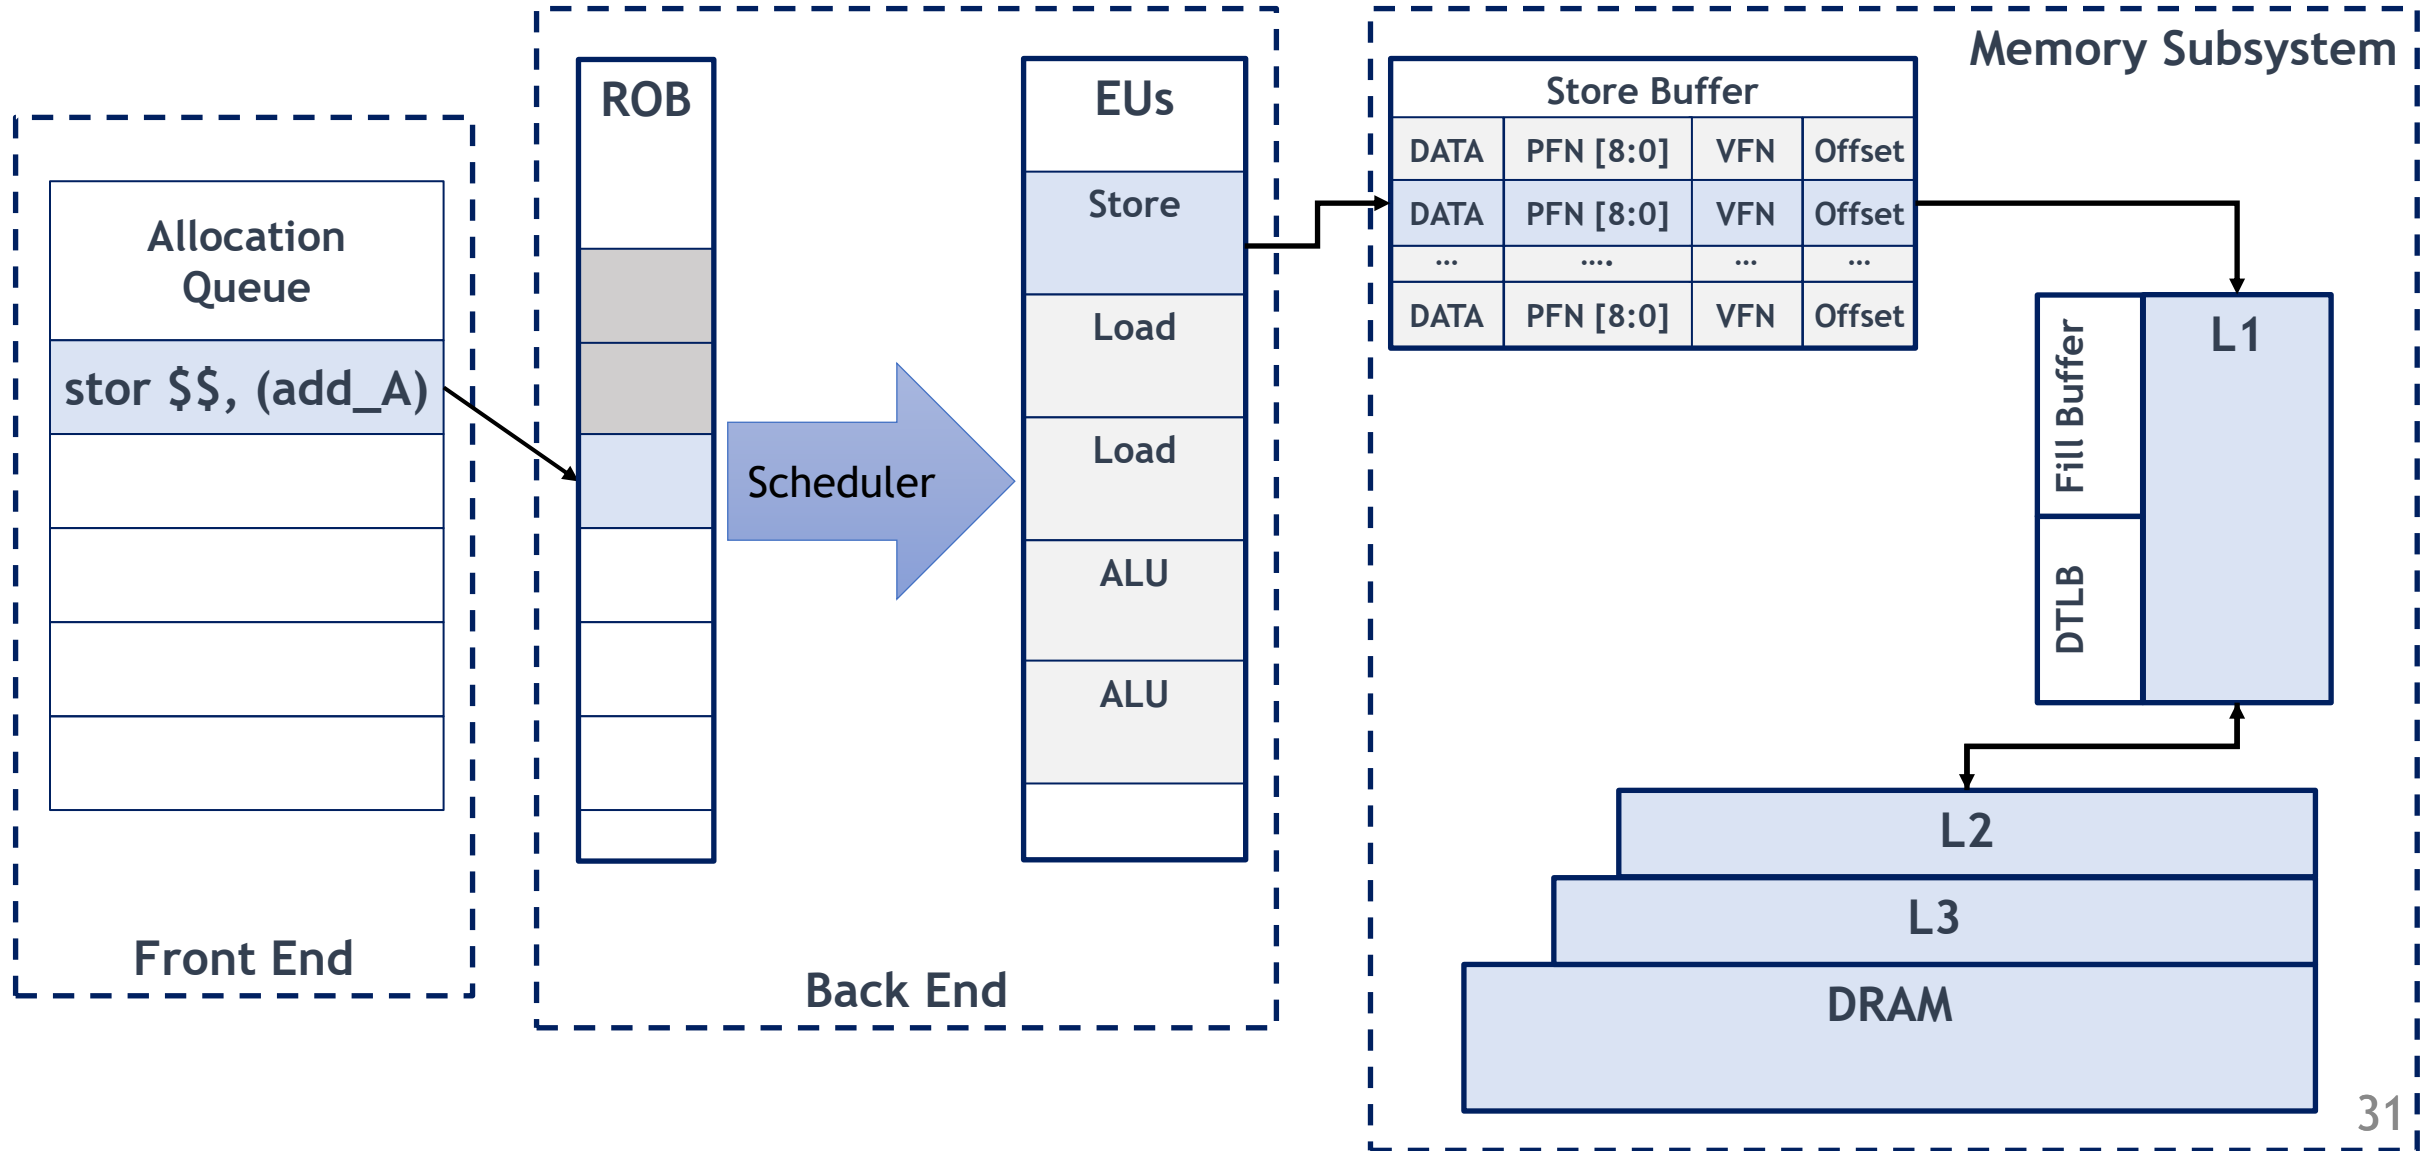
# CPU Memory Subsystem



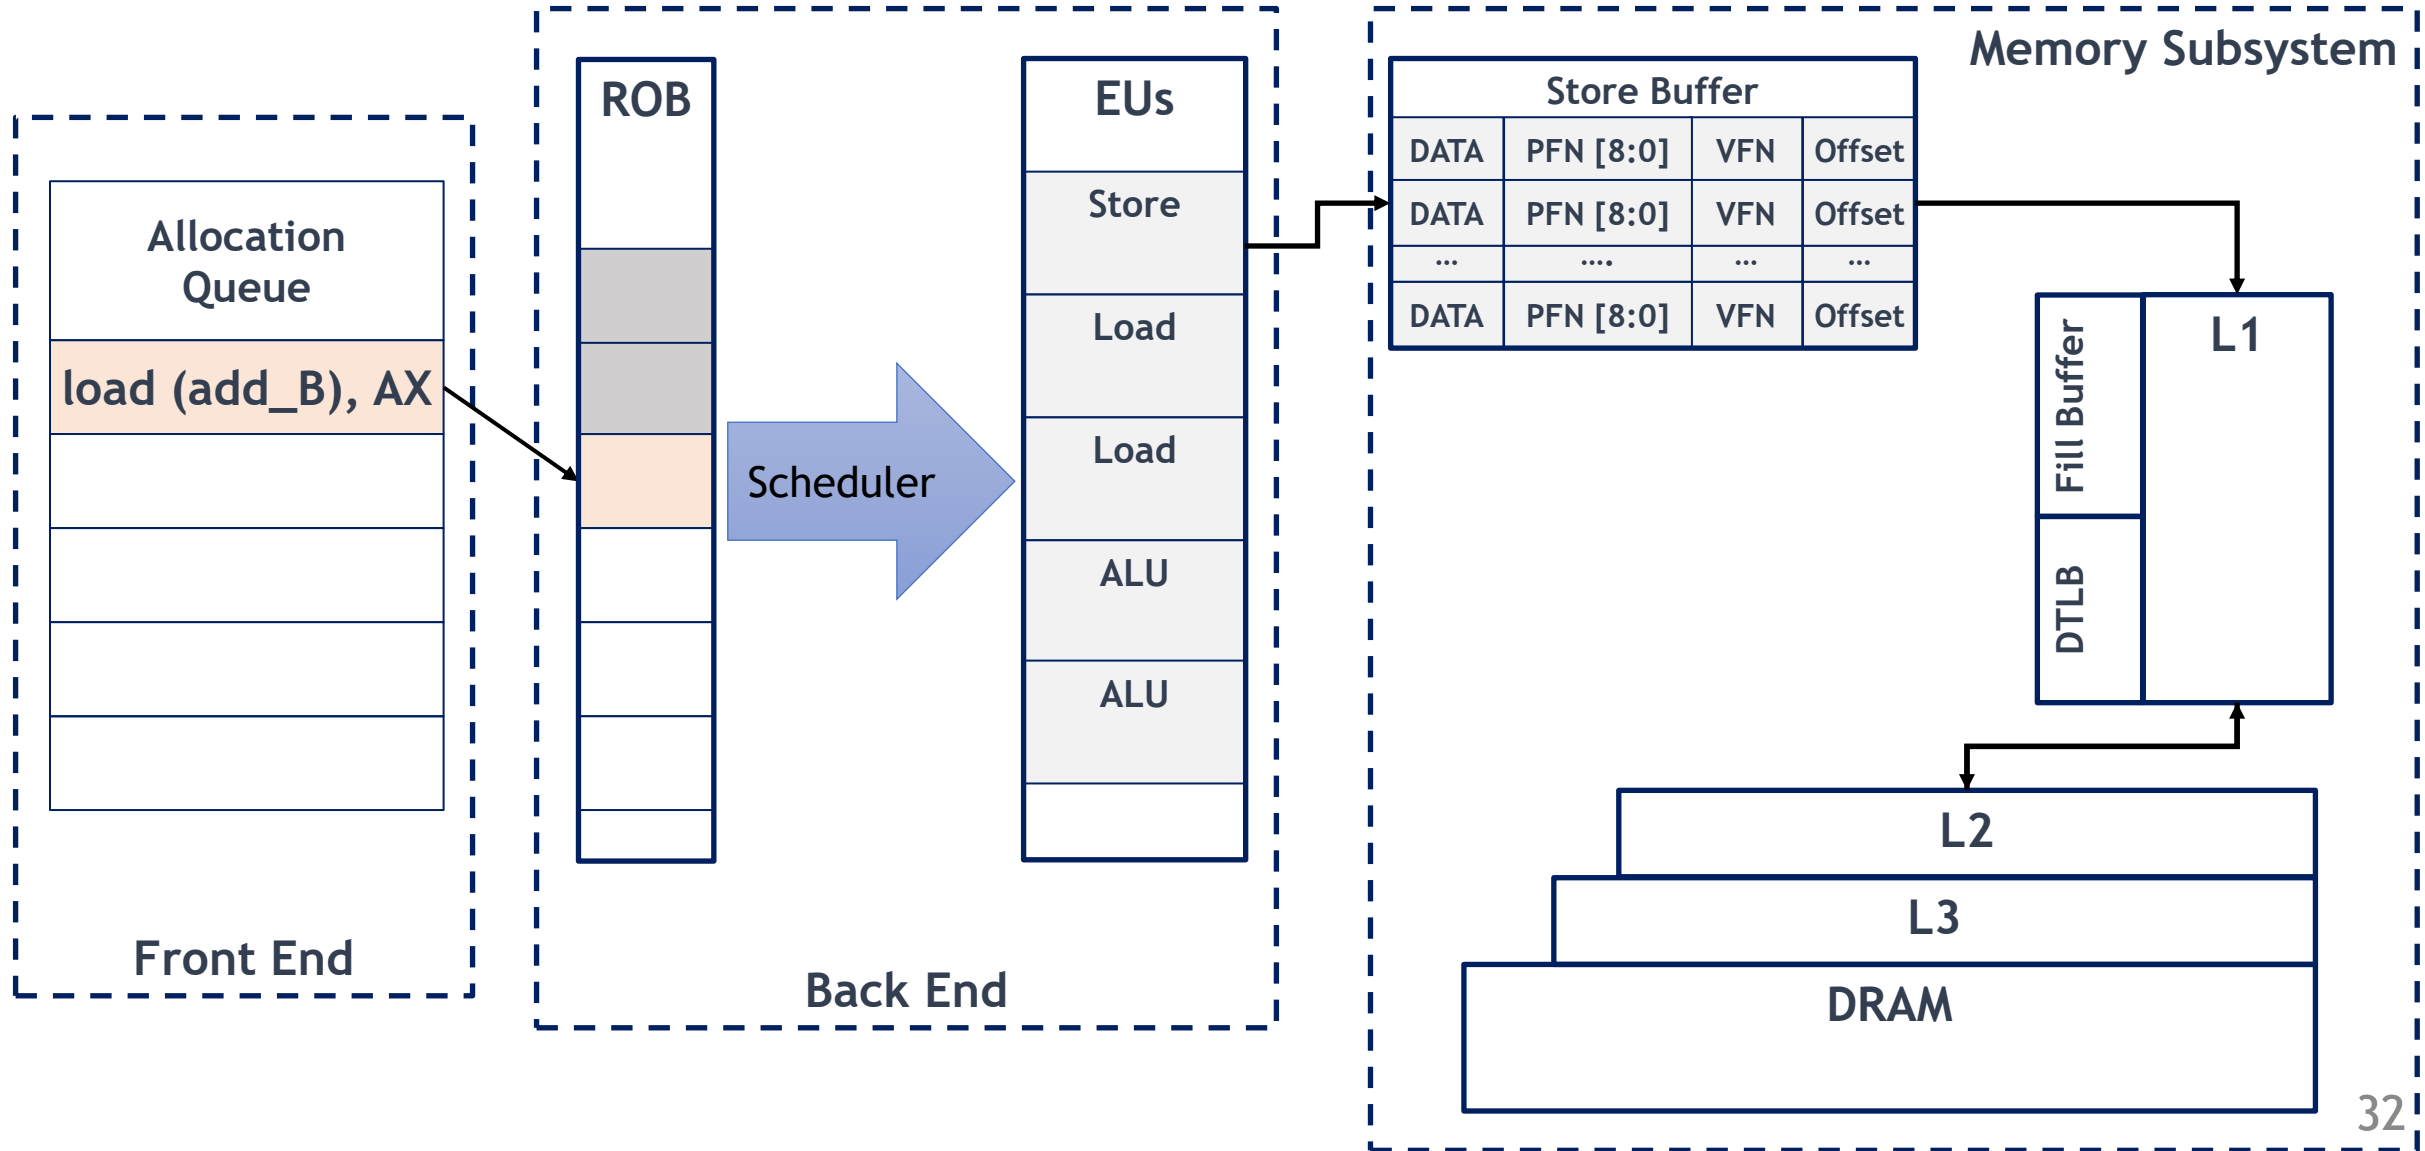
# CPU Memory Subsystem



# CPU Memory Subsystem

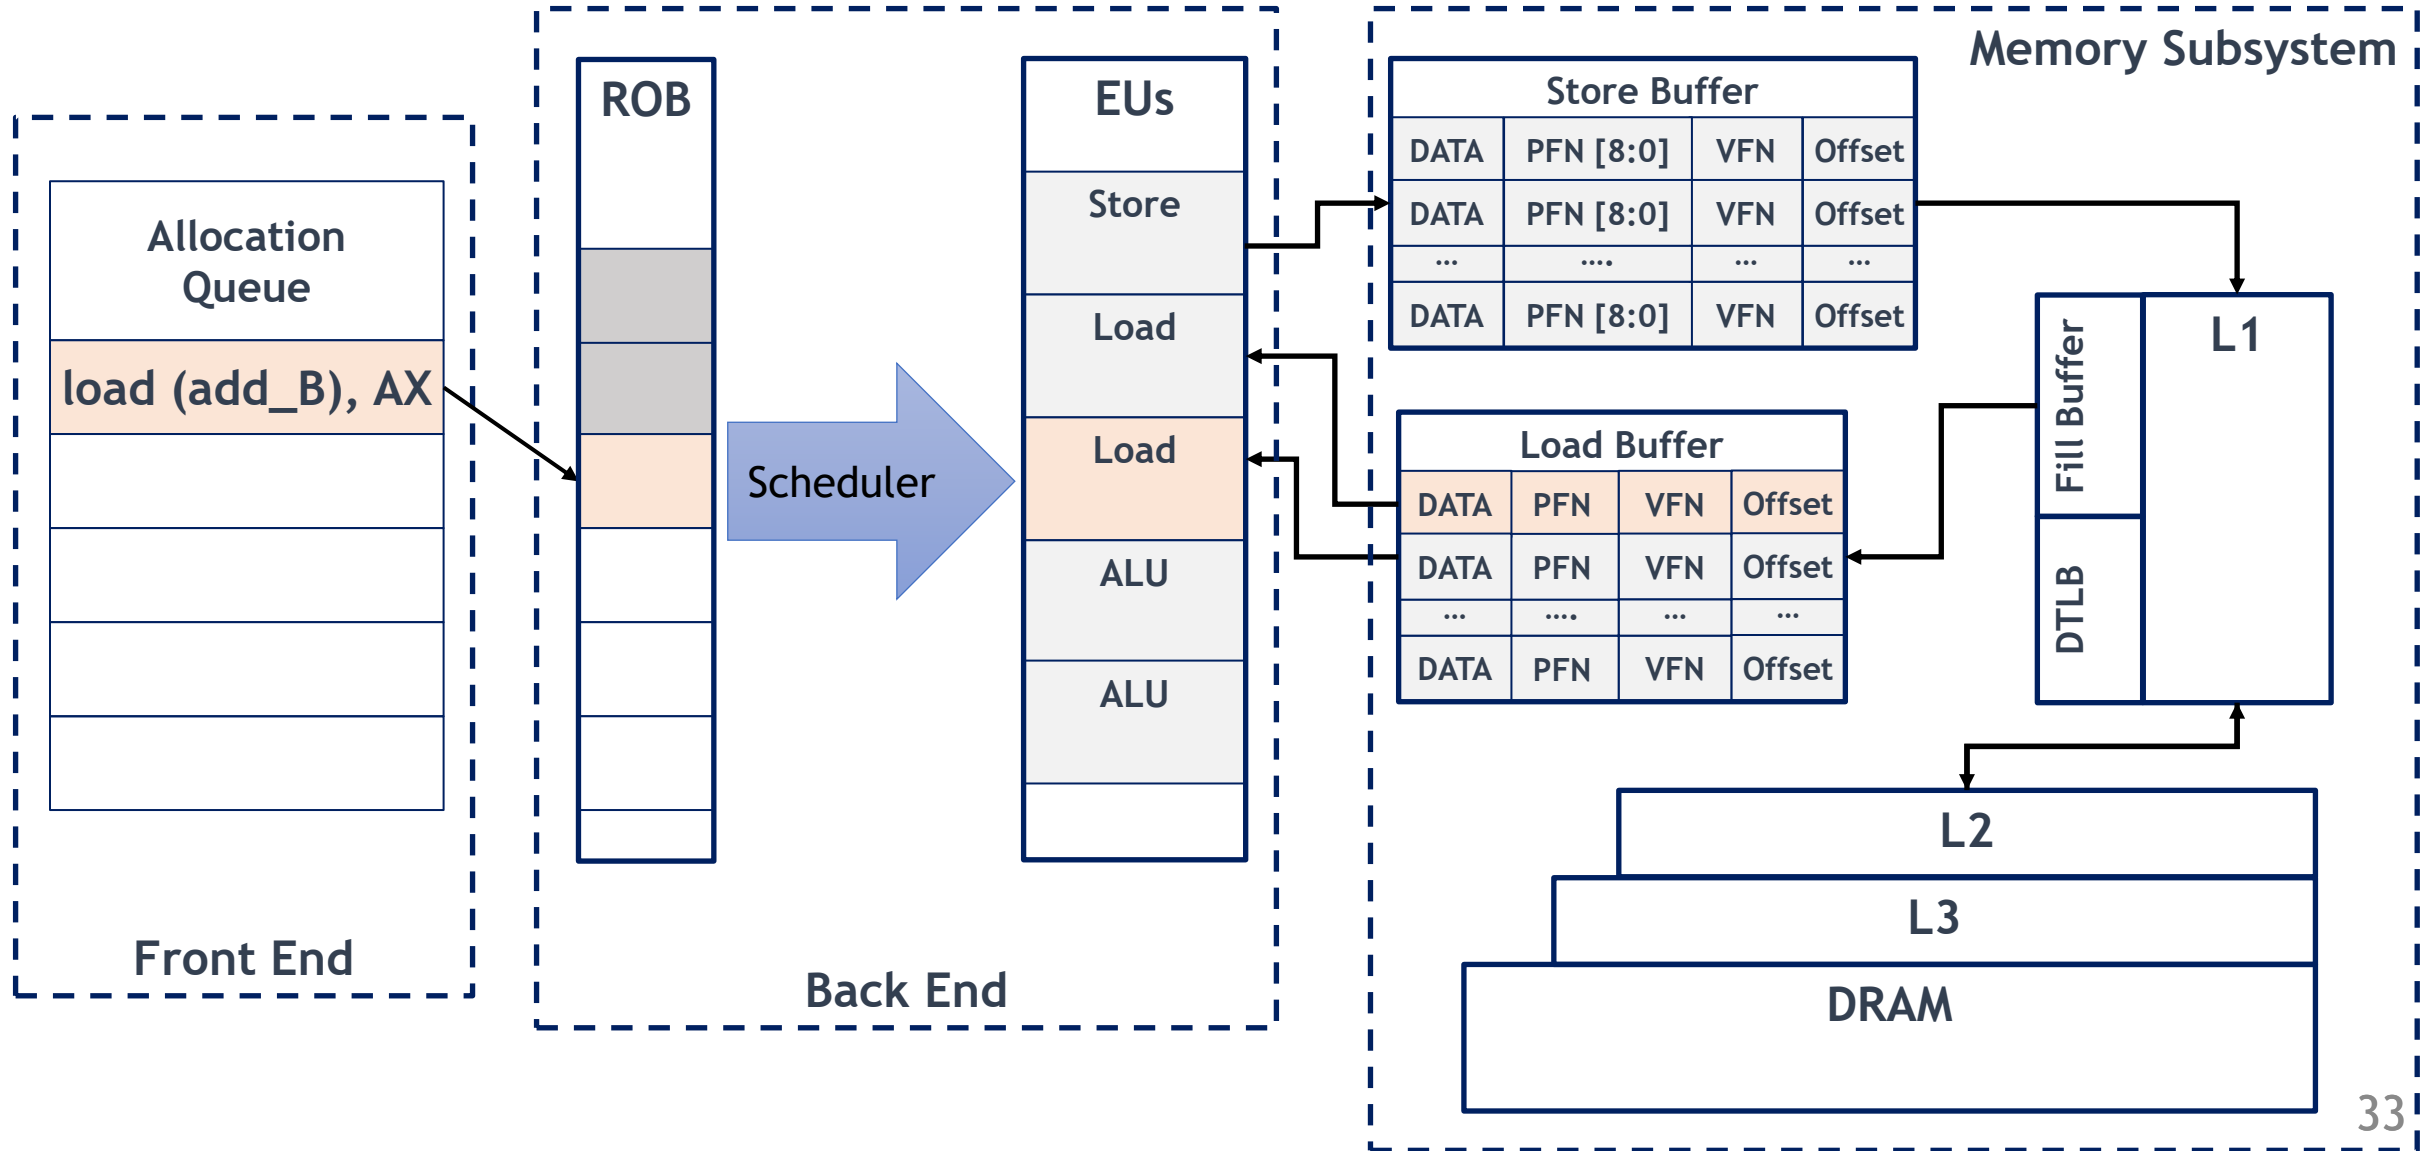


# CPU Memory Subsystem

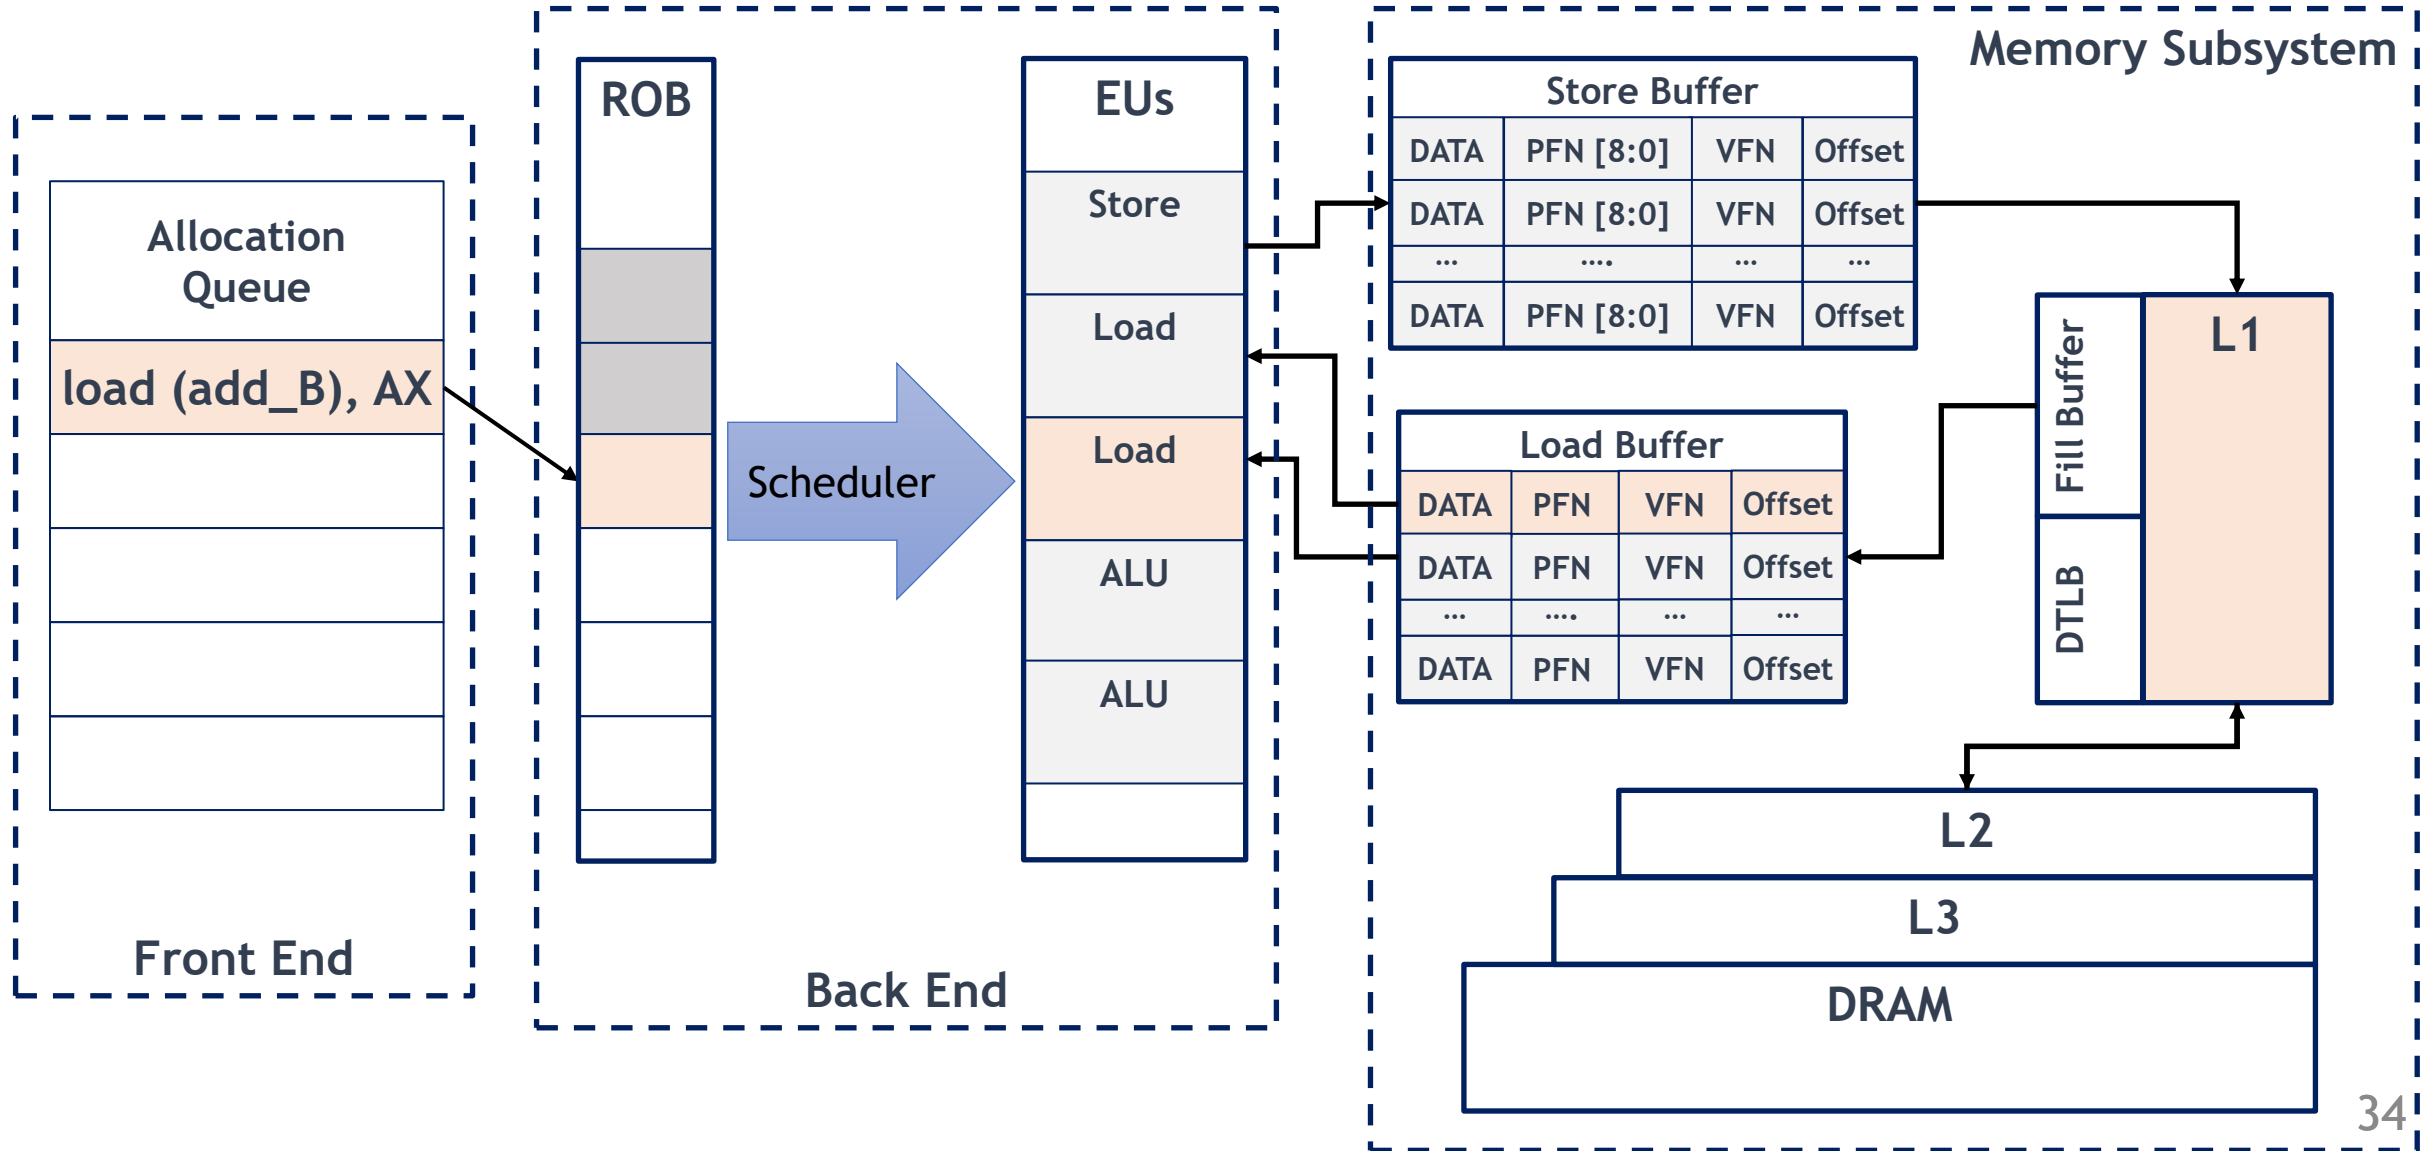




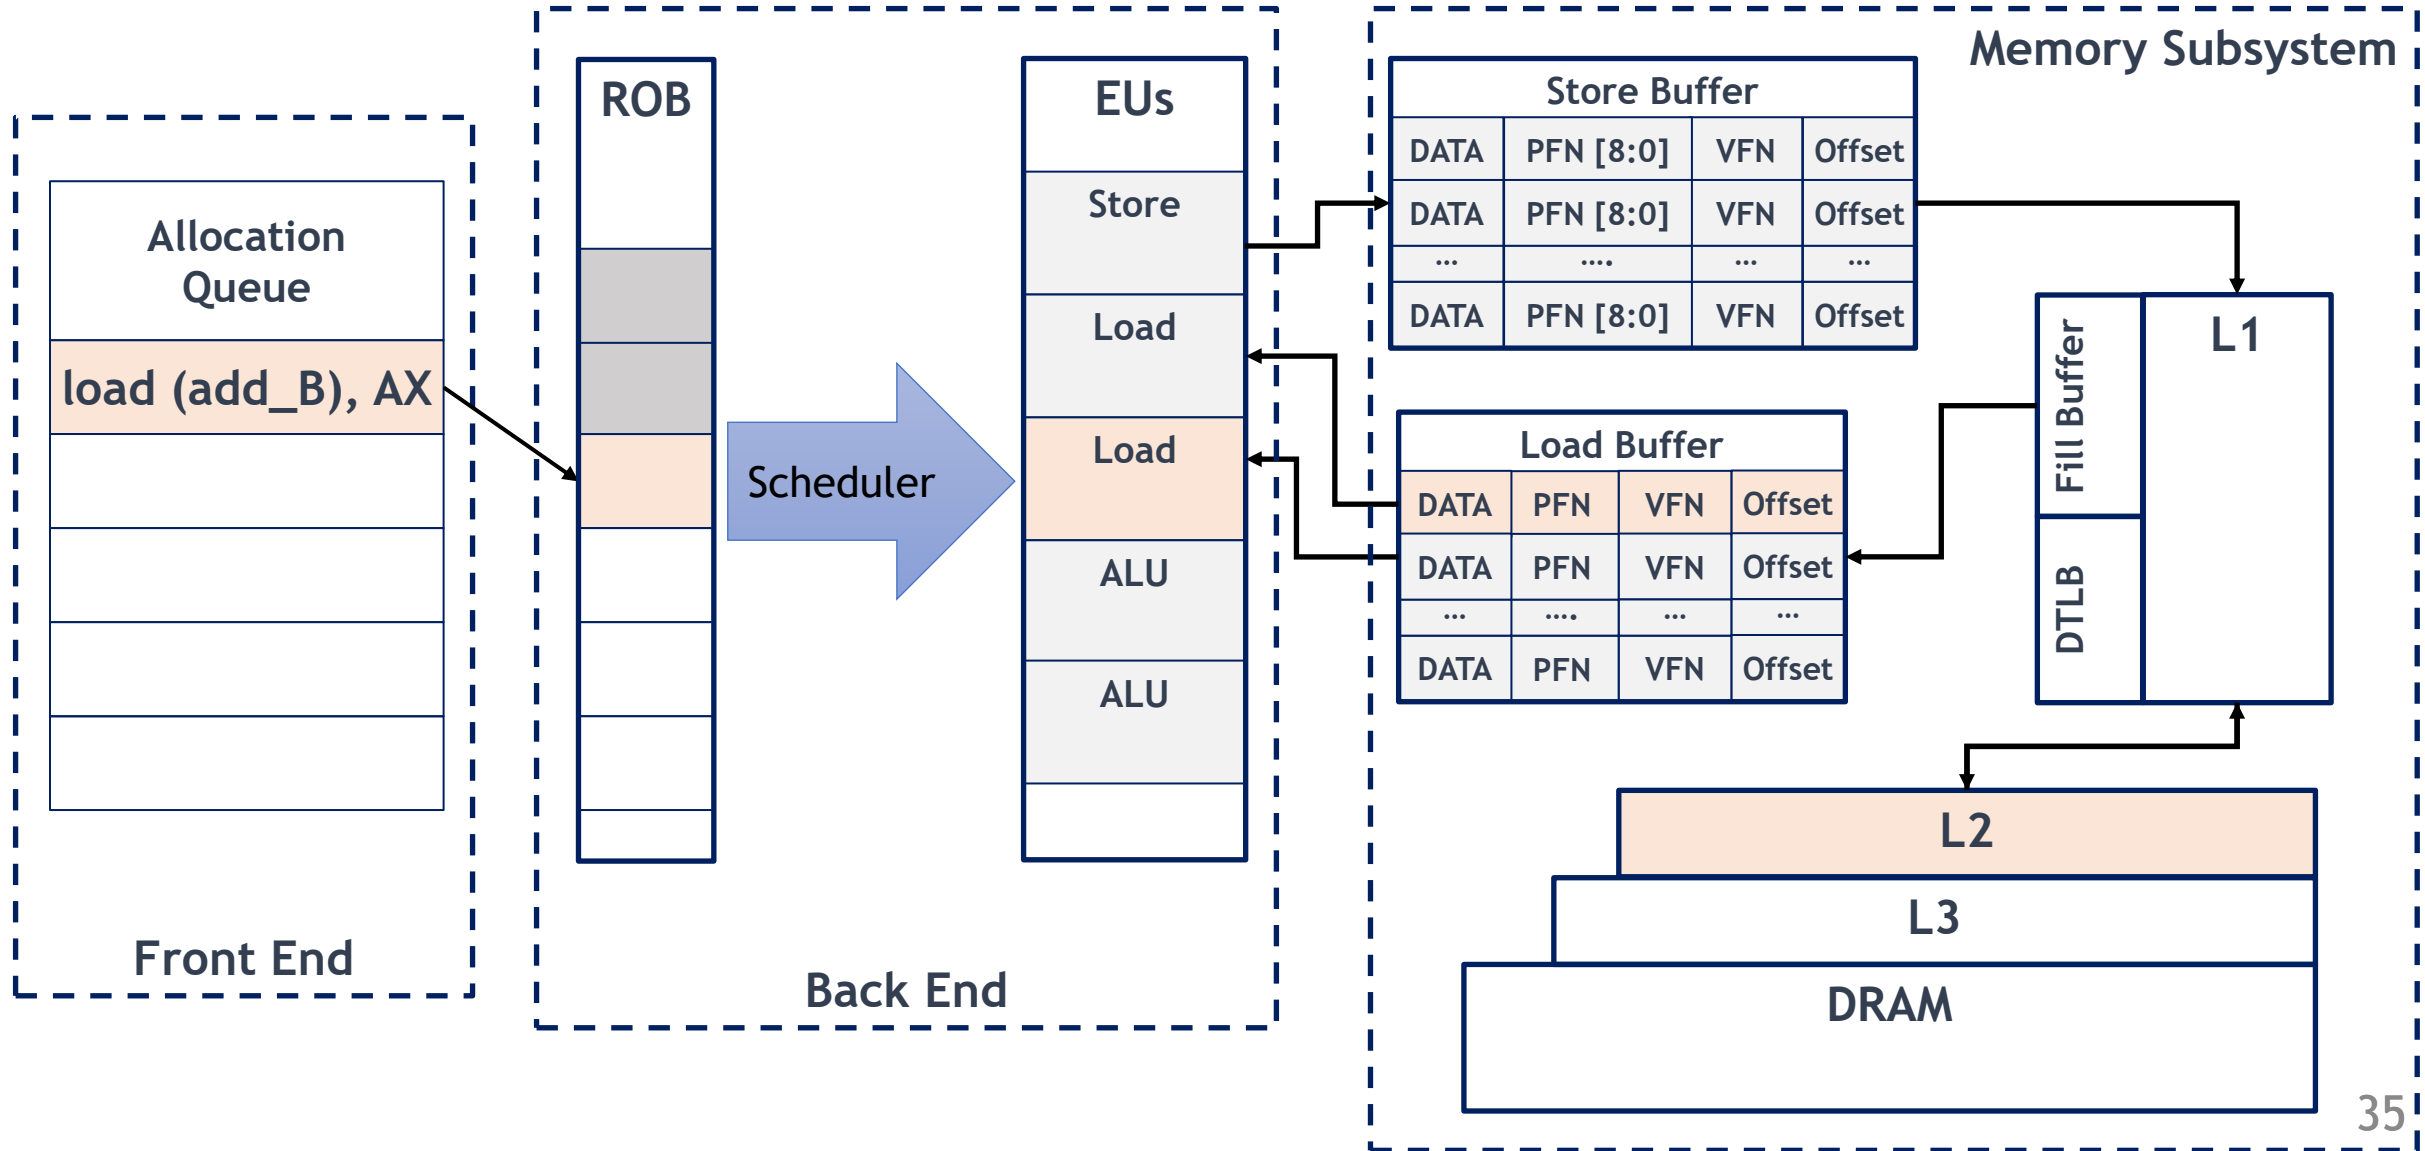
# CPU Memory Subsystem



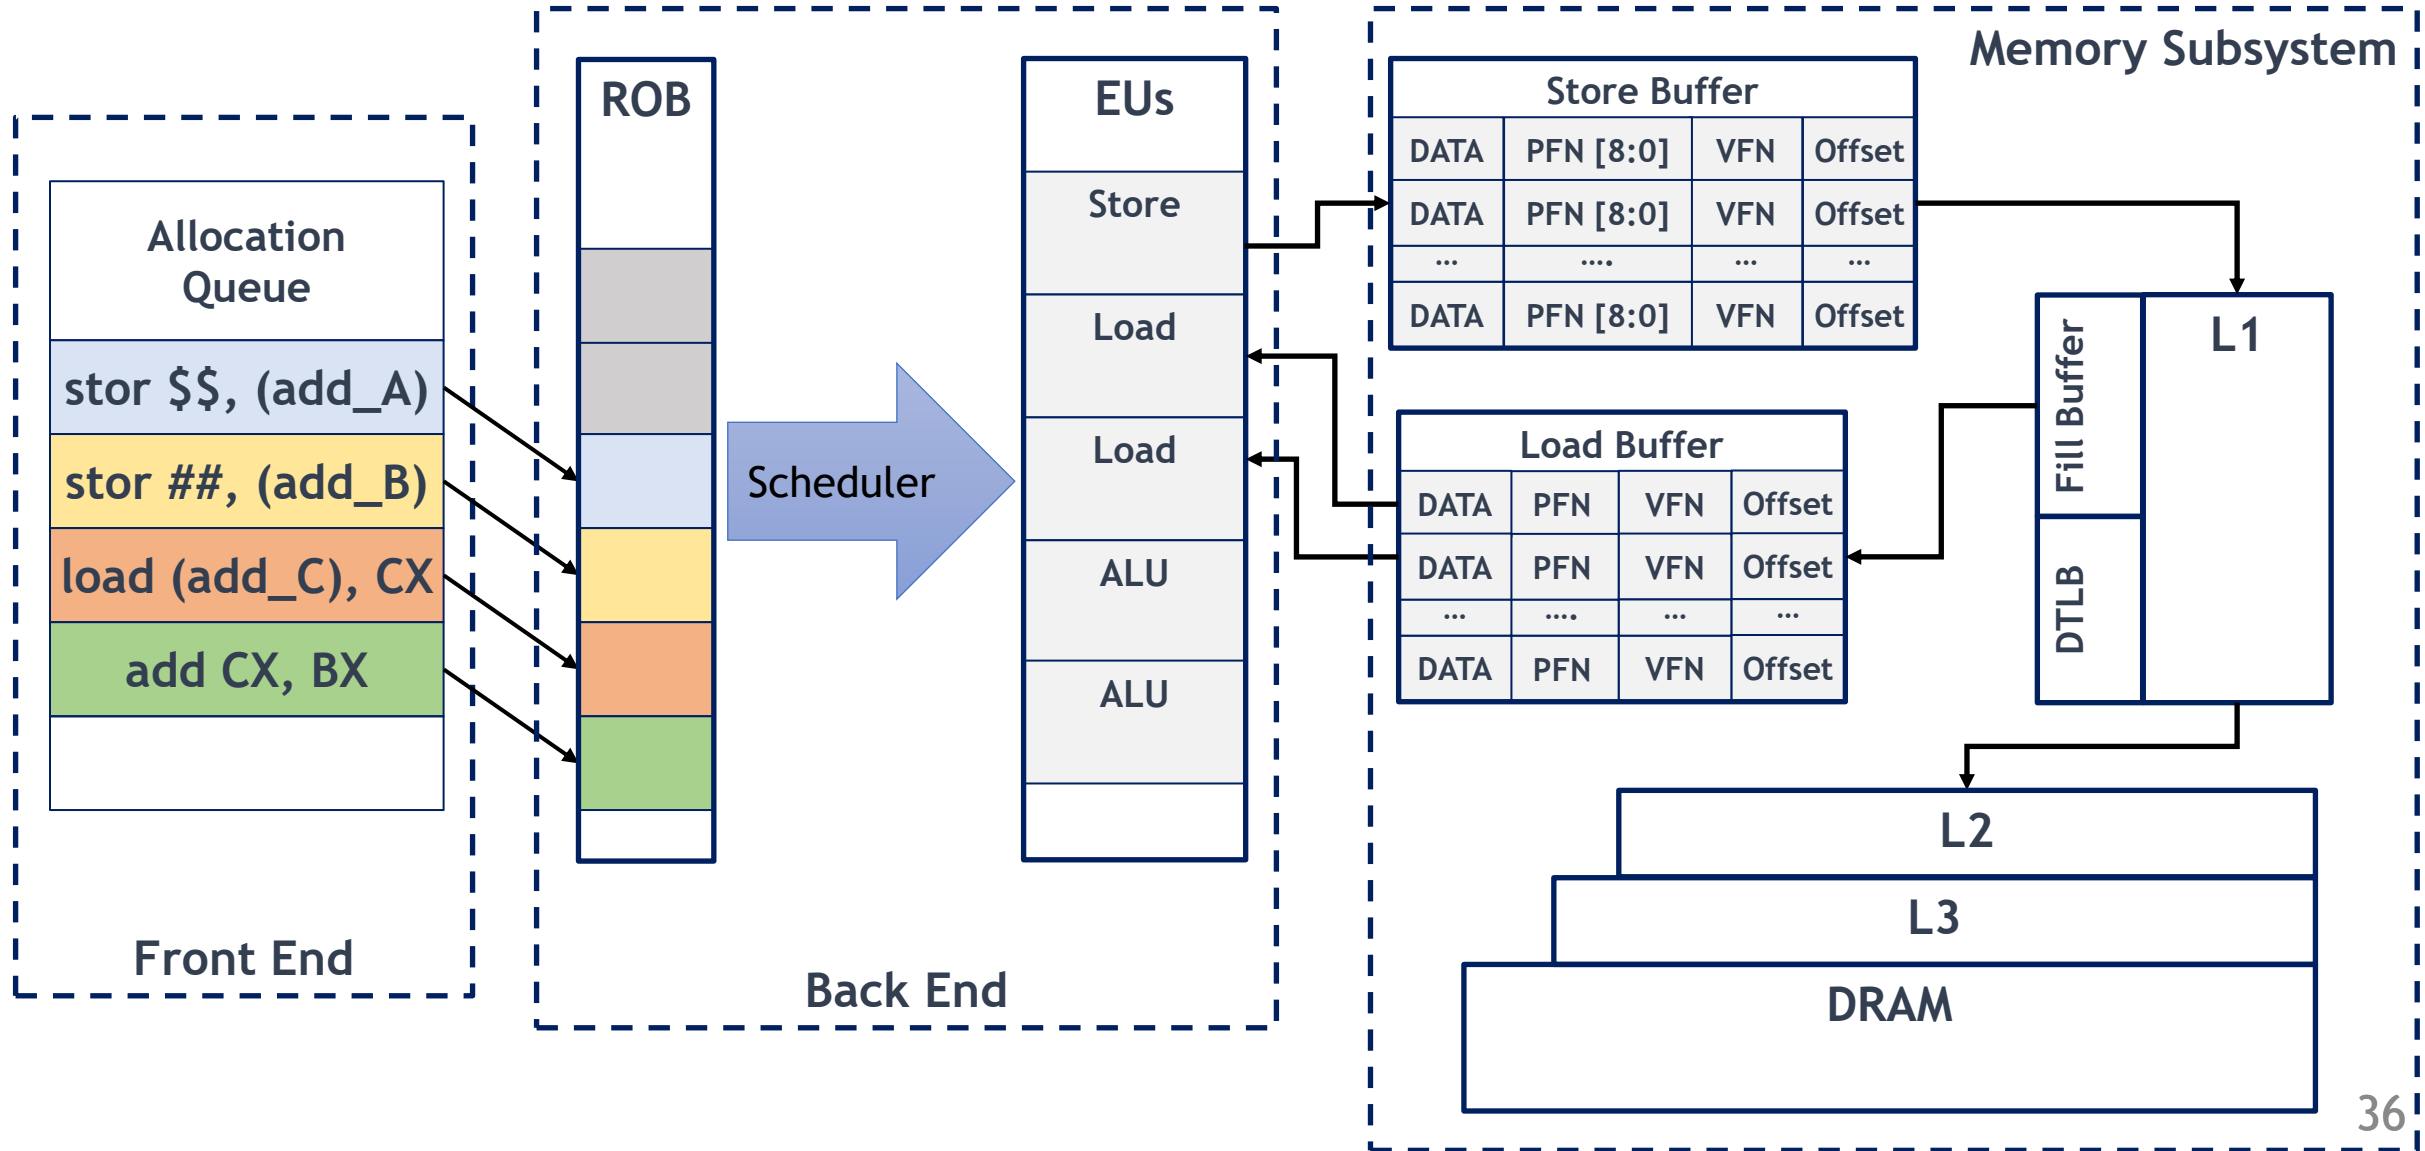
# CPU Memory Subsystem



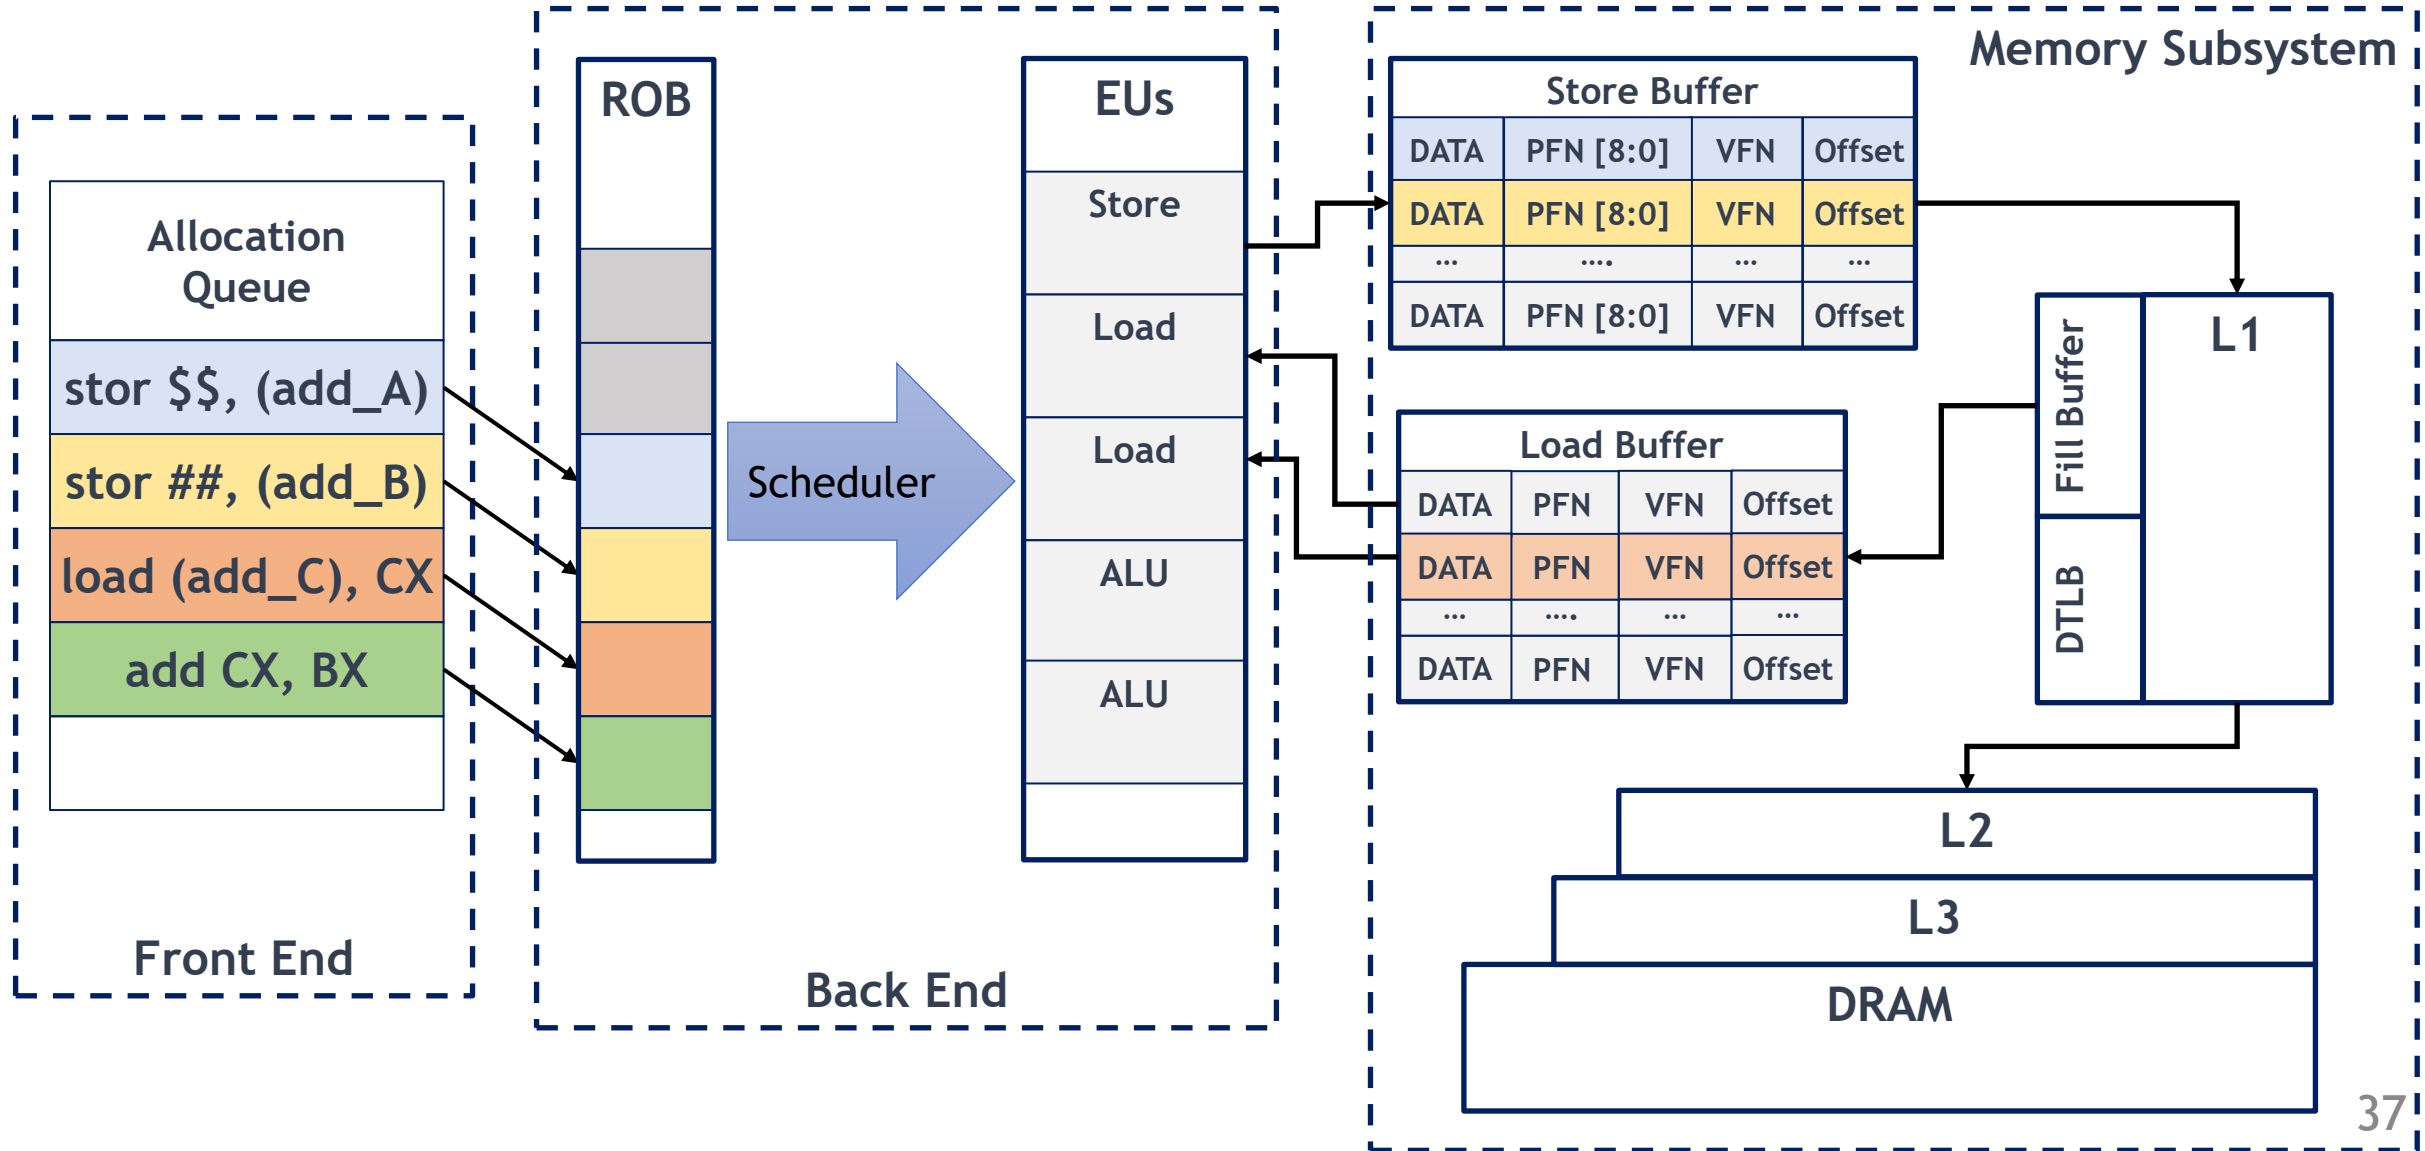
# CPU Memory Subsystem



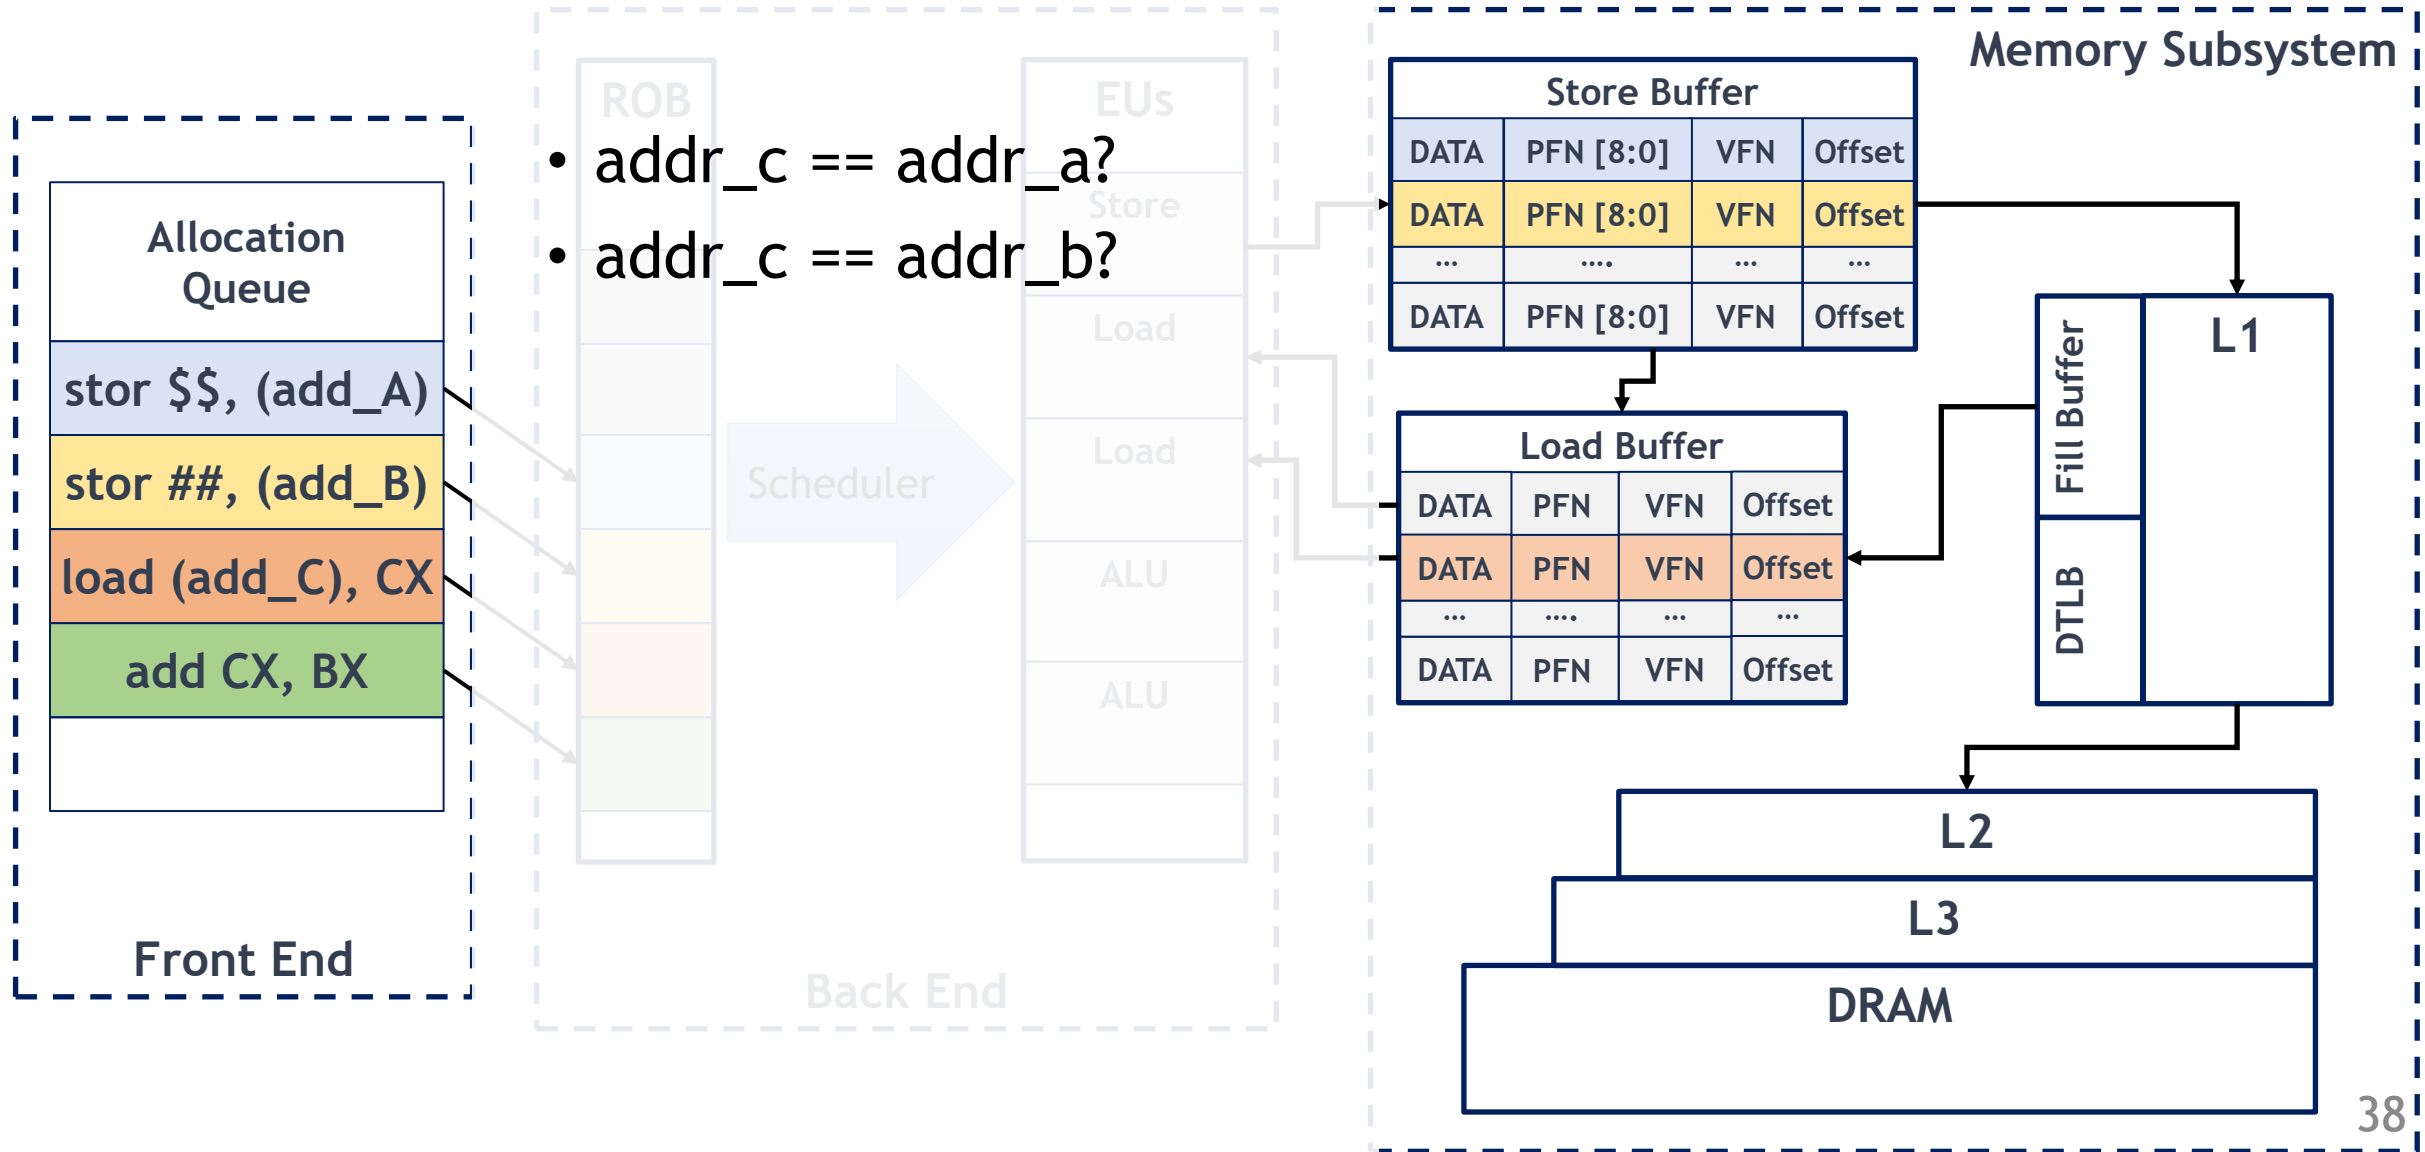
# CPU Memory Subsystem




# CPU Memory Subsystem - Store Forwarding




# CPU Memory Subsystem - Store Forwarding






Virtual Address

VFN	Offset
-----	--------



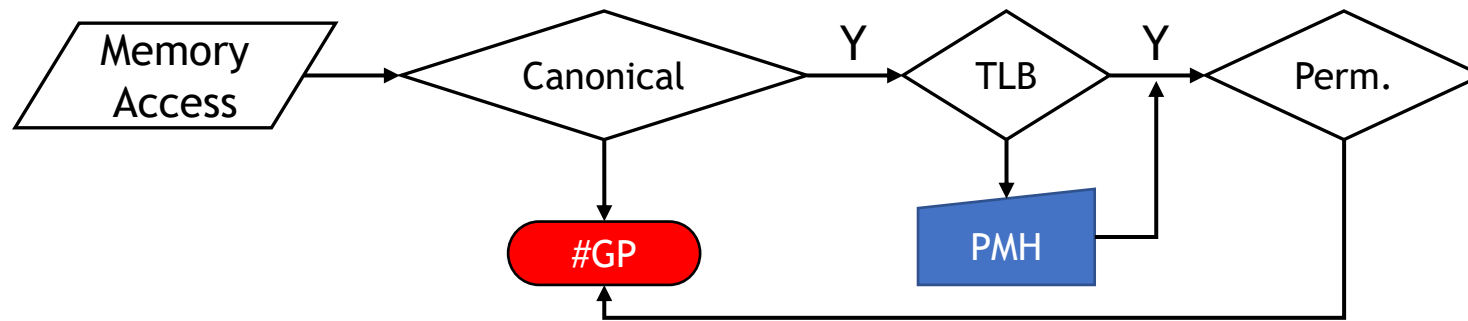
Virtual Address



PTE





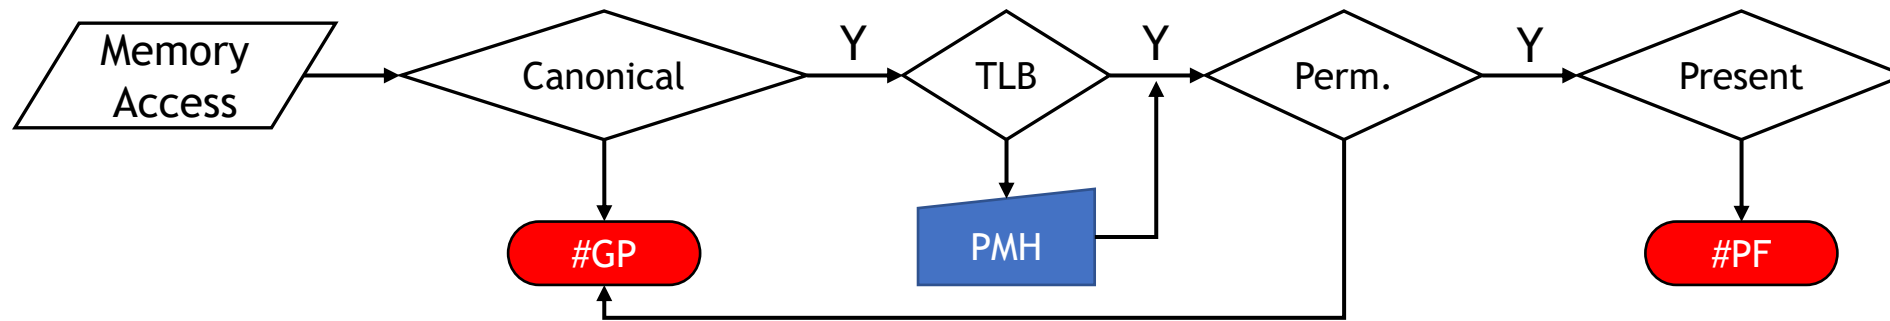


Virtual Address

VFN	Offset
-----	--------

PTE

P	RW	US	...	A	...	Physical Page Number	...
---	----	----	-----	---	-----	----------------------	-----




Virtual Address

VFN	Offset
-----	--------

PTE

P	RW	US	...	A	...	Physical Page Number	...
---	----	----	-----	---	-----	----------------------	-----

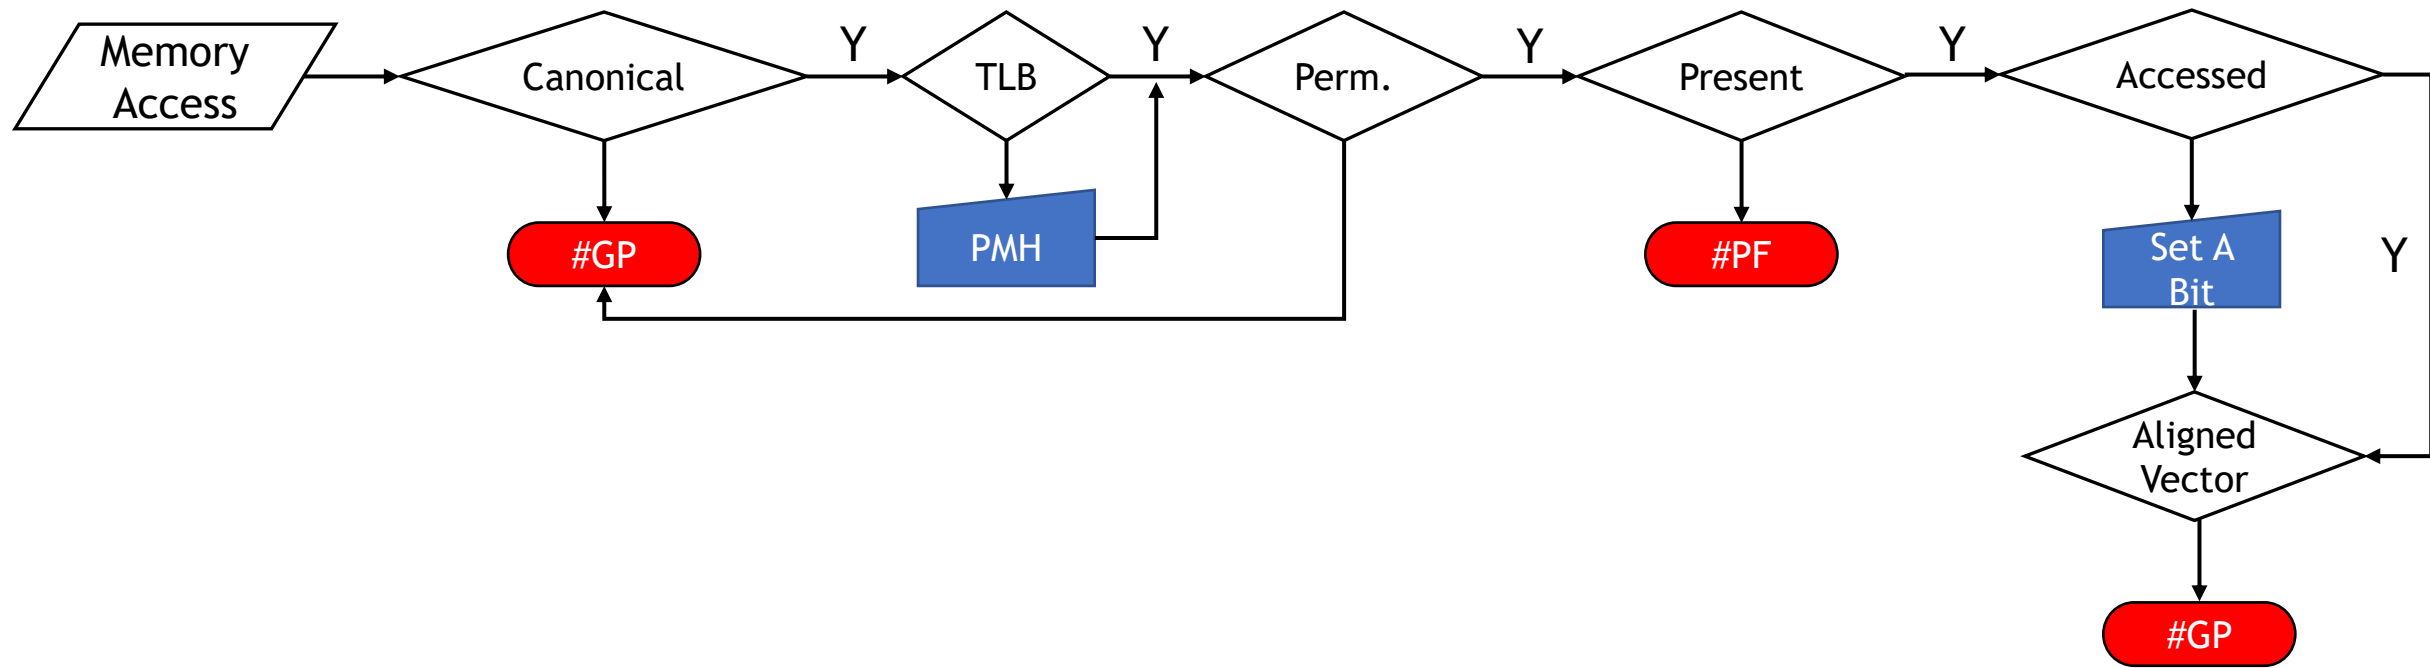


Virtual Address

VFN	Offset
-----	--------

PTE

P	RW	US	...	A	...	Physical Page Number	...
---	----	----	-----	---	-----	----------------------	-----

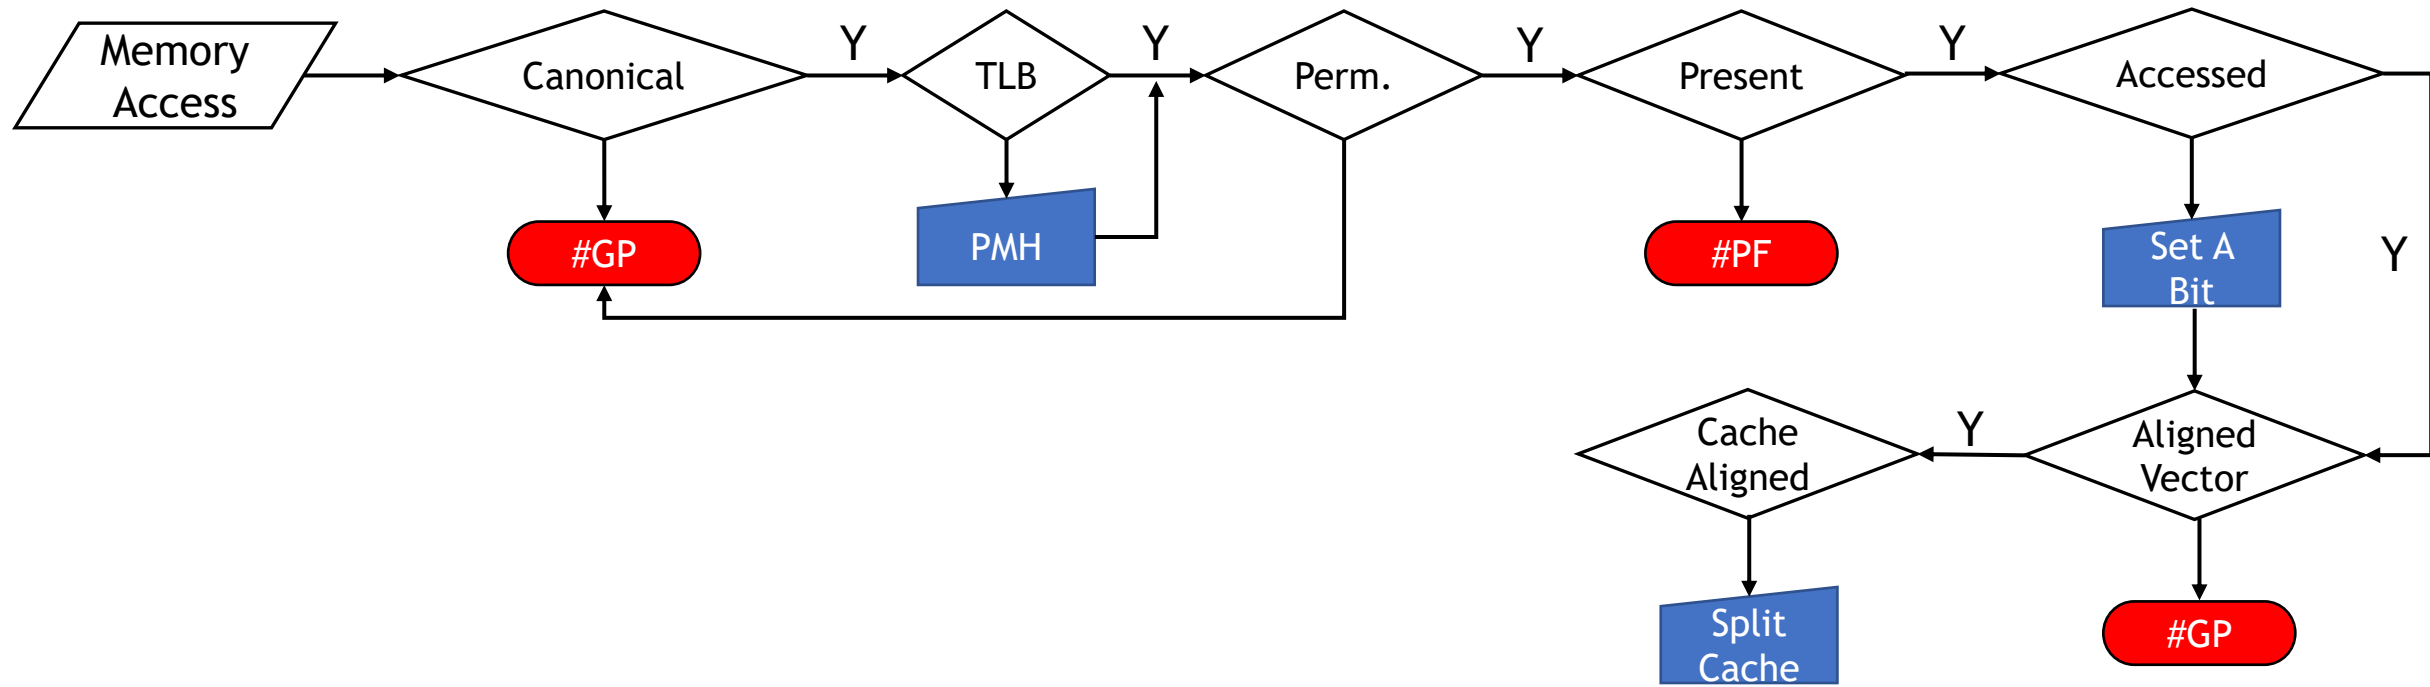


Virtual Address


VFN	Offset
-----	--------

PTE

P	RW	US	...	A	...	Physical Page Number	...
---	----	----	-----	---	-----	----------------------	-----

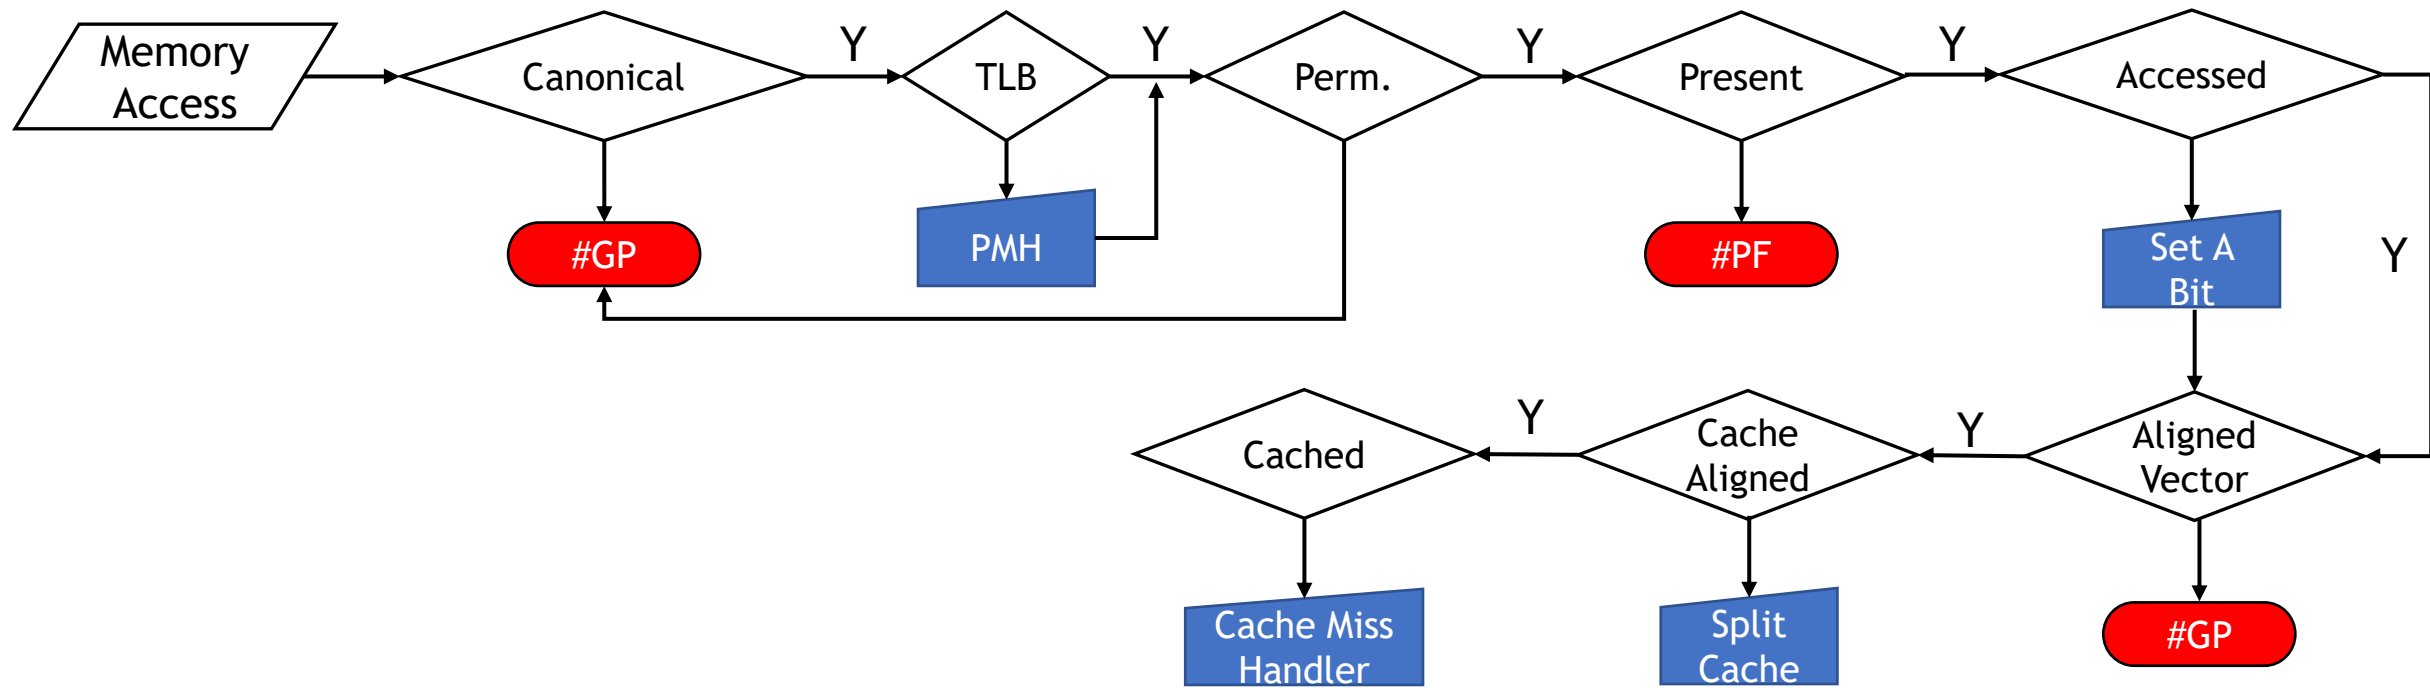


Virtual Address




PTE



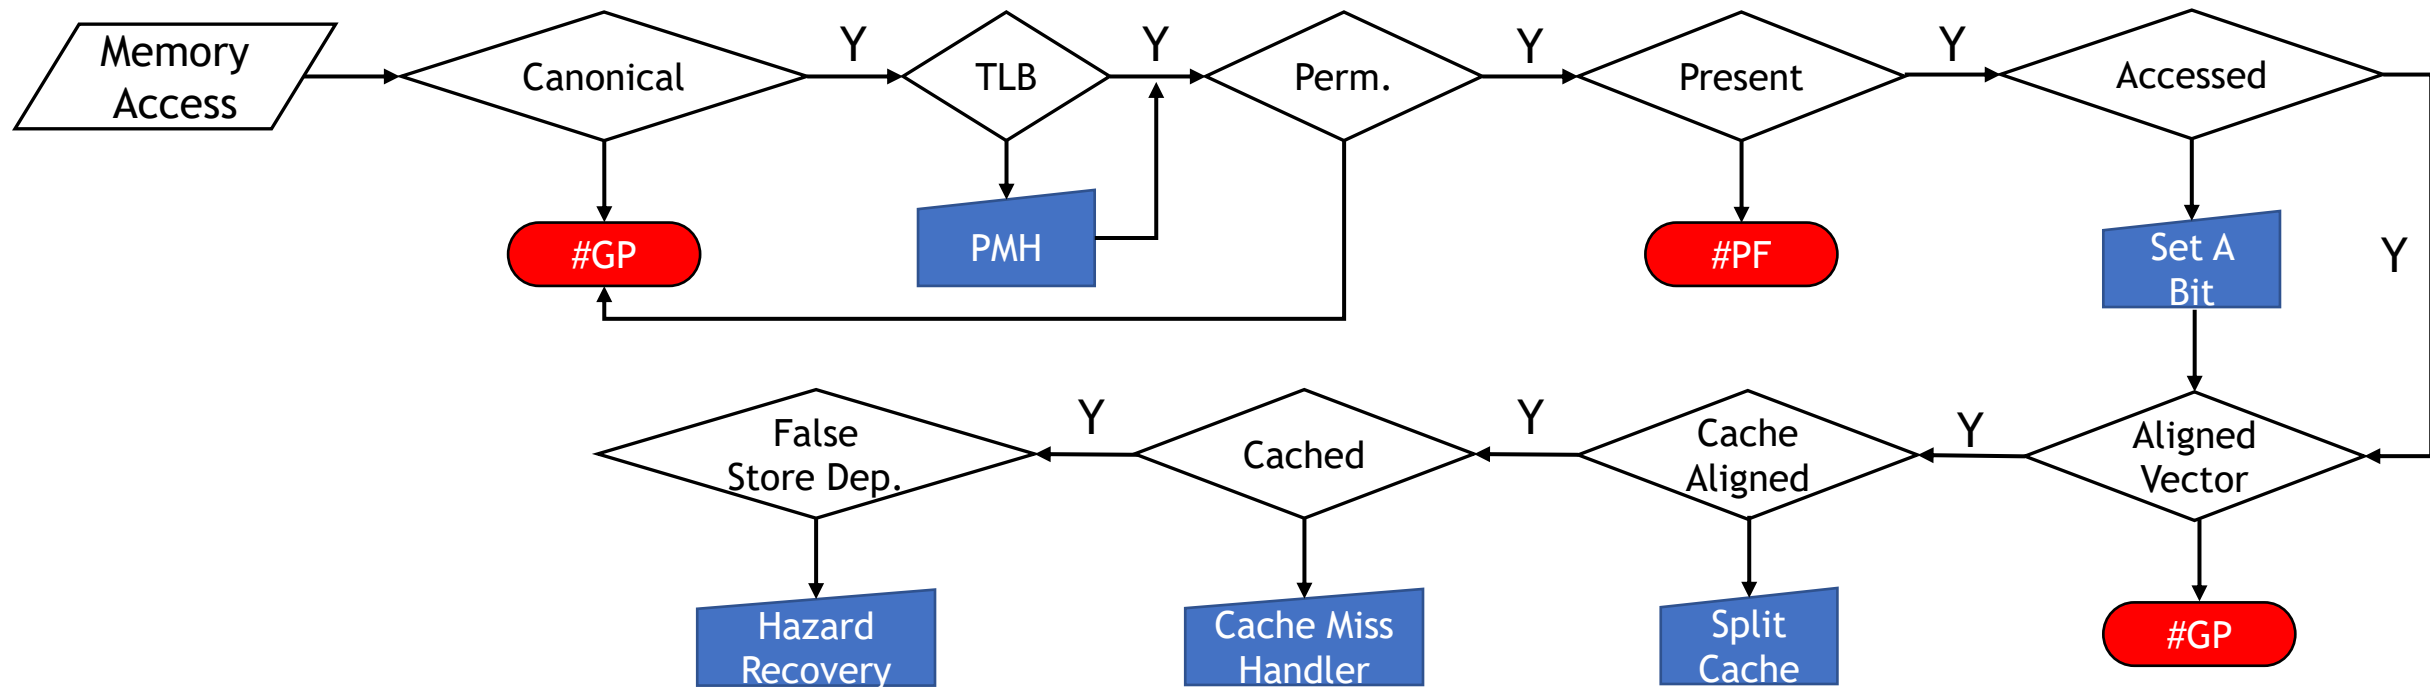


Virtual Address




PTE



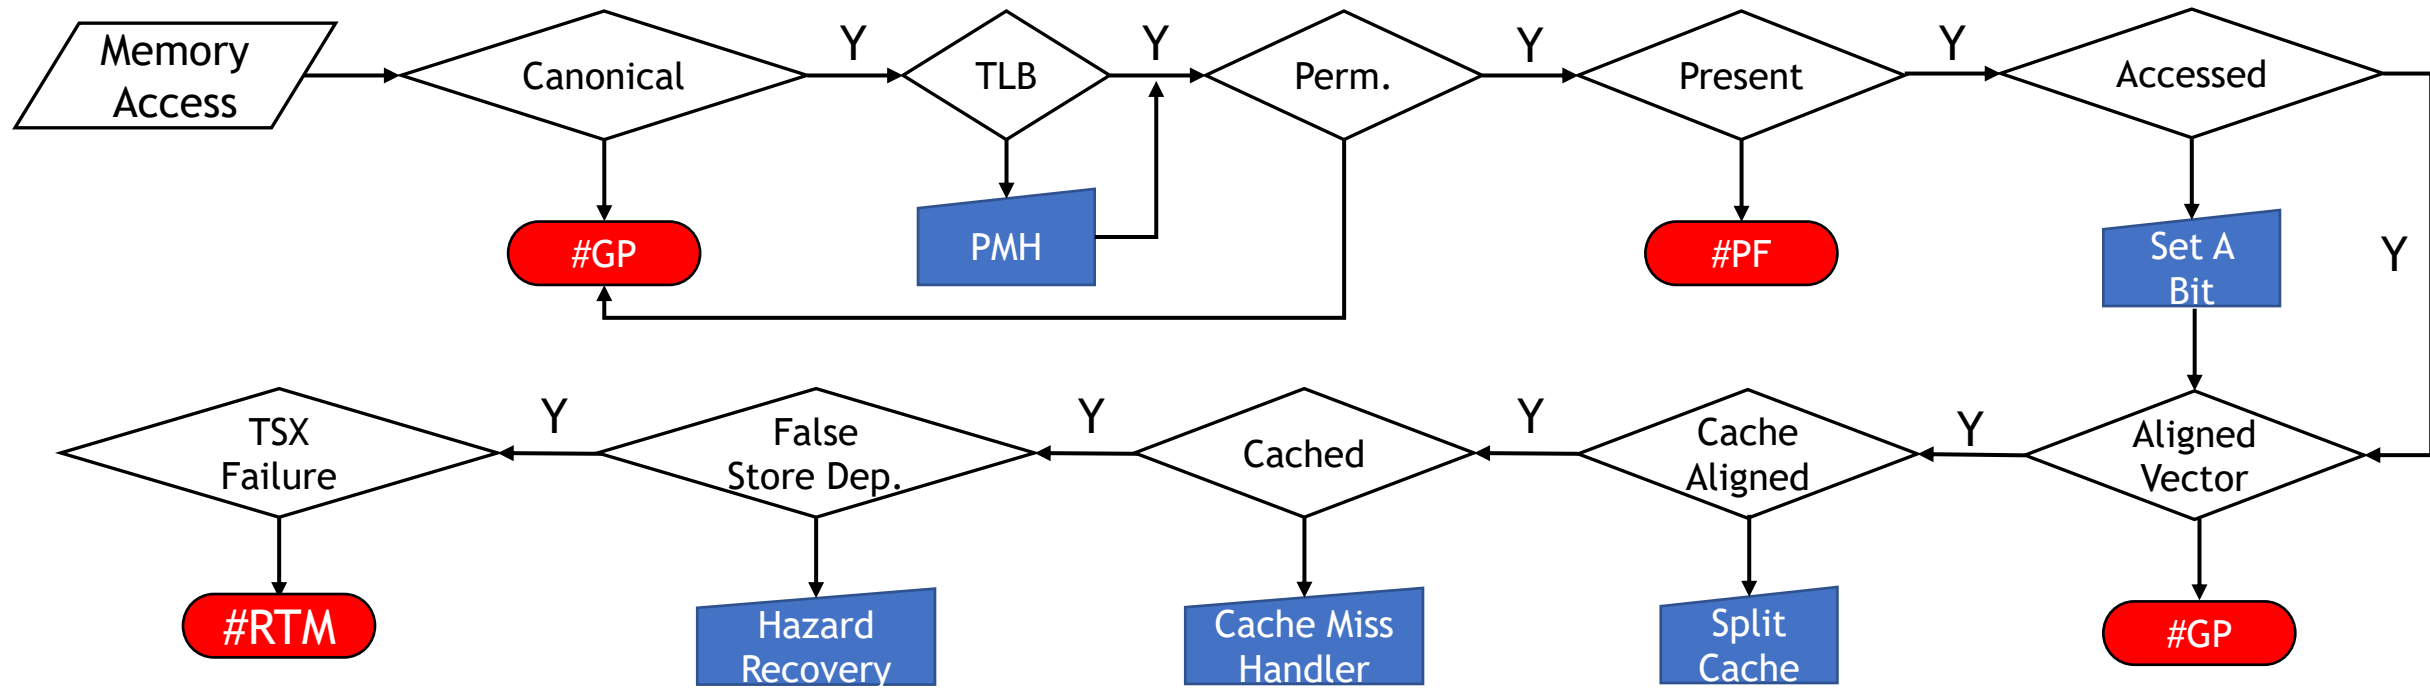


Virtual Address



PTE





Virtual Address

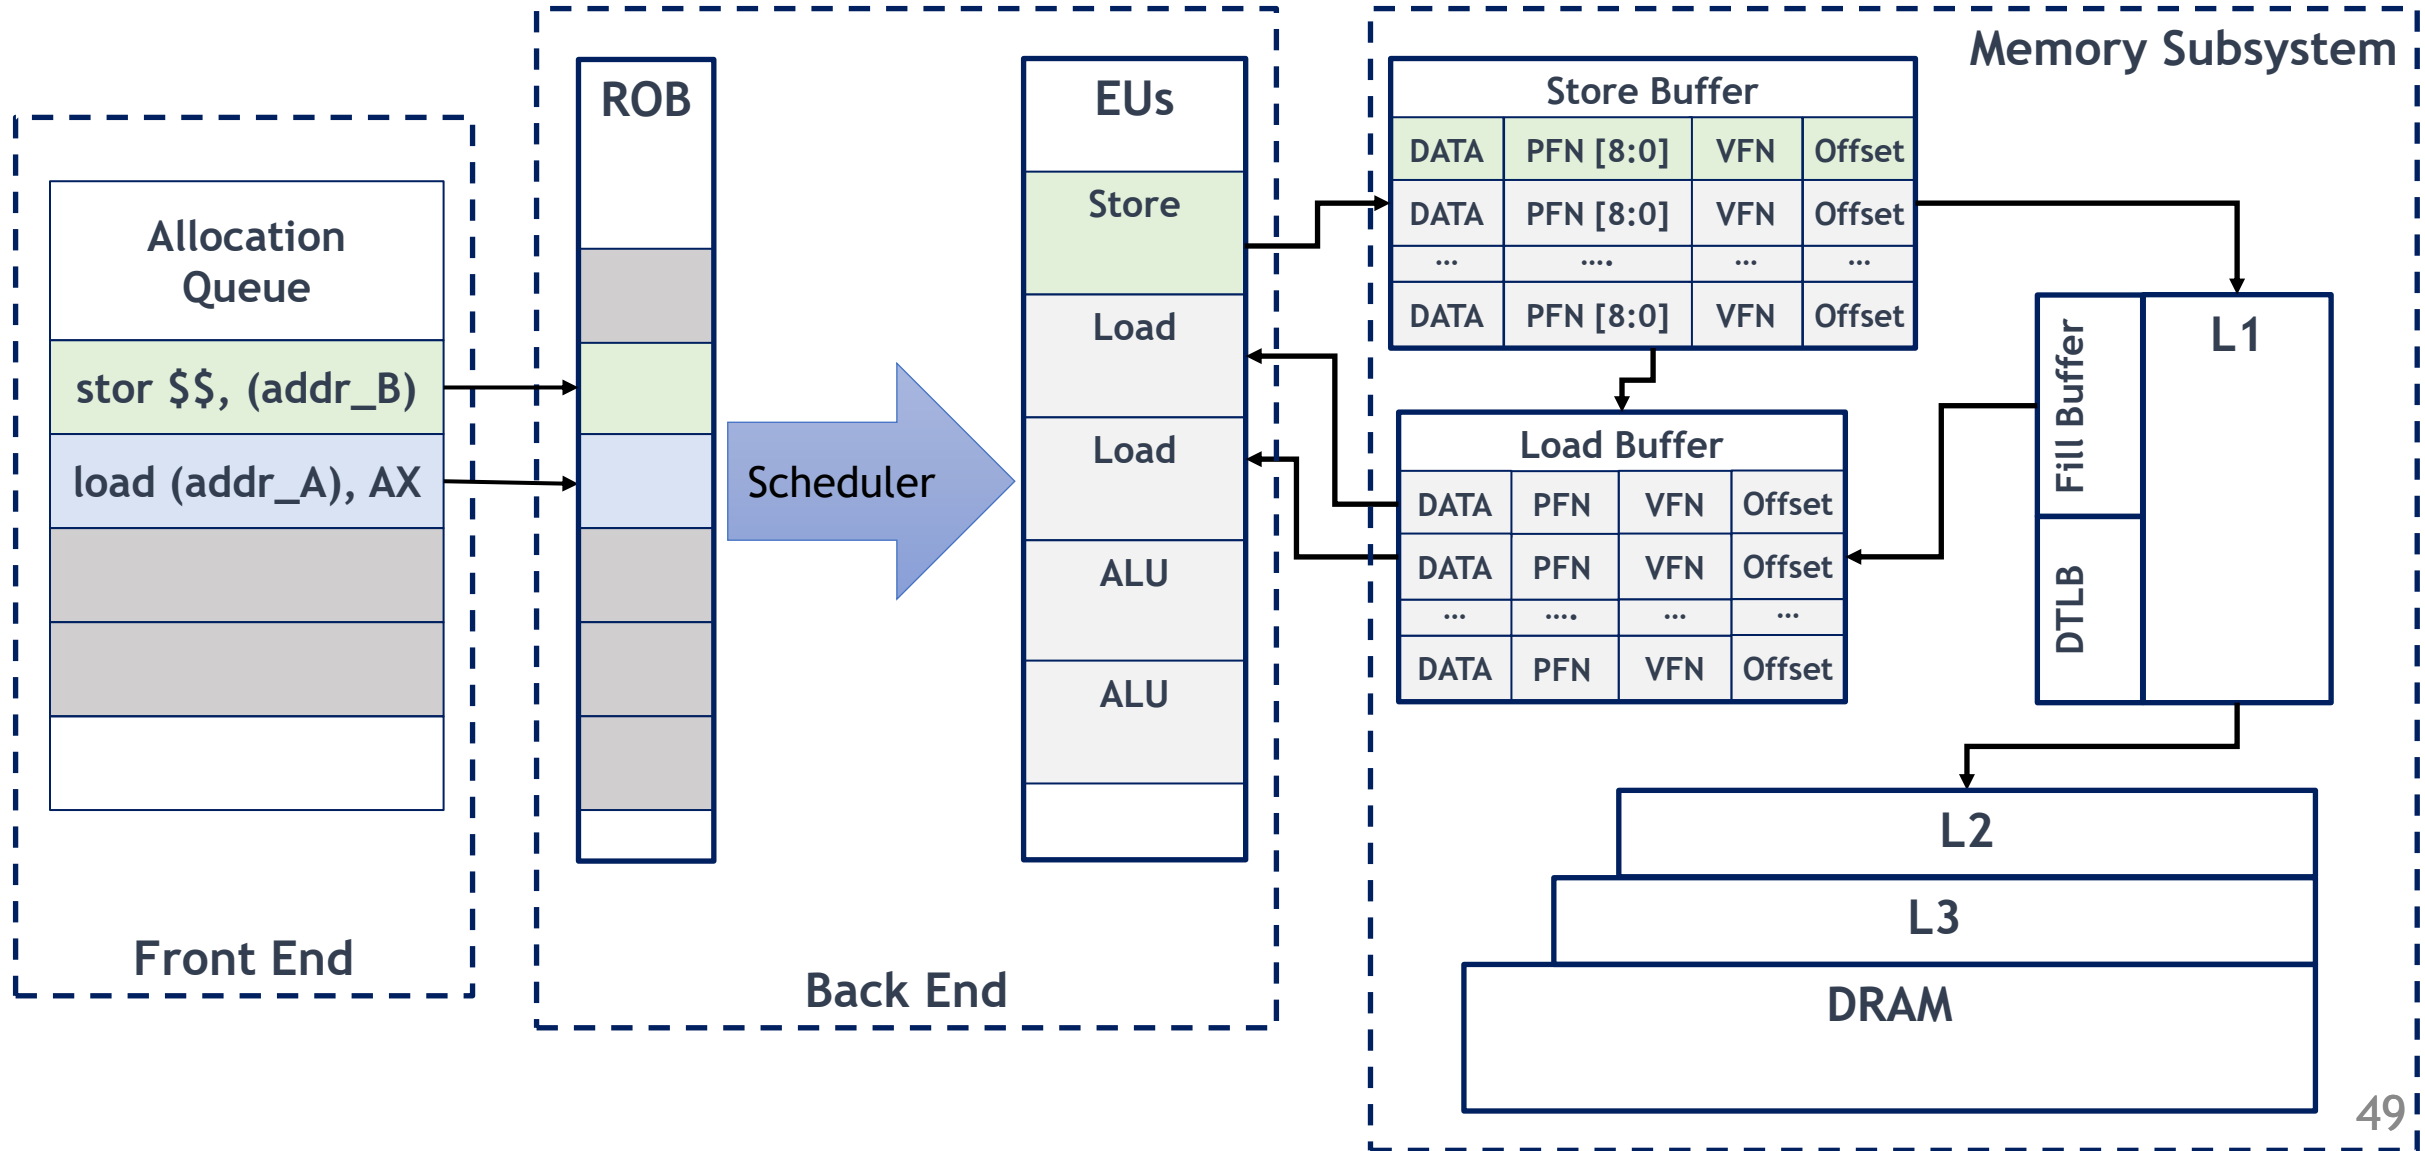


PTE

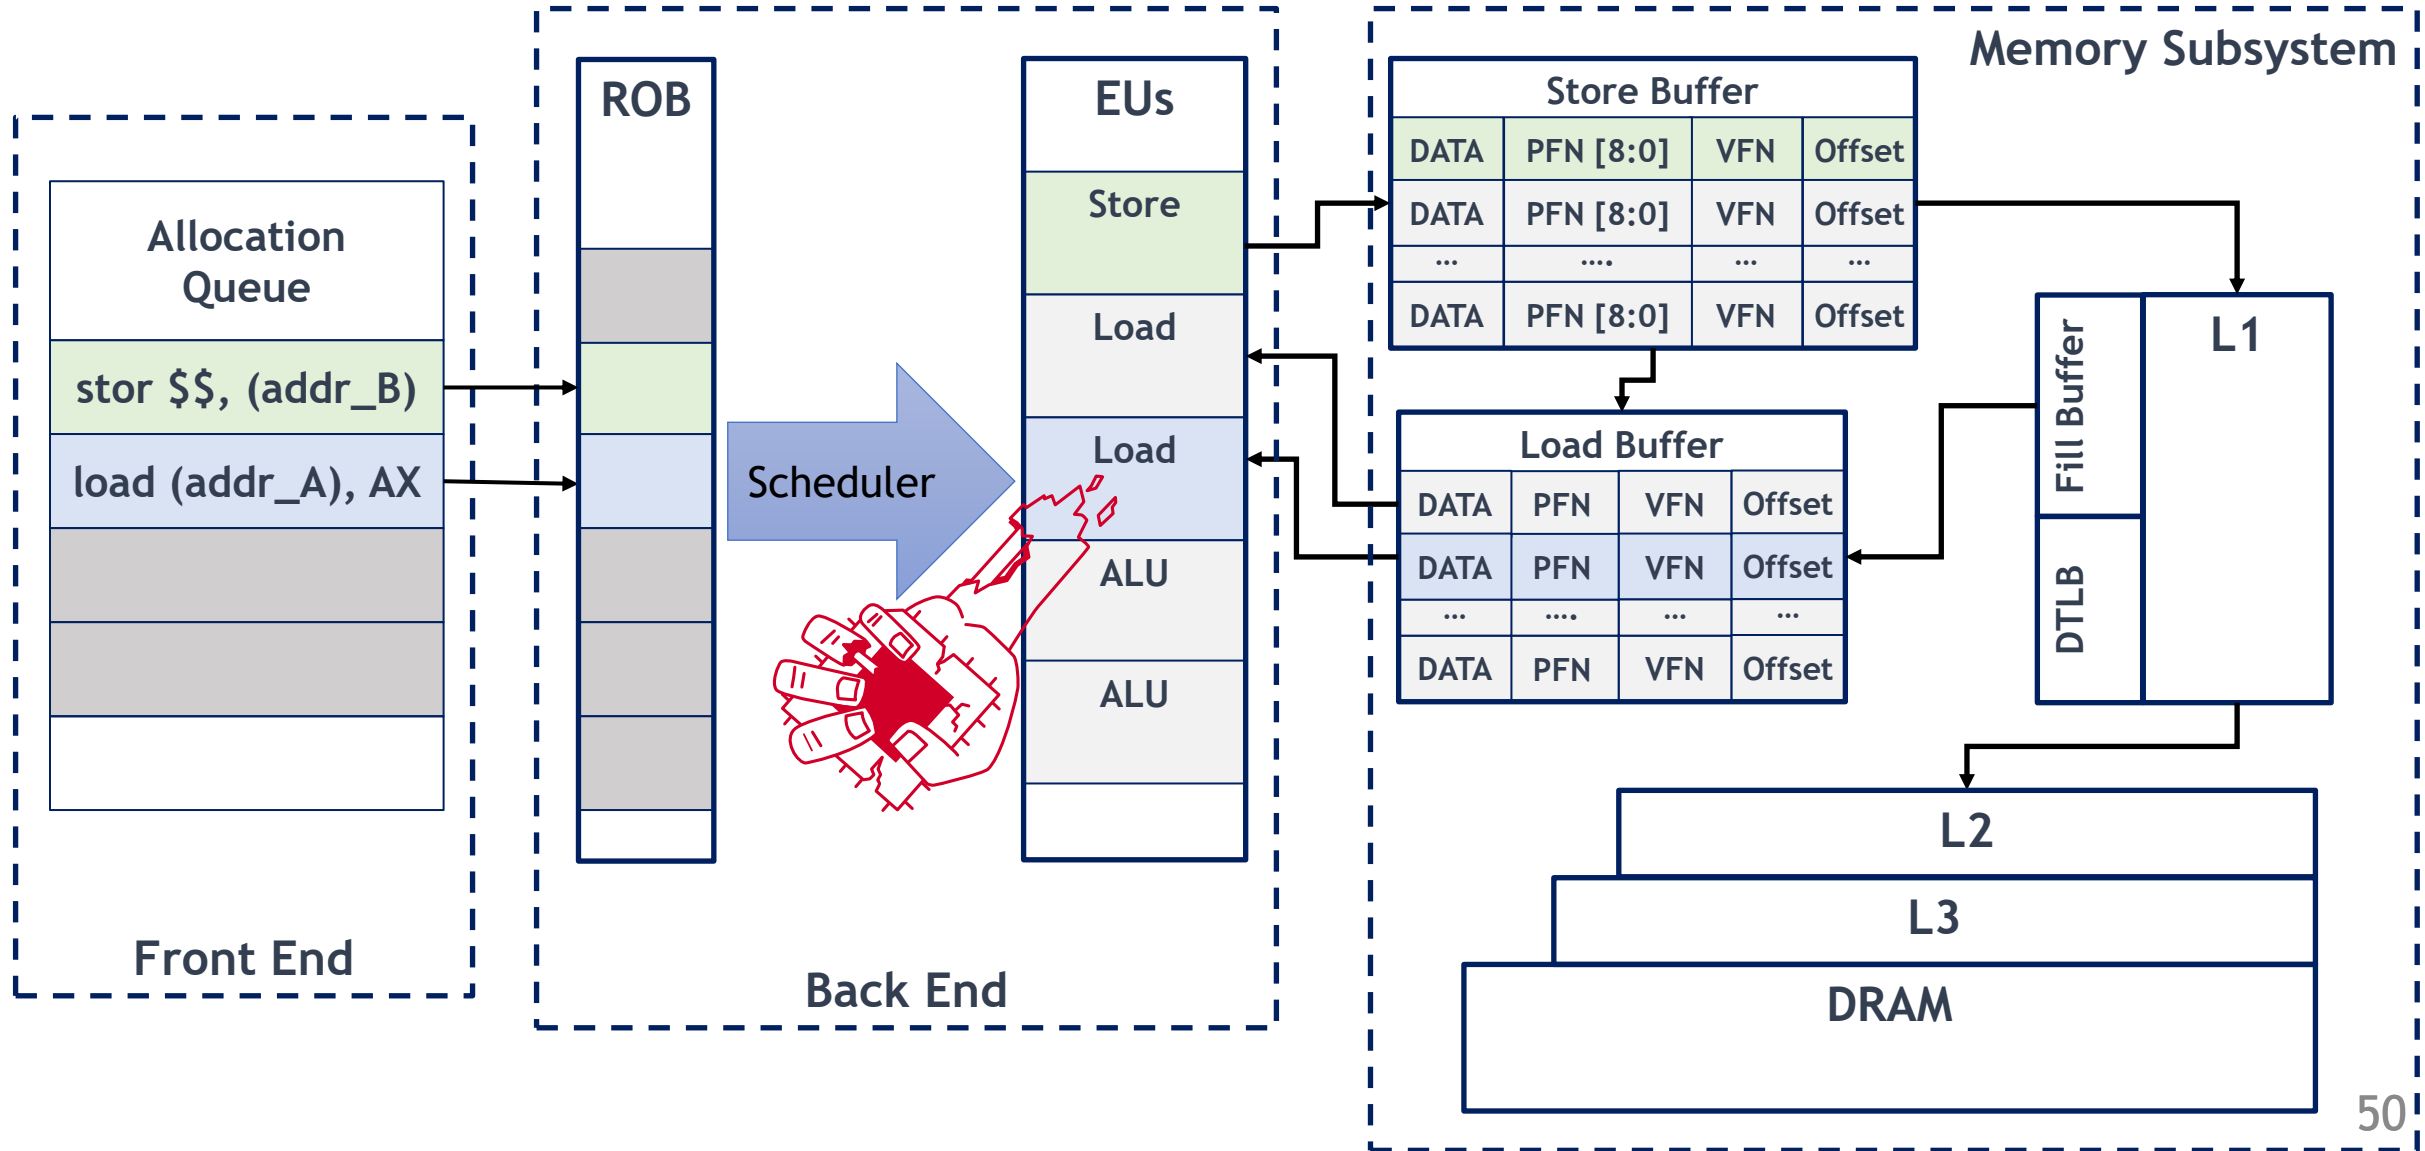




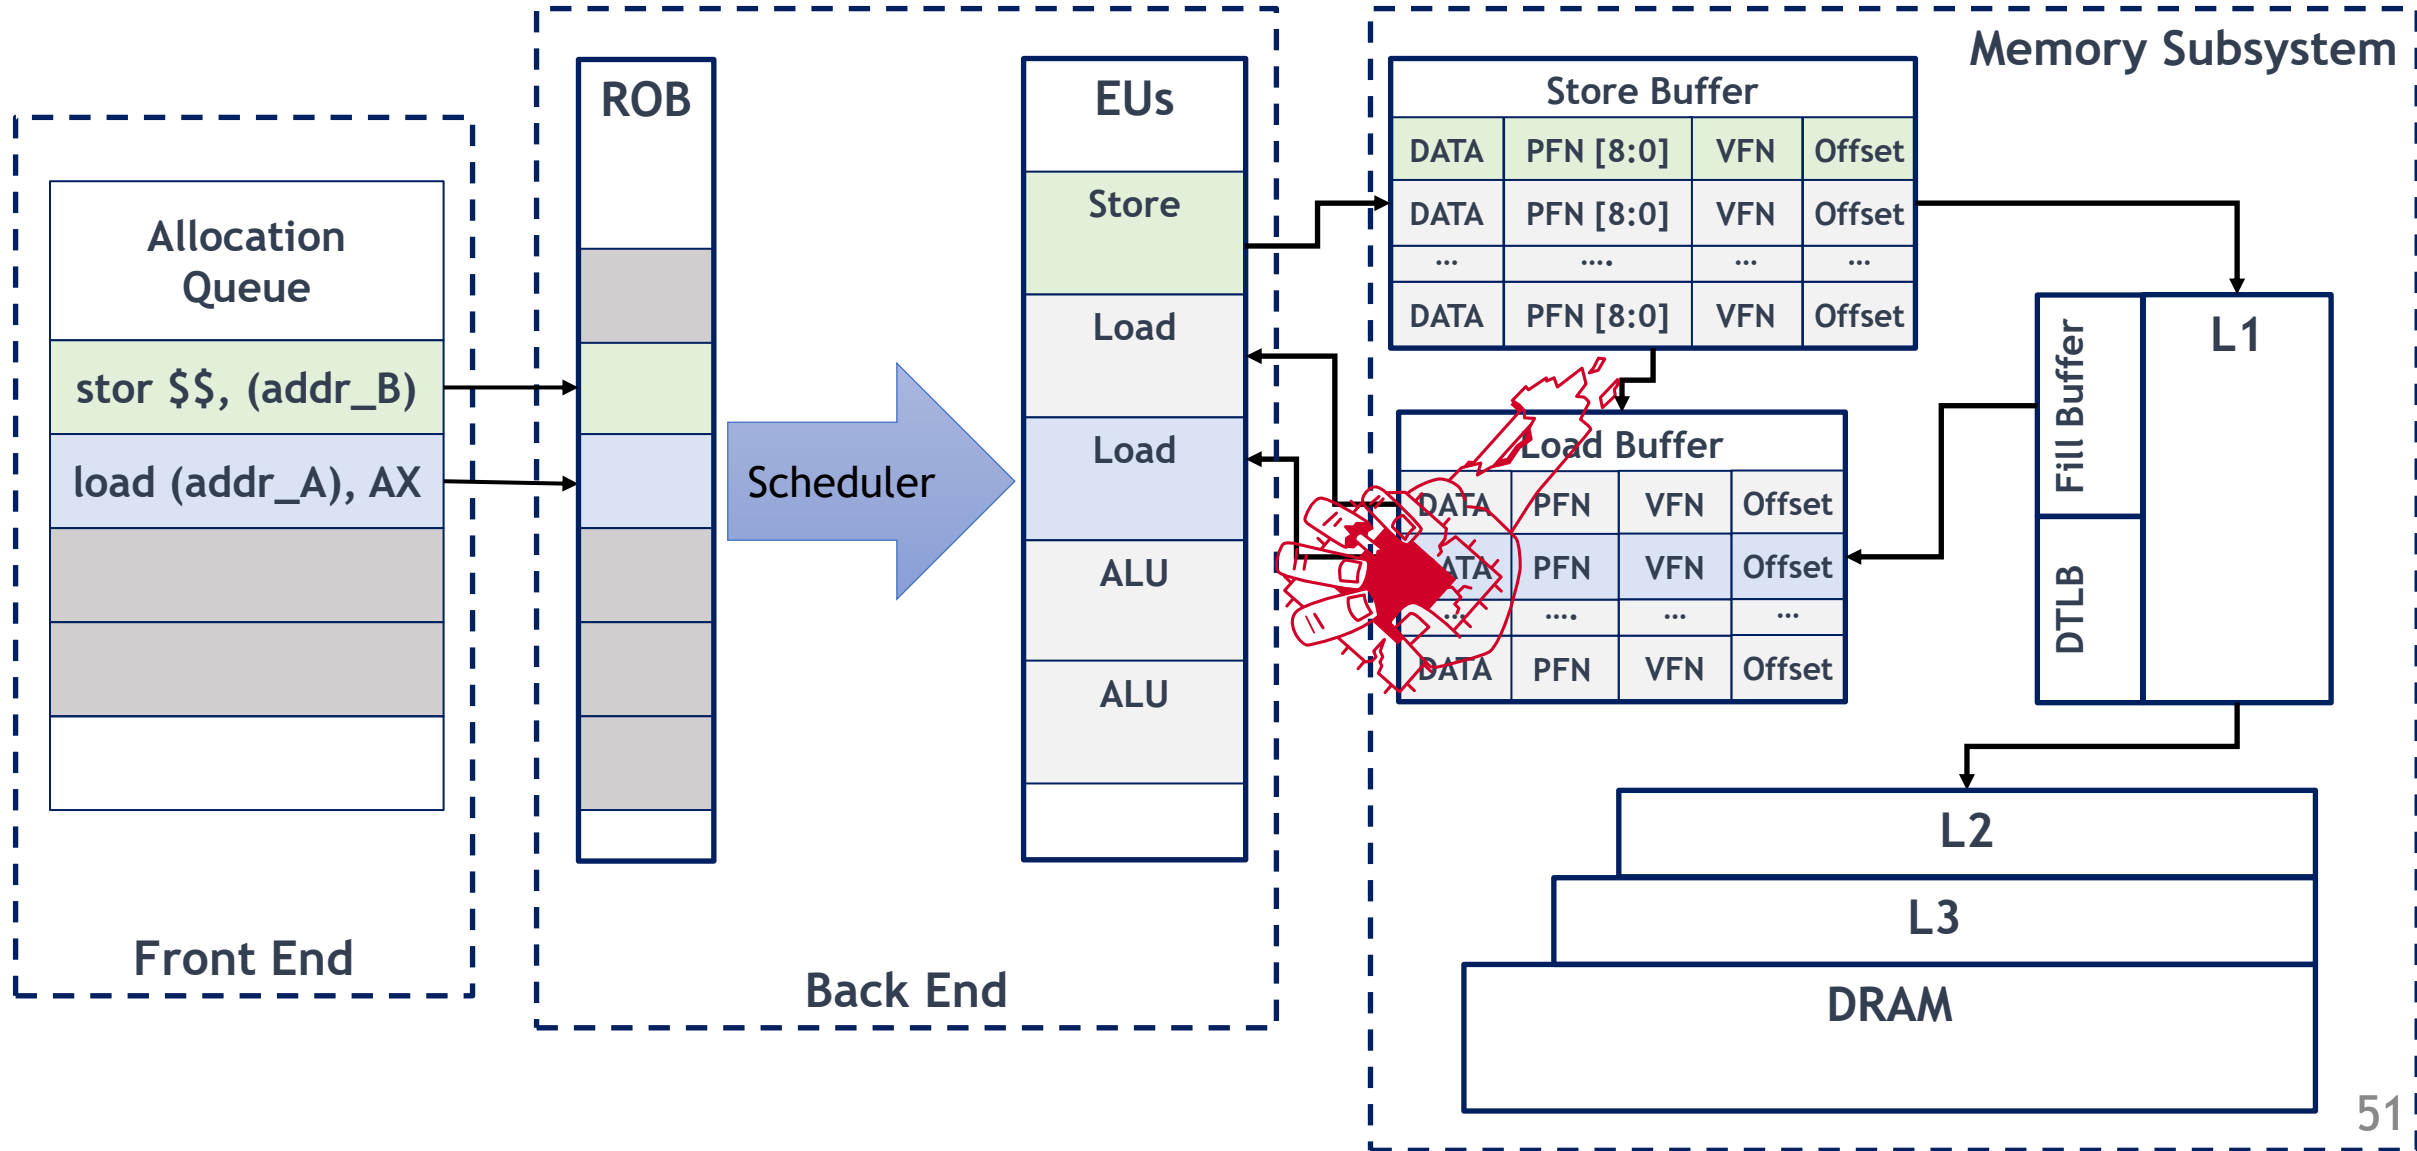
# CPU Memory Subsystem - Hazard Recovery



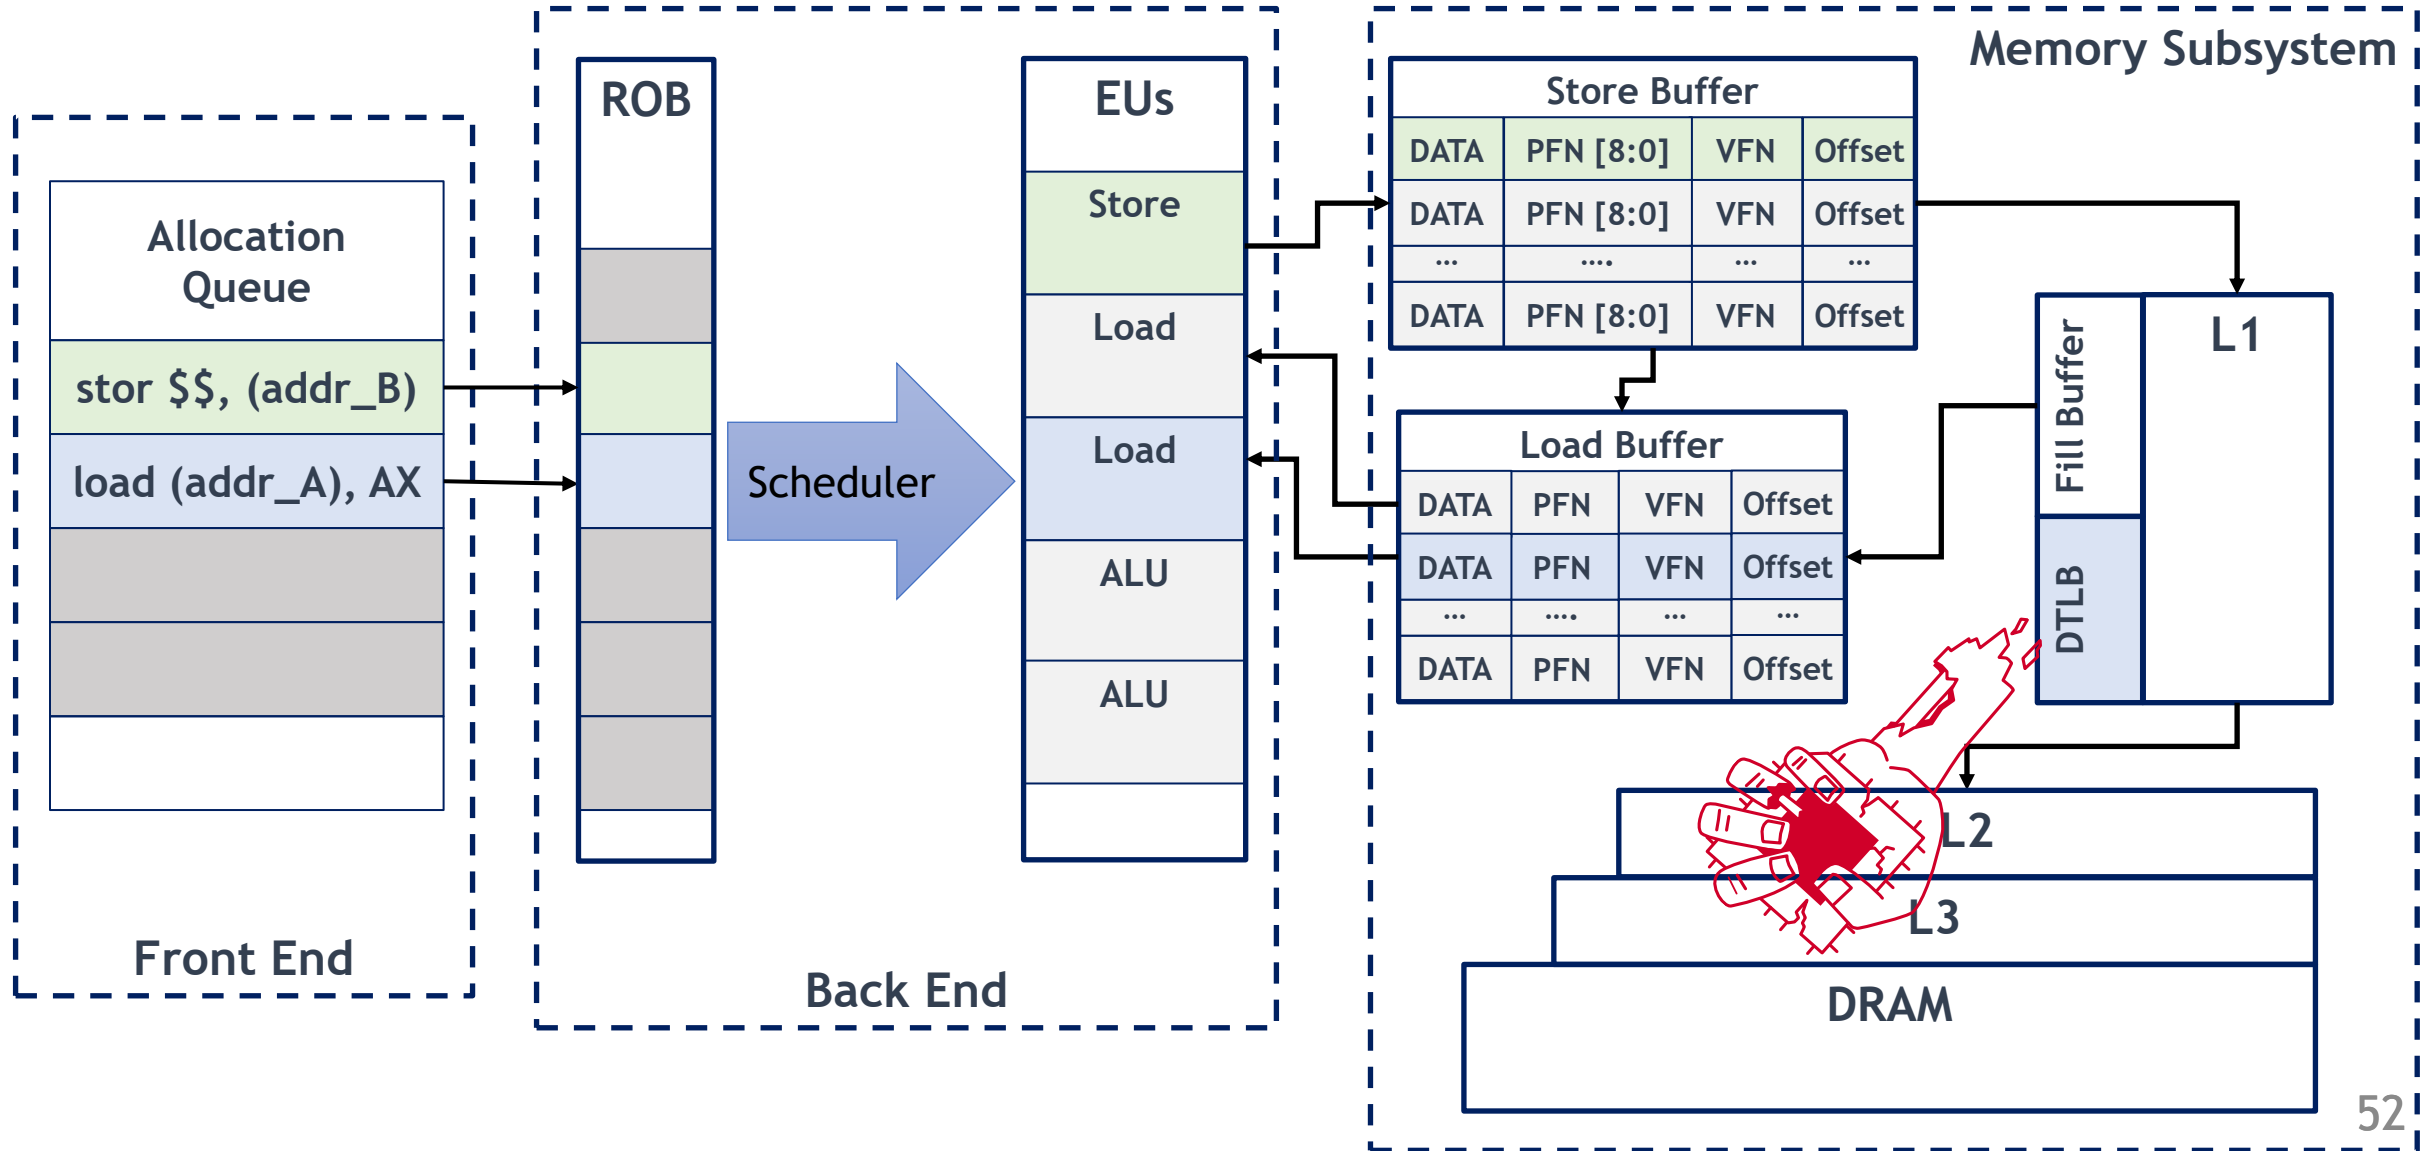
# CPU Memory Subsystem - Hazard Recovery



# CPU Memory Subsystem - Hazard Recovery




# CPU Memory Subsystem - Hazard Recovery



# Challenges with MDS Testing?

- Reproducing attacks is not reliable. It may depend on:
  - massaging the pipeline with other instructions
  - CPU configuration (generation, frequency, microcode patch and etc)
- No public tool to find new variants or to verify hardware patches:
  - Too many things to test (Addressing mode, cache state, assists, and faults)
  - Previous POCs may not work after MC update, but what does it mean?
- Impossible to quantify the impact of leakage:
  - We should care about leakage rate and what data is leaked.
  - My POC is faster than your POC!!

Let's see this problem in action?! (Demo)



# Transynther

# Transynther (Fuzzing-based Random MDS Testing)


Step 1:

```
char secret = *(char *) 0xffffffff81a0123;
```

Step 2:

```
char x = oracle[secret * 4096];
```

Step 3:



256 different CPU Cache Line



# Transynther (Fuzzing-based Random MDS Testing)

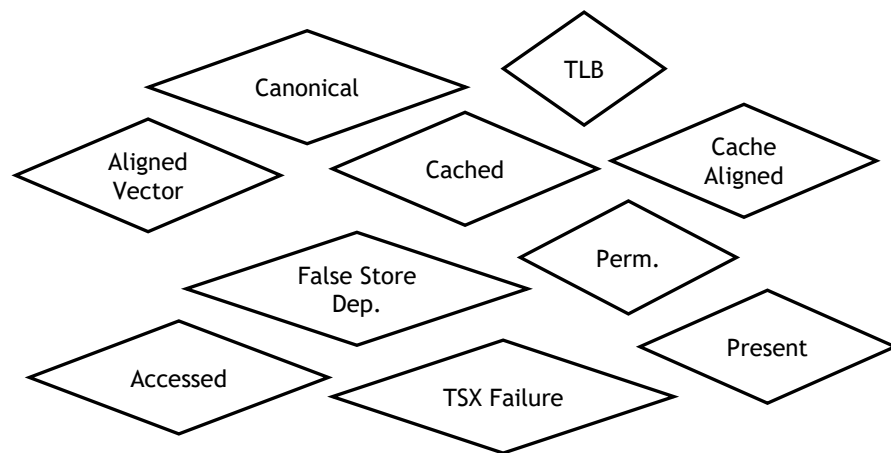
Step 1:



Step 2:



Step 3:



```
char x = oracle[secret * 4096];
```

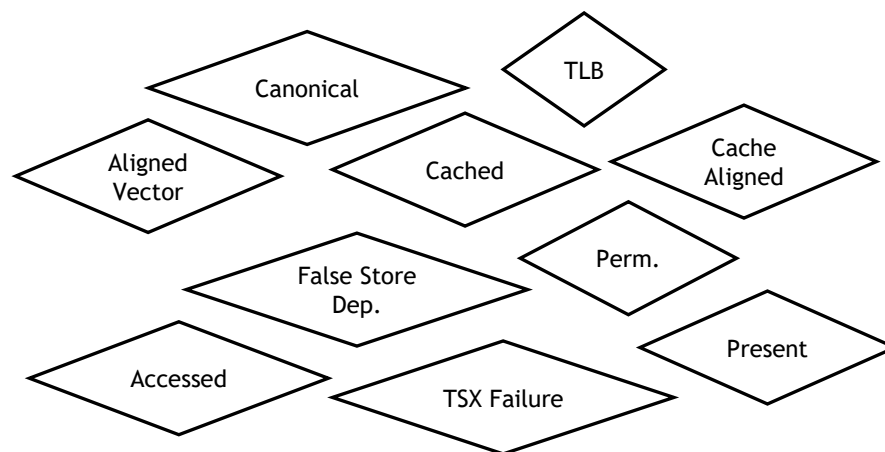

'P' = 0x50



256 different CPU Cache Line

# Transynther (Fuzzing-based Random MDS Testing)

Step 0:  
Buffer  
Grooming



Step 1:



Step 2:



Step 3:

```
char x = oracle[secret * 4096];
```

'P' = 0x50



256 different CPU Cache Line


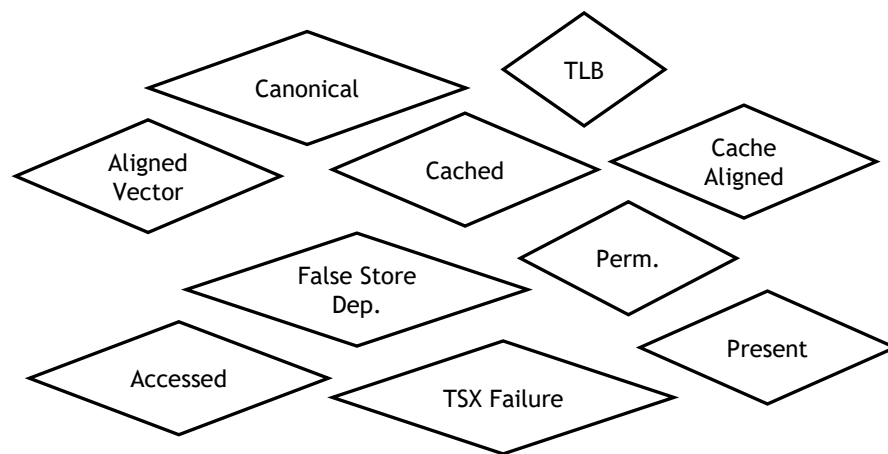

# Transynther (Fuzzing-based Random MDS Testing)

Step 0:  
Buffer  
Grooming

Step 1:

Step 2:

Step 3:




```
char x = oracle[secret * 4096];
```

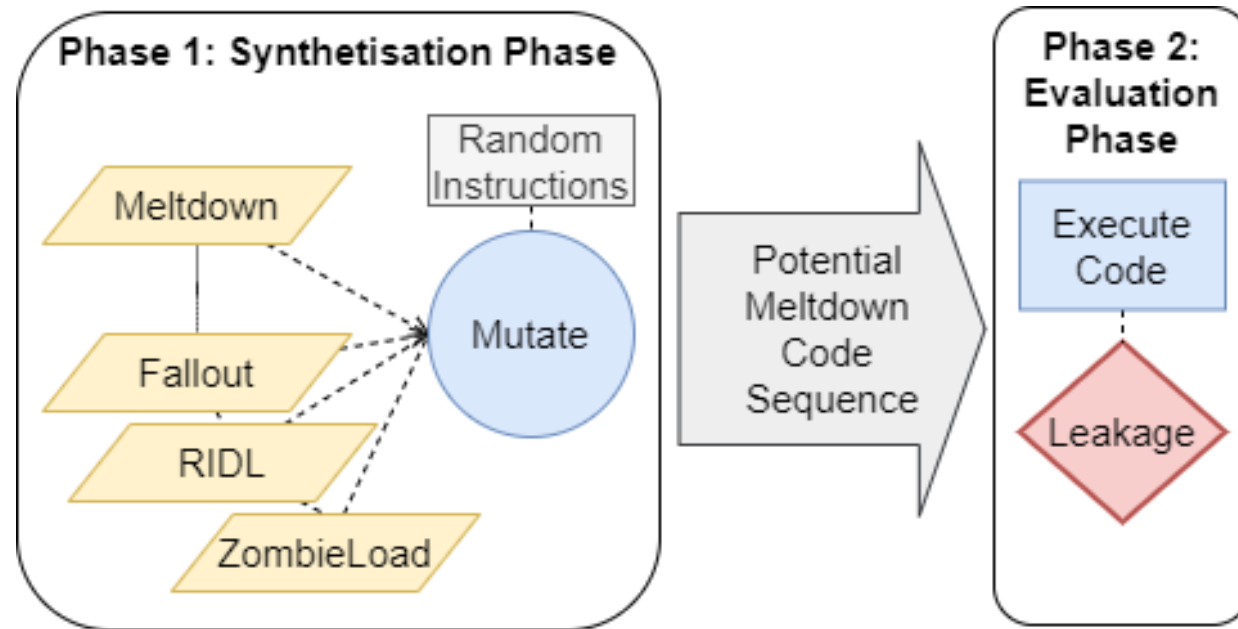
'P' = 0x50



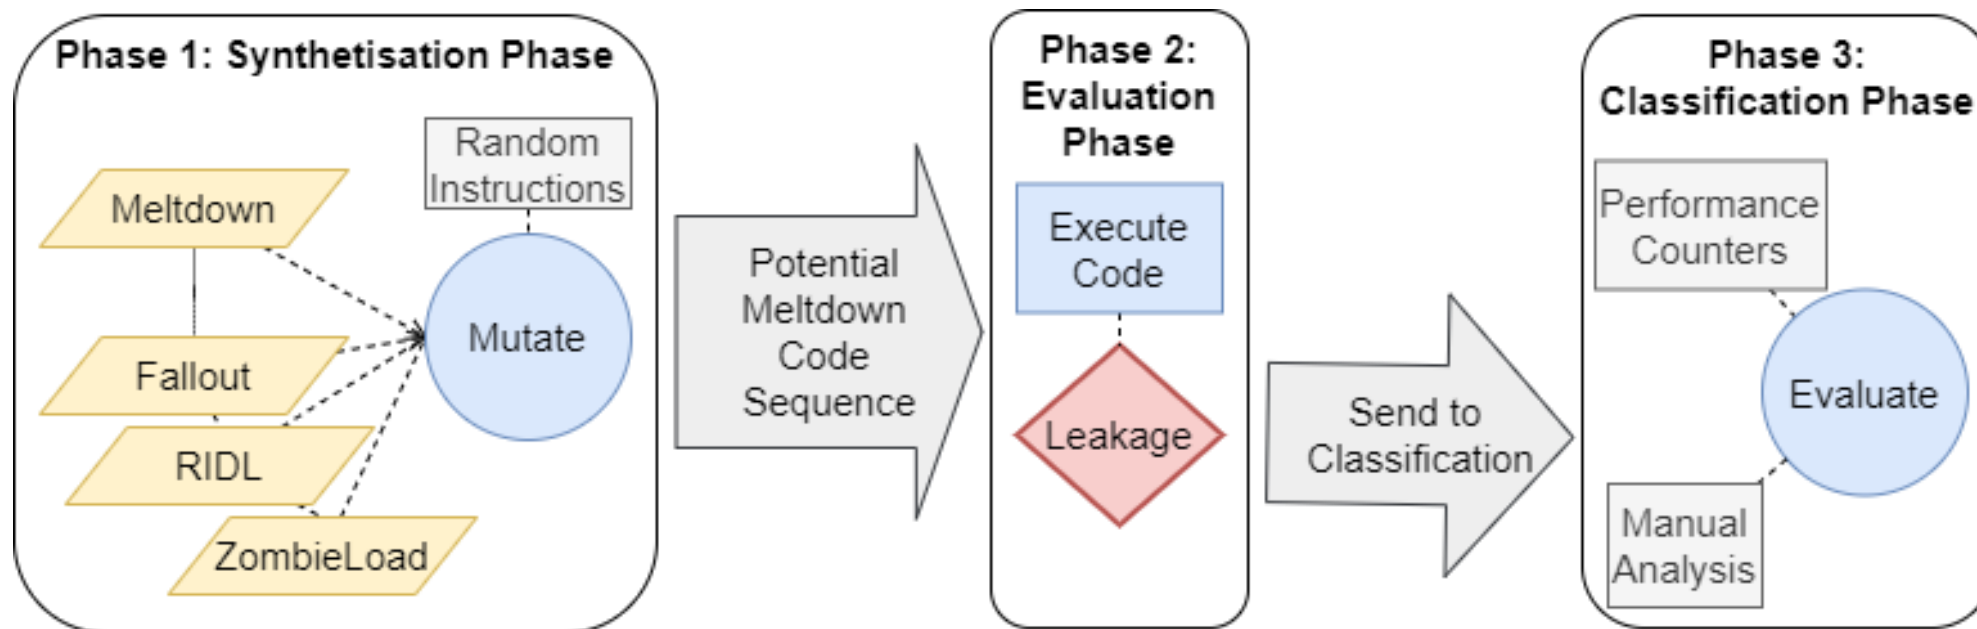
# Transynther (Fuzzing-based MDS Testing)



# Transynther (Fuzzing-based MDS Testing)



# Transynther (Fuzzing-based MDS Testing)



# Transynther Demo

Table 2: Leakage variants discovered by Transynther.

Case	Preparation	Store	Load	Name
①	(access $\emptyset$ , random instructions)	-	$\leftarrow + \text{lock} / \text{Access-bit Assist} / \emptyset$	MLPDS
②	(access $\emptyset$ , random instructions)	-	AVX $\leftarrow + \text{lock} / \text{Access-bit Assist} / \emptyset$	MLPDS
③	(access $\emptyset$ , random instructions)	-	AVX $+ \text{lock} / \text{Access-bit Assist} / \langle x \rangle$	Medusa
④	(access $\emptyset$ , random instructions)	-	AVX $\Rightarrow + \text{lock} / \text{Access-bit Assist} / \emptyset / \langle x \rangle / \checkmark$	Medusa
⑤	-	store (to load)	$\text{lock} / \text{Access-bit Assist} / \langle x \rangle / \checkmark$	S2L
⑥	(rep mov + store, store + fence + load)	store (to load)	$\text{lock} / \text{Access-bit Assist} / \langle x \rangle / \checkmark$	-
⑦	-	store (4K Aliasing) + $\text{lock} / \text{Access-bit Assist} / \emptyset / \langle x \rangle / \checkmark$	$\text{lock} / \text{Access-bit Assist}$	MSBDS
⑧	-	store (4K Aliasing, to load) + $\text{lock} / \text{Access-bit Assist} / \emptyset / \langle x \rangle / \checkmark$	AVX $\Rightarrow + \text{lock} / \text{Access-bit Assist} / \emptyset / \langle x \rangle / \checkmark$	MSBDS, S2L
⑨	(Sibling on/off)	store (random address) + $\emptyset$	$\text{lock} / \langle x \rangle$	MSBDS
⑩	(Sibling on/off + clflush (store address))	store (Cache Offset of Load) + $\emptyset$	$\text{lock} / \langle x \rangle$	MSBDS
⑪	(Sibling on/off + repmov (to Load))	store (to Load)	AVX $\Rightarrow + \text{lock} / \text{Access-bit Assist} / \emptyset / \langle x \rangle / \checkmark$	Medusa, MLPDS
⑫	-	Store (Unaligned to Load)	$\text{lock} / \text{Access-bit Assist} / \langle x \rangle$	Medusa
⑬	(random instructions)	AVX Store (to Load)	$\langle x \rangle$	Medusa, MLPDS, MSBDS
⑭	-	random fill stores	$\langle x \rangle$	MSBDS

$\langle x \rangle$  Non-canonical Address Fault     $\emptyset$  Non-present Page Fault     $\text{lock}$  Supervisor Protection Fault     $\Rightarrow$  AVX Alignment Fault

$\text{Access-bit Assist}$      $\leftarrow$  Split-Cache Access Assist     $\checkmark$  Access without fault or Assist



# Demo some Interesting Leakage Pattern

# MDS Attacks (ZombieLoad, RIDL, Fallout, ...)

- The CPU must flush the pipeline before executing an assist.
- Upon an Exception/Fault/Assist on a Load, Intel CPUs:
  - Execute the load until the last stage.
  - Flush the pipeline at the retirement stage (Cheap Recovery Logic).
  - Continue the load with some data to reach the retirement stage.
- Which data? (Fill buffer, Store Buffer, Load Buffer)
- Which one will be leaked first? (First come first serve)

# MEDUSA





# Medusa Attack

- Write Combining fills up the entire Data Bus.
- Medusa only leaks the Upper-half of the Data Bus.
- Implicit WC, i.e., 'rep mov', 'rep stos', can be leaked.
- Served by a Write Combining Buffer (or just the the Fill Buffer)
- Advantages:
  - Prefiltered data
  - Less Noise
  - More targeted (maybe also a disadvantage)

# Medusa Attack - V1 Cache Indexing


Cache Line Index



An invalid (Non-canon) address:  
0x5550000000000000008-20

# Medusa Attack - V1 Cache Indexing


Cache Line Index




An invalid (Non-canon) address:  
0x5550000000000000008-20

# Medusa Attack - V1 Cache Indexing

Cache Line Index



An invalid (Non-canon) address:  
0x5550000000000000008-20




# Medusa Attack - V1 Cache Indexing

Cache Line Index



Common Data Bus?!





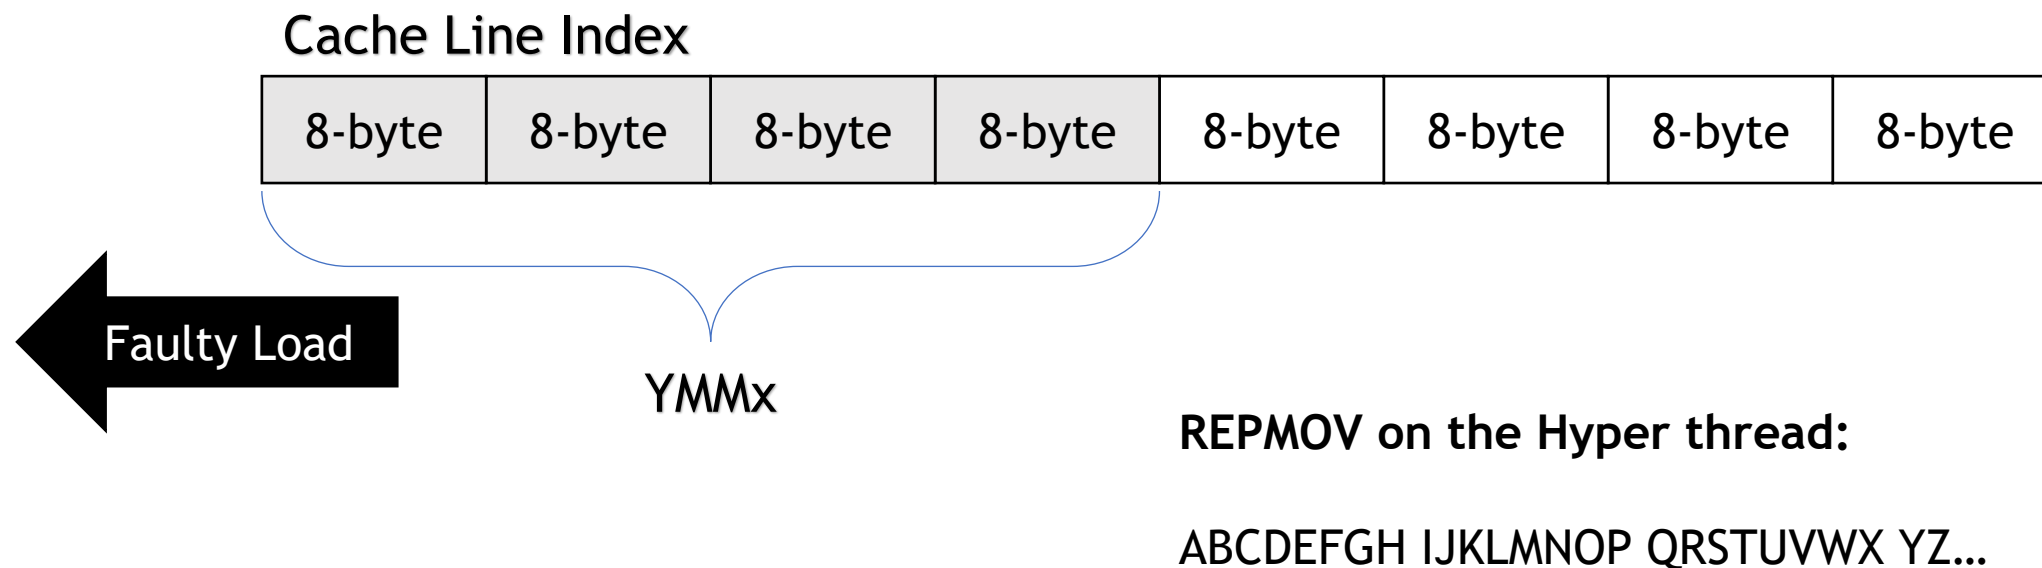
# Medusa Attack - V2 Unaligned S2L Forwarding

Cache Line Index

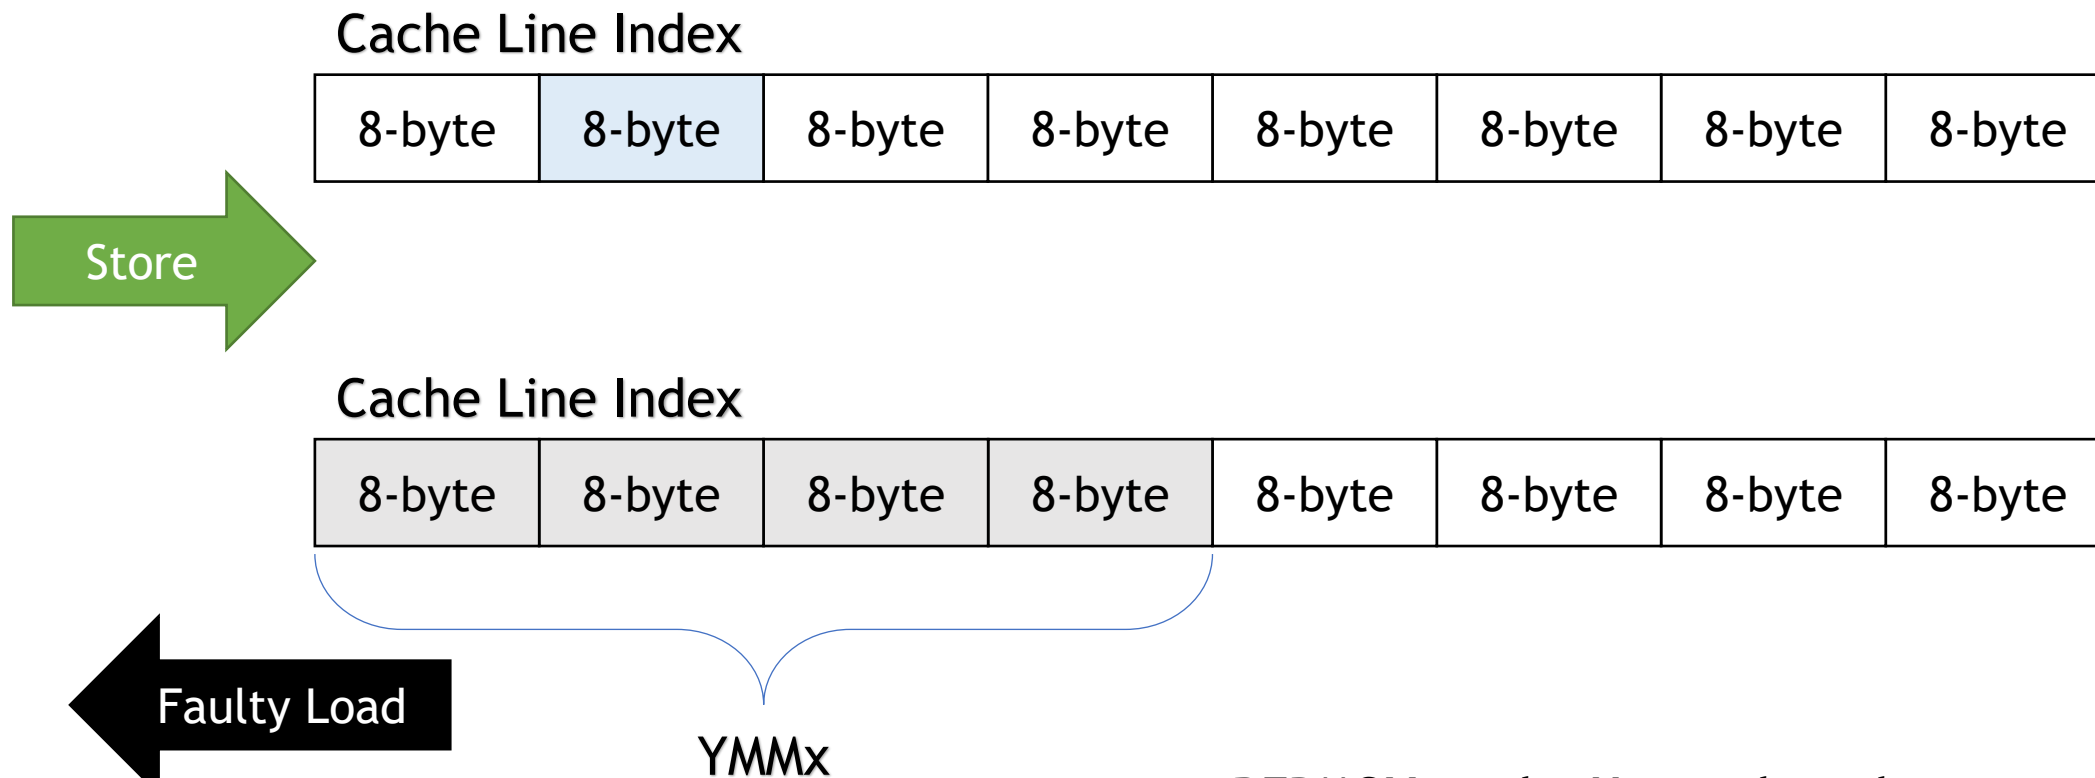
8-byte	8-byte	8-byte	8-byte	8-byte	8-byte	8-byte	8-byte
--------	--------	--------	--------	--------	--------	--------	--------



# Medusa Attack - V2 Unaligned S2L Forwarding



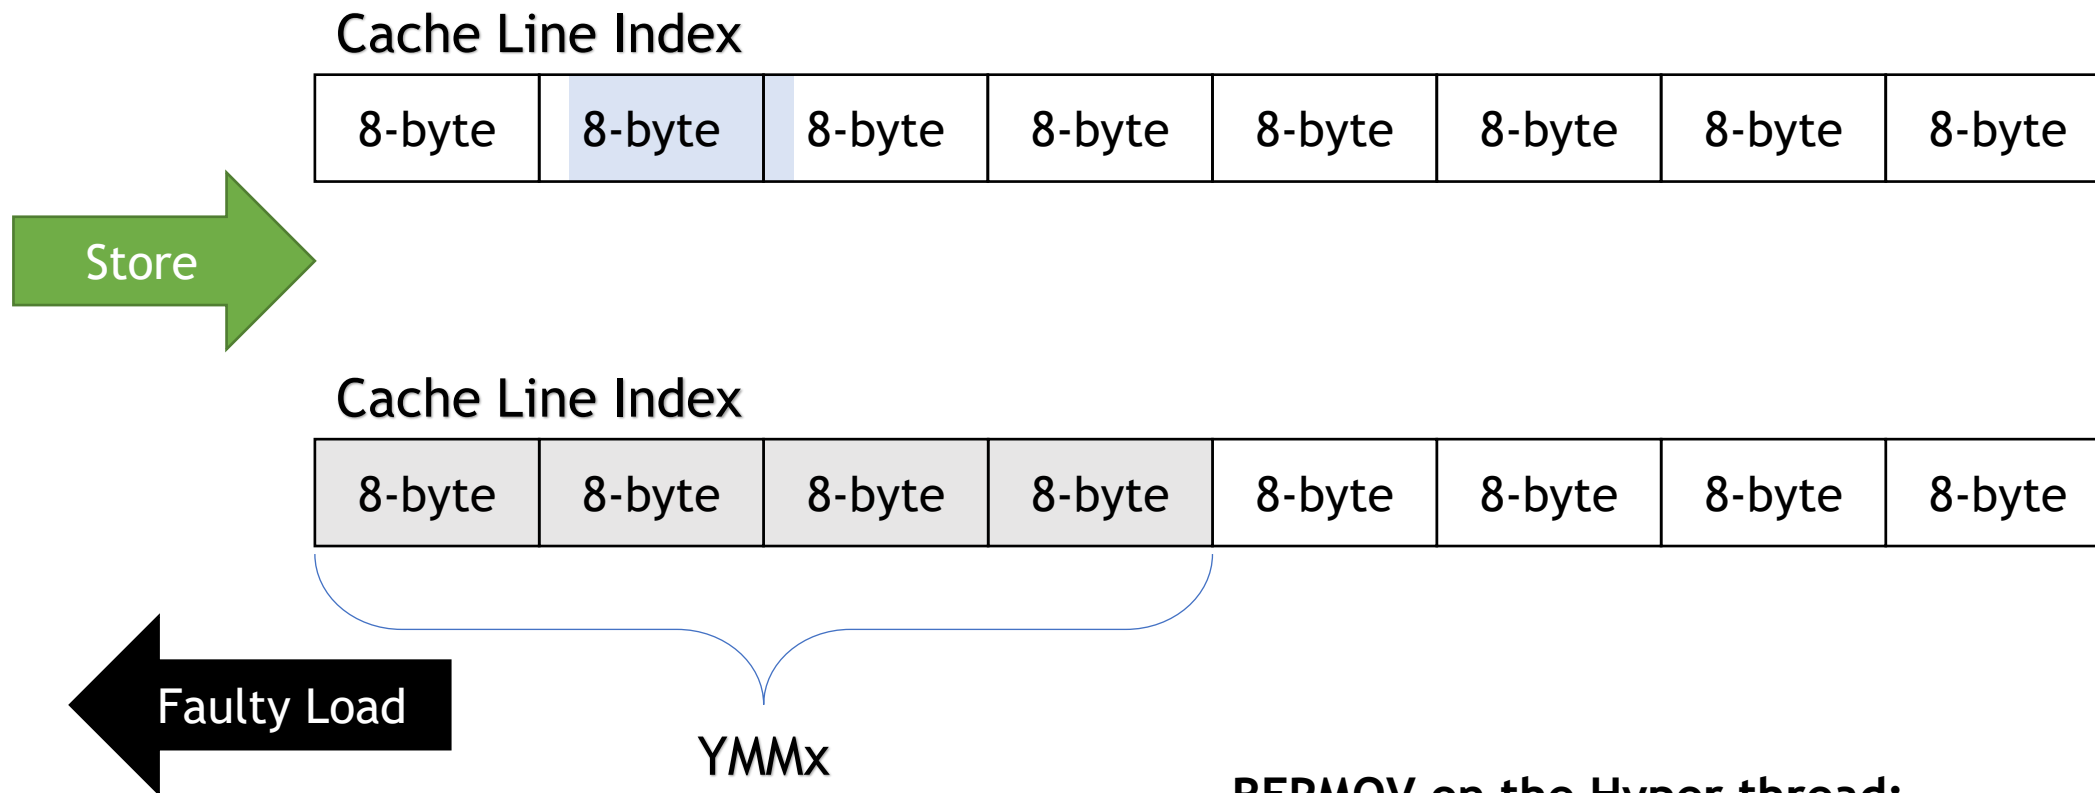
# Medusa Attack - V2 Unaligned S2L Forwarding



REPMOV on the Hyper thread:

ABCDEFGH **IJKLMNOP** QRSTUVWX YZ...

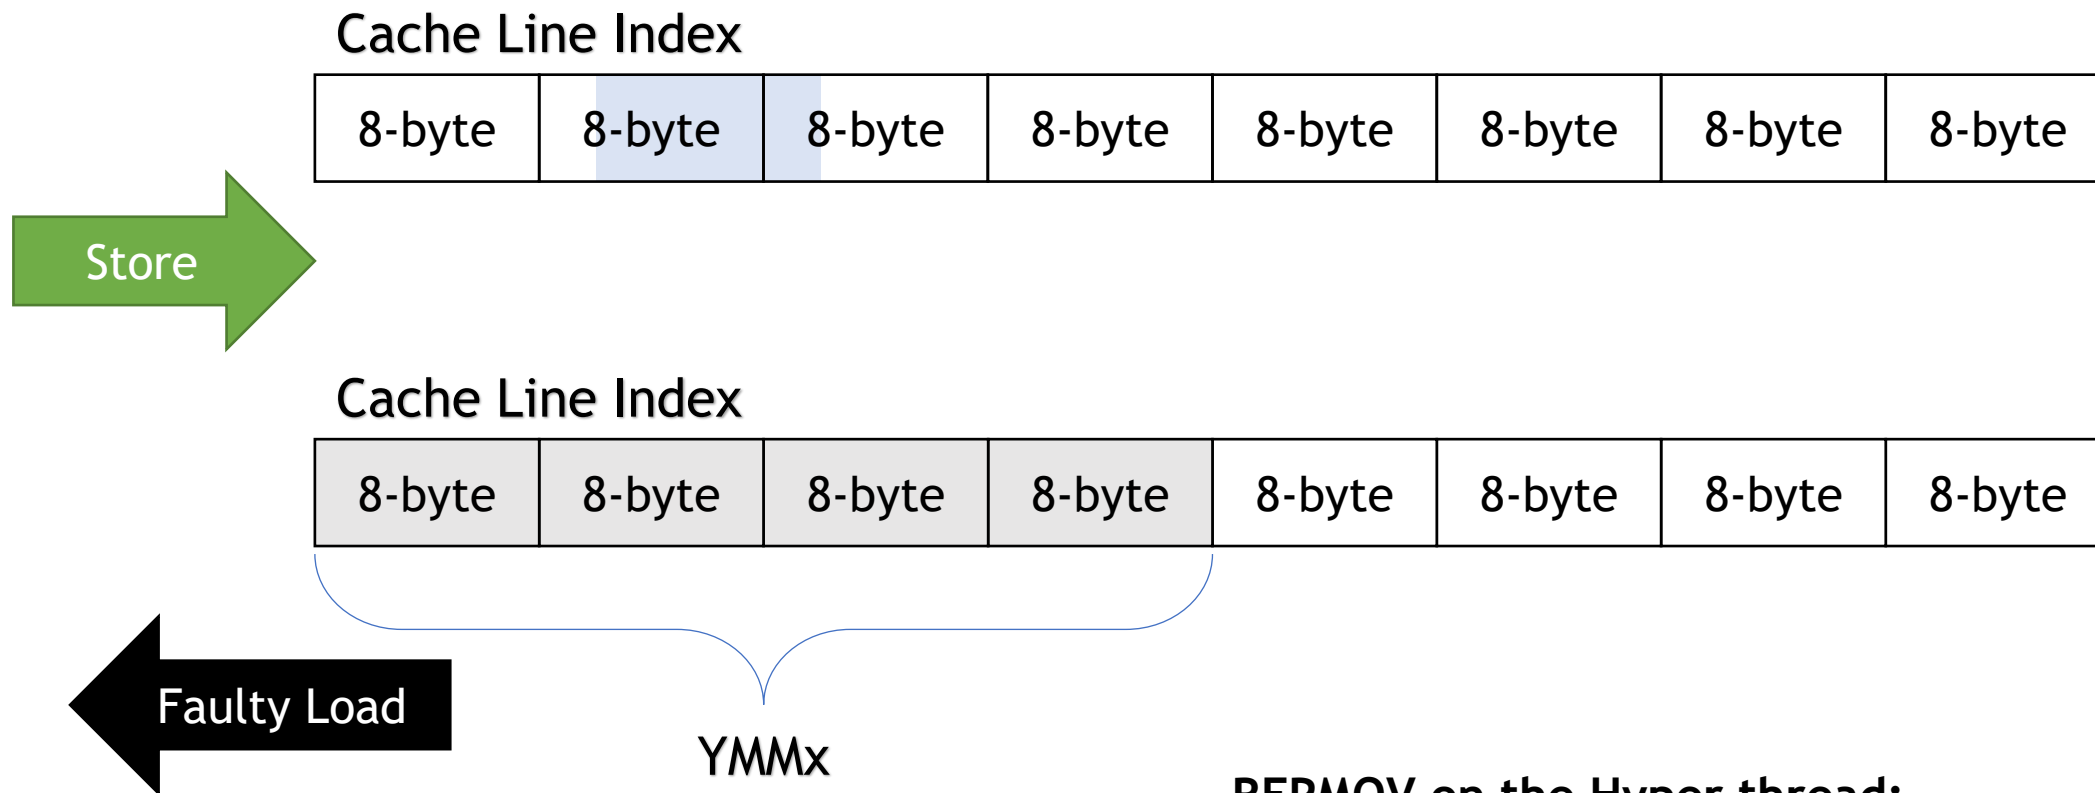
# Medusa Attack - V2 Unaligned S2L Forwarding



REPMOV on the Hyper thread:

ABCDEFGH IJKLMNOP QRSTUVWX YZ...

# Medusa Attack - V2 Unaligned S2L Forwarding



REPMOV on the Hyper thread:

ABCDEFGH IJ**KLMNOP** **QRSTU**VWX YZ...

# Medusa Attack - V2 Unaligned S2L Forwarding


Cache Line Index

8-byte	8-byte	8-byte	8-byte	8-byte	8-byte	8-byte	8-byte
--------	--------	--------	--------	--------	--------	--------	--------



Cache Line Index

8-byte	8-byte	8-byte	8-byte	8-byte	8-byte	8-byte	8-byte
--------	--------	--------	--------	--------	--------	--------	--------



# Medusa Attack - V3 Shadow *REP MOV*

- A *REP MOV* that faults on the load leaks:
  - the data from the legitimate store address
  - but also the data from the *REP MOV* running on the hyper thread

HT 1: *REP MOV*  
Valid Store, Faulty Load

AAAAAAAAA	A	AAAAAAAAA
AAAAAAAAA	A	AAAAAAAAA

MD Leak

HT 1: *REP MOV*  
Valid Store, Faulty Load

ABCDEFGHIJKLMNOP
AAAAAAAAAAAAAAAAAAAA

# Medusa Attack - V3 Shadow *REP MOV*

- A *REP MOV* that faults on the load leaks:
  - the data from the legitimate store address
  - but also the data from the *REP MOV* running on the hyper thread

HT 1: *REP MOV*  
Valid Store, Faulty Load

```
AAAAAAAAA AAAAAAAAAA  
AAAAAAAAA AAAAAAAAAA
```

MD Leak

HT 1: *REP MOV*  
Valid Store, Faulty Load

```
ABCDEFGH IJKLMNOP  
AAAAAAAAA AAAAAAAAAA
```

AAAAAAAAAAAA || AAA | A | AAA | A | A || AAAAAA...



# OpenSSL RSA Key Recovery

- OpenSSL Base64 Decoder uses inline Memcpy(-oS)
- Triggered during the RSA Key Decoding from the PEM format:

-----BEGIN RSA PRIVATE KEY-----

```
MIICXQIBAAKBgQDmTvQjttGtnlqMwmmaLW+YjbYTsNR8PGKXr78iYwrMV5Ye4VGy
BwS6qLD4s/EzCzGIDwkWCVx+gVHvh2wGW15Ddof0gVAtAMkR6gRABY4TkK+6YFSK
AyjmHvKCfFHvc9loeFGDyjmWFFkfdwzppXnH1Wwt00lNyCU1GbQ1w7AHuwlDAQAB
AoGBAMyDri7pQ29NBIfMmGQuFtw8c0R3EamlldQbX7qUguFEoe2YHqjdrKho5oZj
nDu8o+Zzm5jzBSzdf7oZ4qaeekv0fO+ZSz6CKYLbuzG2IXUB8nHJ7NuH3lacfivD
V4Cfg0yFnTK+MDG/xTVqywrCTsslkTCYC/XZOXU5Xt5z32FZAkEA/nLWQhMC4YPM
0LqMtgKzfgQdJ7vbr43WVVNpC/dN/ibUASI/3YwY0uUtqSjillghIY7pRohrPJ6W
ntSJw0UAhQJBAOe2b9cfiOTFKXxyU4j315VkulFfTyL6GwXi/7mvpcDCixDLNRyk
uRigmdKjtlUrAX0pwjgXa6niqJ691jExez8CQQCcMZZAvTbZhHSn9LwHxqS0SIY1
K+ZxX5ogirFDPS5NQzyE7adSsntSioh6/LQKBX6BAR9FwtXBPACtwz5F9geZAKA8
a3z0SlvG04aC1cjkgUPsx6wxxbl79F2RhmSKRbv7JiYk3RQ+L7vJgmWPGu5AcLM
oVPsjmbbkKfJZNTyVOW/AkABepEi++ZQQW0FXJWZ3nM+2CNcXYCtTgi4bGkvnZPp
/1pAy9rjeVJYhb8acTRnt+dU+uZ74CTtfuzUTZLOluVe
```

-----END RSA PRIVATE KEY-----

# OpenSSL RSA Key Recovery

- OpenSSL Base64 Decoder uses inline Memcpy(-oS)
- Triggered during the RSA Key Decoding from the PEM format:

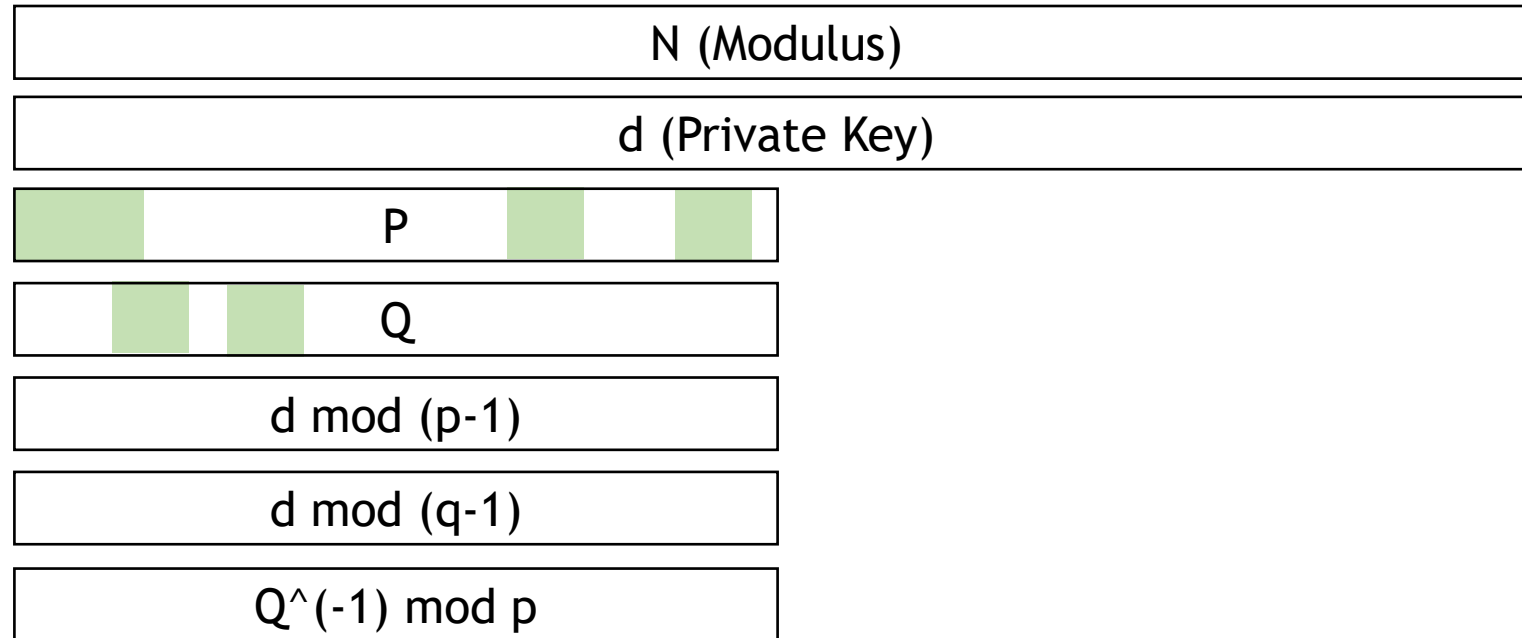
-----BEGIN RSA PRIVATE KEY-----

```
MIICXQIBAAKBgQDmTvQjttGtnlqMwmmaLW+YjbYTsNR8PGKXr78iYwrMV5Ye4VGy
BwS6qLD4s/EzCzGIDwkWCVx+gVHvh2wGW15Ddof0gVAtAMkR6gRABY4TkK+6YFSK
AyjmHvKCfFHvc9loeFGDyjmWFFkfdwzppXnH1Wwt00lNyCU1GbQ1w7AHuwlDAQAB
AoGBAMyDri7pQ29NBIfMmGQuFtw8c0R3EamlldQbX7qUguFEoe2YHqjdrKho5oZj
nDu8o+Zzm5jzBSzdf7oZ4qaeekv0fO+ZSz6CKYLbuzG2IXUB8nHJ7NuH3lacfivD
V4Cfg0yFnTK+MDG/xTVqywrCTsslkTCYC/XZOXU5Xt5z32FZAkEA/nLWQhMC4YPM
0LqMtgKzfgQdJ7vbr43WVVNpC/dN/ibUASI/3YwY0uUtqSjillghIY7pRohrPJ6W
ntSJw0UAhQJBAOe2b9cfiOTFKXxyU4j315VkulFfTyL6GwXi/7mvpcDCixDLNRyk
uRigmdKjtlUrAX0pwjgXa6niqJ691jExez8CQQCcMZZAvTbZhHSn9LwHxqS0SIY1
K+ZxX5ogirFDPS5NQzyE7adSsntSioh6/LQKBX6BAR9FwtXBPACtwz5F9geZAKA8
a3z0SlvG04aC1cjkgUPsx6wxxbl79F2RhmSKRbvH7JiYk3RQ+L7vJgmWPGu5AcLM
oVPsjmbbkKfJZNTyVOW/AkABepEi++ZQQW0FXJWZ3nM+2CNcXYCtTgi4bGkvnZPp
/1pAy9rjeVJYhb8acTRnt+dU+uZ74CTtfuzUTZLOluVe
```

-----END RSA PRIVATE KEY-----


# OpenSSL RSA Key Recovery

- OpenSSL Base64 Decoder uses inline Memcpy(-oS)
- Triggered during the RSA Key Decoding from the PEM format:



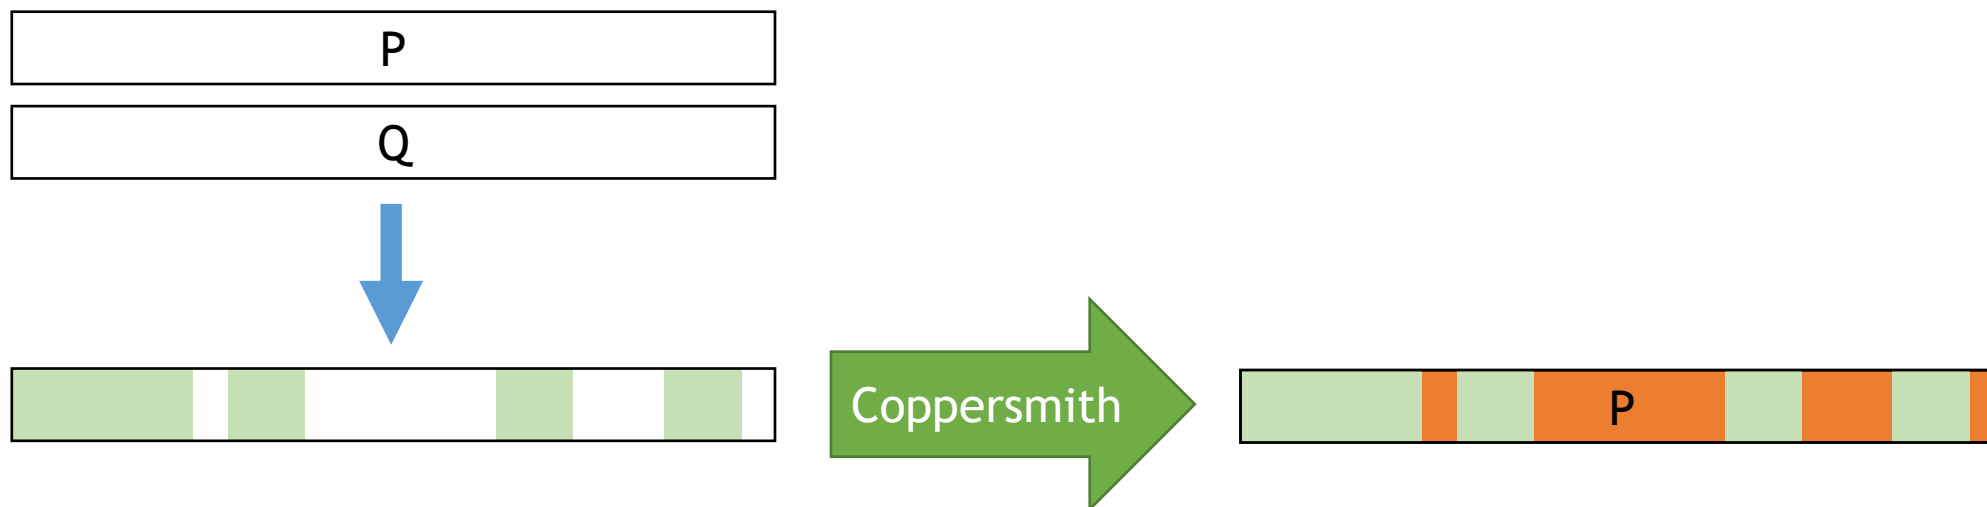
# OpenSSL RSA Key Recovery - Coppersmith

- Knowledge of at least  $\frac{1}{3}$  of  $P+Q$
- Create a  $n$  dimensional hidden number problem where  $n$  is relative to the number of recovered chunks
- Feed it to the lattice-based algorithm to find the short vector



# OpenSSL RSA Key Recovery - Coppersmith Attack

- Knowledge of at least  $\frac{1}{3}$  of  $P+Q$ .
- Creating a  $n$  dimensional hidden number problem where  $n$  is relative to the number of recovered chunks.
- Feeding it to the lattice-based algorithm to find the short vector.



# Conclusion

- Automated Testing for CPU Attacks,
  - helps us to understand the root cause of these issues better.
  - can be used to verify hardware mitigations (e.g., Fallout on ICL).
  - can help us to improve the leakage rate and understand the impact of attacks better.
- The impact of attacks depend also on the exploitation technique.
- Potentials and Future work:
  - Can we integrate such tools with feedback from hardware/simulator?


# Questions?!



**Daniel Moghimi**

@danielmgmi

- Moghimi, D., Lipp, M., Sunar, B., & Schwarz, M. (2020). Medusa: Microarchitectural Data Leakage via Automated Attack Synthesis. In 29th USENIX Security Symposium (USENIX Security 20).



# CPU Memory Subsystem

